

Individual Decision Making

Colin Camerer

I. Introduction

I will review experimental studies of individual decision making, with their implications for economists in mind. Decision making is increasingly important for economics for at least two reasons.

First, in many economic settings individuals make decisions by and among themselves: consumers save, sell their labor, buy houses and durable goods, form economic and social relationships, and bargain. In these cases the institutional veils separating people from others are thin. A couple decides whether to buy a house from another couple; Mario hires a college student to work in the grocery store he owns; a daughter borrows money from her mother. (In other settings the institutional veil separating individuals doing business is thick: Monique lends money to the shareholders of General Motors through the concrete veil of Citibank, where she has a savings account and GM has a line of credit.)

The thickness of institutional veil is important because there is a strong intuition that institutional forces correct errors people make; the more directly people trade with each other, that intuition implies, the more likely their errors are to persist. Economic analysis has increasingly reached into settings with thin veils recently. Judges are presumed to make law as if they had economic efficiency in mind; there are models in which people optimize: marriages, sleep, suicide, and extramarital affairs; the household is modeled as a unit of production; and so forth. In these settings systematic errors by individuals may not be corrected by institutional forces. Studies of individual decision making can help predict when market prices may be wrong and allocations inefficient and suggest ways to improve efficiency.

Second, economic analysis has also reached into increasingly complicated domains recently. Until thirty years ago there were few formal models with any uncertainties. Weak assumptions about agent rationality were adequate to generate strong market level results (e.g., Pareto optimality). Now many models presume agents can make choices under risk and uncertainty, over time, keeping in mind subtle game-theoretic effects. As the models grow more and more complicated, agents are assumed to have more and more rationality. Then it is more likely individual agents violate the models; studies like the ones I describe may tell us how and why.

A. Limited Rationality and Decision Research

For the last thirty years or so, most research on individual decision making has taken nonnative theories of judgment and choice (typically probability rules and utility theories) as null hypotheses about behavior, and tested these hypotheses in psychology experiments. Much of this work is called "behavioral decision research" (a term coined by Edwards [1961a]) or, sometimes, 'cognitive illusions' or "cognitive misperceptions." The goal is to test whether nonnative rules are systematically violated and to propose alternative theories to explain any observed violations.

The most fruitful popular alternative theories spring from the idea that limits on computational ability force people to use simplified procedures or "heuristics" that cause systematic mistakes (biases) in problem solving, judgment, and choice. The roots of this approach are in Simon's (1955) distinction between *substantive* rationality (the result of nonnative maximizing models) and *procedural* rationality—people behave coherently by following reasonable procedures but sometimes make suboptimal decisions as a result.

I. Why Study Errors in Decision Making?

Cataloguing systematic violations of rational models was not always the theme of the psychologists' efforts. In 1967, Peterson and Beach wrote a review of research on intuitive statistical judgment and concluded that people obeyed nonnative laws rather well. Psychologists began focusing on judgment errors in the 1970s because they thought judgment errors might reveal how people generally make judgments, just as optimal illusions tell us about perception and forgetting ellipses about memory (Kalmeman and Tversky 1982). The same scientific heuristic is used in other fields. The Great Depression, the stock market crash of 1987, and the savings-and-loan crisis are carefully studied for clues about the general behavior of economies and markets. Engineers study bridge collapses and airplane crashes to learn how to build sturdier bridges and planes.

Whether people make judgment errors frequently or not is difficult to judge and to most psychologists beside the point. Psychologists study errors because if people use simplified procedures to judge and choose, those procedures may be seen most clearly through the errors they cause. For economists, the frequency of errors is important because errors might affect economic efficiency, and methods for removing errors could be useful policy tools.

B. Two Controversies: Methods and Implications

Since many of the psychologists' studies can be seen as direct attacks on assumptions of individual rationality, the studies are sometimes hotly debated. There are two kinds of debates: methodology and implications.

The conventional *methods* used in psychological studies of decision making are often different than the conventions established by experimental economists (detailed throughout this handbook). In the psychology experiments, subjects are often not paid according to their performance, or are paid small amounts; stimuli have natural labels that may induce nonmonetary utilities; subjects do not always make repeated choices under stationary replication; treatments are sometimes created by deceiving subjects; and so forth. As a result, many economists discount evidence from the psychologists' studies. Replication of findings using the methods of experimental economics is therefore popular and tests robustness of results.

The *implication* of evidence of irrationality is another source of controversy. Despite the psychological evidence, economists have been cautious about reconsidering the presumption in their work that agents maximize choices based on well-informed preferences. Their caution is often defended by a tenet of "positive economics" (Friedman 1953): the market-level predictions may be approximately right even if the model of individuals from which the predictions are derived is wrong. Thus, better models of individual decision making *may not* improve market-level prediction; whether they do is fundamentally an empirical question that economics experiments help answer (e.g., Plott 1986). Experiments are helpful because they naturally give simultaneous observations of individual and aggregate activity, which are the best raw material for judging whether individual errors are present and important for aggregate behavior. There are a few studies of this sort, comparing individual and aggregate behavior within an experiment. They are reviewed below, in sections II (C,D,F.2) and III (I.3, J.1).

C. A Map and Guidebook

This chapter is organized in two sections: judgment (II) and choice (III). The study of judgments is almost purely psychological, *except* for a few replications and market studies by experimental economists. Studies of choice have had more interplay between axiomatic theories (mostly, though not exclusively, generated by economists) and experimental data gathered by economists and psychologists alike.

I have tried to weave the many studies by psychologists and the relatively few by experimental economists into whole cloth depicting classes of systematic mistakes and the procedures people use that seem to create the mistakes. I say the most about ongoing debates in which several studies have cumulated knowledge—aggregation of Bayesian errors in markets, utility theory, preference reversals, buying-selling price gaps. But in many places, economists have not joined the debate because they are not familiar with psychological results, do not appreciate their impact on economics, or think the results are unlikely to replicate. I try to remedy unfamiliarity by discussing a broad array of results in minimal detail, to encourage appreciation by providing recipes for expressing psychological

findings in theoretical terms familiar to economists, and to provoke replication with numerous suggestions for further research.

The methodological range of studies summarized in this chapter is perhaps as wide as in any chapter in the handbook. I sprinkled brief digressions about methodology throughout the chapter, at points where they illuminate debate and where the debate provides a context that adds flavor to an otherwise bland discussion.

Other sources include Thaler (1987), who presents much of the same evidence organized as a critique of economic tenets (and see his *Journal of Economic Perspectives* columns, collected in Thaler [1992]). Edited collections of important articles in behavioral decision theory are Kahneman, Slovic, and Tversky (1982). Ares and Hammond ([1987] and Beu, Raiffa, and Tversky (1988)). There are graduate level textbooks by Dawes (1988) and Hogarth (1987). Texts by Dzhurmal (1990) and Rumelhart and Schoemaker (1989) are easier. Yates (1990) has a series of articles in the *Annual Review of Psychology* (most recently Payne, Bettman, and Johnson [1992]) provide an authoritative overview of psychological decision research. The chapter by Abelson and Levi (1988) is a rough equivalent of this chapter, aimed at psychologists. New work on models of choice is reviewed by Machina (1987), Fishburn (1988), Weber and Camerer (1987), and Camerer and Weber (1992).

For psychology-economic nexus is covered by the book edited by Hogarth and Reder (1987) (reprinting a 1986 *Journal of Business* special issue, critically reviewed by Smith (1991)). Cox and Isaac (1986) cover a small patch of similar ground.

II, Judgment

A. Calibration

Good probability judgments should match actual relative frequencies. The match is shown in a "calibration curve."¹ For example, in 1955 the National Weather Service began requiring in meteorologists to announce numerical judgments of the probability of precipitation. Figure 8.1 shows a calibration curve for one forecaster, using actual forecasts from several days. On the Y axis is the relative frequency of events (proportion of days with precipitation) for each category of probability forecast shown on the X axis. The number of events in each forecast category is indicated by the size of each point (and written alongside it). The forecaster shown in Figure 8.1 said there was a 30 percent chance of rain on 160 days; it actually rained slightly more than 30 percent of those days.

Accuracy of probability judgments has two distinct components, calibration and resolution (sometimes called "calibration-in-the-small" and "calibration-in-the-large"). Calibration is how well the event forecasts (in a particular category (all events with 3 probability) matches the actual relative frequency of those events (50 of 160 occurred). In a calibration curve like the one shown in Figure 1U, calibration is measured by how close points are to the identity line (adjusting for




Figure 1U, Calibration graph for forecaster A. A probabilistic forecast of precipitation from the Chicago area. Numbers are relative frequencies. Source: Murphy and Winkler 1977.

sampling error). Resolution (also called "discrimination") is how well probabilities enable one to discriminate between likely and unlikely events. A high-resolution forecaster will have many forecasts in the extreme categories near zero and one. When making predictions is difficult--in long-term economic forecasting, for example--resolution may only be achieved at the expense of calibration, by confidently making high and low guesses that are only partly right.

The judgments of the weather forecaster in Figure 1U show terrific calibration (the points are close to the line) and good resolution (most of the observations are between zero and .2). Calibration as good as the weather forecasters' seems to be rare (see Uchtemstein, Fischhoff, and Phillips 1982). Some empirical calibration curves, based on students' judgments, are shown in Figure S.2. These are "half-range" curves: from two possible answers to a general interest question--did potatoes come from Ireland or Peru?--subjects pick the more likely answer and judge its probability (which must be at least .5).

In general subjects are overconfident. They are insufficiently regressive in judging the likelihood of events. Events they say are certain happen only 80 percent of the time. "Full range" curves, with subjective probabilities from zero to one, show overconfidence too. (Events judged to be impossible happen 20 percent of the time.) However, subjects are often underconfident when questions are easy (i.e., when the percentage of people answering the questions correctly is high).




Figure 82. Calibration for half-range knowledge. Reprinted from Lichienst(1, Hsdihoff, and Phillips (1992) with permission of Cambridge University Press. For further details, see *ibid.*

L. Scoring Rules

In most of the studies above, subjects were not directly rewarded according to the accuracy of their probabilities. "A proper scoring rule" is a scheme that rewards truthful judgments, depending on the judgment and the outcome of the event being forecasted, in a way that induces truthful revelation of probabilities. An example is the quadratic scoring rule: If the subject reports a probability p , pay her $\$(2p - p^2)$ if the event occurs and $\$(1 - p^2)$ if the event doesn't occur.³ (Or define $p = 1$ if the event occurs and $p = 0$ if not, and pay $\$(1 - (p - ti)^2)$.) For example, a subject who reports $p = .10$ earns either $\$.19$ if the event occurs or $\$.99$ if not; a report $p = .20$ earns $\$.36$ or $\$.96$. If the subject thinks the true probability is $.20$, then the $p = .20$ bet earns an expected value of $.2 (\$.36) + .8 (\$.96) = \$.92$; and the $p = .10$ bet earns $\$.83$, so the subject should report $p = .20$.

Besides being incentive-compatible, scoring rules enable judgments of probability to be elicited without mentioning the word "probability" or defining it.

Instead, a Subject expresses a probability judgment implicitly, by choosing among various bets. (Scoring rules are sometimes used to grade students for probabilistic answers to multiple-choice questions.)

However, scoring rules assume risk neutrality; if a subject is risk averse her expressed probabilities will be biased toward .5. (Allen [1987], suggests paying subjects in lottery tickets; cf. the discussion of the binary lottery procedure in chapter 1.) And the payoff function has a flat maximum around the true subjective probability p , as the example above indicates, so subjects are not penalized much for misreporting.

The calibration studies described in this section did not use proper scoring rules. However, there seems to be little difference between judgments motivated by scoring rules and unincentivated judgments (Beach and Phillips 1967; Jensen and Peterson 1973). The main difference is that when subjects use extreme probabilities too frequently (and inappropriately), scoring rules punish their mistakes severely and reduce them (Fischer 1982). When rewarded with "improper" scoring rules that do not penalize misreports, subjects learn to misreport probabilities (Nelson and Bessler 1989). For example, suppose subjects who report an event probability p are paid p if the event occurs and $1 - p$ otherwise. Then subjects quickly learn to exaggerate their beliefs, reporting $p = 1$ if their true belief is above .5 and reporting 0 otherwise.

Scoring rules could be useful in a wide range of economics experiments to facilitate probabilistic beliefs in an incentive-compatible way. For example, game theories often make sharp predictions about beliefs that are difficult to test indirectly (e.g., beliefs after "out-of-equilibrium" events that should never occur and that actually occur only rarely). Some researchers have used scoring rules to elicit beliefs of subjects in games. The only published economics experiment that uses them, that I know of, is McKelvey and Page (1990).

2. Confidence Interval

Other studies have elicited confidence intervals for quantities (the length of the Amazon river, next month's spot oil price), instead of probabilities for events. In these studies confidence intervals are typically too narrow; subjects seem to anchor on a point estimate, then adjust upward and downward by too little. Fifty percent intervals included the true quantity only about 30 percent of the time; 98 percent intervals, only 60 percent of the time (Alpert and Raiffa 1982). Subjects can learn to widen their intervals out with intensive feedback and training, but they never get high probability intervals quite wide enough.

Many studies have examined the effects of expertise on overconfidence. A lot of these studies were motivated by the impressive performance of weather forecasters, shown in Figure 8.1. Researchers then became curious whether other experts were equally well calibrated. Professional accountants' intervals around estimated account balances of client firms are good (Tomassini et al 1982). Weather forecasters' intervals of high and low temperatures are precisely the right

width (Murphy and Vinkler 1974). But intervals around estimates of physical constants, published in physics journals, are systematically too narrow (Henrion and Mschhoff 1984).

There are mixed effects of expertise in studies of event-probability calibration too. Students and professionals forecasting outcomes of basketball and baseball games are poorly calibrated (Yates 1982; Rosis and Yates 1987). Novices, statistical experts, and blackjack dealers are equally well calibrated (Keien 1988) and expert bridge players are better calibrated than novices (Keren 1987) at judging the probability of winning a hand given certain cards. Physicians are accurate in some settings and poor in others, especially diagnosing rare diseases (Yates 1990, table 4.1; see the discussion of base-rate fallacy in section LC.1 later). Betting odds at horse racing tracks are well calibrated, with a slight but persistent tendency to overestimate the chance that longshots will win and underestimate the chances of favorites (Ziemba and Hausch 1986). (Curiously, the opposite betting pattern occurs at Hong Kong racetracks; see Busche and Hall 1988.) Forecasts by professional economists of the date of economic downturn are pretty well calibrated one quarter ahead, but the calibration gets much worse as the forecast horizon extends out to four quarters (Braun and Yaniv 1992).

There are a few cross-cultural studies of calibration. Asians seem to have high resolution they use extreme probabilities a lot but are very badly calibrated (Wright et al. 1978; Yates et al. 1989). Some psychologists think differences in the role of chance and bravado in Asian and Western philosophies and culture might account for the differences.

Recent studies found an important difference between "local confidence," the appropriateness of a single confidence interval for a single quantity, and "global confidence," the fraction of several intervals that contain their true quantities. Subjects were university employees and students who were asked ten questions about local operation (e.g., what is the current value of the university's land holdings?). The subjects' 90 percent confidence intervals were too narrow (as usual) but their global confidence was not bad: they guessed that about five often intervals contained true quantities, when only three of ten actually did (Sniezek and Buckley 1991 see also May 1986). These results suggest an important difference between the psychological process of constructing a judgment about a single quantity (or event) and making a collective guess about several such judgments. Most of us are probably overconfident about the chance of publishing our next article in a leading journal or teaching a brilliant class tomorrow, but are more level-headed about how many of our next ten articles or ten classes will be similarly successful.

The pervasive finding that subjects are (locally) overconfident may have important economic implications. If people underestimate the width of distributions of future quantities, they will underinvest in flexibility and insurance. This might have implications for equilibrium models of rental and ownership of housing, choices of mortgage terms (adjustable vs. fixed-rate), marriage and divorce rates, managerial investments in manufacturing flexibility, and so on. Underestimation of variation might help explain why so many small businesses fail

of insufficient cash flow (stemming from overly narrow planning, perhaps; cf. Kahneman and Lovallo 1993).

Recent studies of calibration and confidence have rekindled debate along three lines. The first idea is that part of the apparent overconfidence could be caused by probability judgments that are correct on average but contain error (Erev, Wallsten, and Budesru 1992; Soll 1993). The second claim is that calibration researchers may have selected sample questions nonrandomly, oversampling "vicarious" questions in which natural cues yield the wrong answer (much as the Peru-Ireland potato question), and hence producing more overconfidence than is present in natural settings (Gigerenzer, Hoffrage, and Kleinbolting 1991; Justin, in press). Some new studies sample questions differently and reduce apparent overconfidence, but Griffin and Tversky (1992) and Soll (1993) sampled randomly and still observe overconfidence.

Third, Griffin and Tversky (1992) suggest a framework to organize many empirical results on confidence results. They point out that evidence has both strength (or extremeness) and weight. In several studies they find that judgments of confidence overemphasize the strength of evidence (compared to a Bayesian probability benchmark) and underemphasize its weight. Their framework can explain the observed difference in calibration for hard and easy questions (people underweight the strong weight of evidence in easy questions), conflicting results on expert calibration (experts will be highly overconfident in unpredictable environments, when they overweight weak evidence), and predicts some other phenomena.

B Perception and Memory Biases

Machines are natural metaphors and benchmarks for human perception and thinking. The metaphor of man as an information-processor now dominates cognitive psychology (e.g., Lachman, Lachman, and Butterfield 1979). It has proved fruitful by suggesting coherent theory and many empirical tests. Can people order events as memory does? Are memories stored like films in a library? Does information processing proceed in steps like a computer program?

However, much evidence suggests that human perception deviates systematically from the camera benchmark and memory deviates from the computer benchmark. (My goal in this very brief section is to inform readers about some shreds of evidence, to whet their appetites, and to suggest ways the data might matter for economists.) For example, Bruner, Pietschman, and Rodrigues (1951) showed subjects glimpses of playing cards in which colors and shapes were deliberately misrouted—hearts were black instead of the familiar red. Subjects thought they saw the familiar cards (red hearts). Error of this sort is systematic, not random: people more often err by mistaking unfamiliar patterns for familiar ones than vice versa. Put more formally, errors in absorbing information appear to be correlated with how unusual the information is. Misperception of surprising events implies that agents will misperceive outliers that signal regime switches or turning points in a time series. Their expectations will not be rational (in the sense of efficiently

using available information) because the process of new information depends on the stock of old information, or familiar images.

There are many biases in memory too. When guessing which cancer claims more lives or which journal to submit an article to, people rumple their memories. Sampling memories is a natural and reasonable heuristic because our memories are a sample of life. But even if our life sample is random, the sample we retrieve from memory will not be random because memories are not equally retrievable or "available" (Frederick and Kahneman 1973). For example, the most pleasant and pleasant memories are more easily remembered, which creates illusory nostalgia (Holmes 1970). Personal and concrete experiences are often overweighed (Sisbett et al. 1976). For example, Kunreuther et al. (1978) found that the purchase of earthquake insurance rose after a quake (though the probability of a subsequent large quake actually falls, because stress on the fault line is relieved). The availability of personal experiences is thought to create "egocentric" biases in judgments of fault (Bosch and Sicoly 1982); or two sides in an experimental dispute both think a judge's settlement with favor them [see Babcock et al. in press]. Memorable media reports cause biases in judgments because media coverage is not random⁷ (e.g., Greenberg et al. 1989). For example, Combs and Slovic (1979) found that newspapers vastly overreport accidents compared to diseases, and people think deaths from disease and accidents are equally common. (In fact, deaths from disease are 15 times more common.)

Availability can limit imagination and make theories, lists of words, or "fault trees" appear more complete than they really are. In a study by Fischhoff, Slovic, and Lichtenstein (1978), students and automobile mechanics undervalued the probability of "other causes" in an incomplete fault tree listing reasons why a car would not start. Similar biases in imagining contract contingencies might lead contracts to appear overly incomplete.

C. Bayesian Updating and Representativeness

When the probabilities people judge are conditional as to updating belief in X after learning M , they should follow the prescription of Bayes' rule:

$$P(X|M) = \frac{P(M|X)P(X)}{P(M)}$$

Computing probabilities using Bayes' rule is complicated. People seem to use simple heuristics instead: they anchor on $P(X)$ and adjust it to reflect M ; or they judge $P(X|M)$ by how "representative" X is of M (Tversky and Kahneman 1982).

Representativeness will be a useful heuristic because representative values are generally more common than unrepresentative ones. (Eagles are less representative of the set of birds than robins, and less common.) But judging likelihood according to representativeness neglects some features that are normatively important according to Bayes' rule—including the base rates $P(X)$ and $P(M)$. Similarly,

properties, and regression effects. Other features that are not nonnatively important loom large in representativeness-based thinking. Representativeness therefore creates several systematic departures from Bayesian judgment, or biases.

L. Underweighting of Base Rates

A famous problem used to study Bayesian judgment was introduced by Kahneman and Tversky (1972):

A cab was involved in a hit and run accident at night. Two cab companies, the Green and the Blue, operate in the city. You are given the following data:

- (a) 85 percent of the cabs in the city are Green and 15 percent are Blue,
- (b) a witness identified the cab as Blue.

The court tested the reliability of the witness under the same circumstances that existed on the night of the accident and concluded that the witness correctly identified each one of the two colors 80% of the time and failed 20% of the time. What is the probability that the cab involved in the accident was Blue rather than Green?

In the experiments, subjects are given the problem exactly as written above, often as part of a package of problems. Their probability judgments are recorded, and they are paid a small sum for participating (or course credit, in some cases).

The median and modal response is .80. It appears that subjects think the witness' judgment is representative of the actual color of the cab, and its representativeness leads them to confuse $P(\text{identify Blue} | \text{Blue})$ with $P(\text{Blue} | \text{identify Blue})$. According to Bayes' rule, the posterior probability that is asked for, $P(\text{Blue} | \text{identify Blue})$, should reflect the base rate $P(\text{Blue}) = .15$ also; but the base rate plays no role in the logic of representativeness. When the base rate is included too, the correct posterior probability is .41.

In these problems, and others like them, base rates are usually underweighted and often entirely neglected. Studies show that when attention is drawn to base rates, by varying the base rate in several versions of the problem or presenting them in causal forms (15 percent of the cab accidents in the city involve Blue cabs), subjects take base rates into account but still underweight them (Ajzen 1977; Bar-Hillel 1980a; cf. Koehler 1989).

The cab question is typical of stimuli used by psychologists to study judgment. Word problems describing natural events are used to escape from the limits of earlier traditions that emphasized more abstract stimuli, so the sensible presumption that psychological processes people use in everyday life could be better understood by asking people questions drawn from everyday life. The use of the word problems raises some methodological concerns for both economists and psychologists. For example, economists might wonder whether base rates neglect

affects asset prices in a market; some studies answering this question are reported in section II.G below. We return to the methodological concerns after describing a replication of the base rate studies.

Grether (1980) studied base rate neglect in an abstract setting with three bingo cages. A draw from the first cage (whose contents were known) determined whether state A or state B had occurred. The first cage had four N balls and two Gs; the B cage had three Ns and three Gs. Subjects observed a sample of six draws from whichever cage had been chosen (A or B) and were asked to decide which cage was more likely. For example, a subject might observe a sample of one N and five Gs, then choose whether to bet that the draws came from the A or B cage. The process was repeated several times, with fresh samples each time. At the end of the first trial was picked and a subject earned \$10 if they had picked the right cage on that trial.

Using logit estimation, Grether found that subjects weighted the rates $P(A)$ and $P(B)$ less than the likelihoods $P(\text{sample} | A)$ and $P(\text{sample} | B)$, as representativeness predicts, but they did not ignore the base rates entirely. Subjects also thought $P(A | \text{sample})$ was especially high when the sample was four Ns and two Gs, exactly matching the contents of the A cage (and similarly for the sample matching the B cage). Previous experience with a particular sample, or experience combined with monetary incentives for accuracy, reduced representativeness bias slightly but did not eliminate it. Concerned that Grether's subjects were not properly motivated, Harrison (1989b) replicated Grether's experiment with a variety of financial incentives. He found little evidence of representativeness among subjects with experience or financial incentive. There is no obvious way to reconcile the disagreement between his results and Grether's.

Grether (1991) extended his earlier work in three ways. In one experiment he was able to bound the degrees of belief in random events by having subjects choose between a bet on the most-likely cage and a bet on a chance device (so that choices revealed whether beliefs were higher [$> .75$] and lower [$< .25$]). In a second experiment he elicited probability judgments with a variant of the incentive-compatible Becker, DeGroot, and Marschak (1964) procedure (see chapter 1 for a description). Choices in both experiments were affected by representativeness. Probabilities elicited with the BDM procedure in the second experiment were often far too low or too high, but on average they were fairly close to Bayesian (within .05 to .10).

In a third experiment the A and B cages each had ten balls and samples of four balls were drawn. Assuming four ball samples cannot be representative of ten ball cages, representativeness should not affect judgments. In this experiment, judgments were quite different from the first two experiments—sample information was underweighted rather than overweighted (see the next section on conservatism). Grether concluded: "This [difference] suggests that in making judgments under uncertainty individuals use different decision rules in different decision situations." a "contingent-judgment" hypothesis espoused by many psychologists, (e.g., Payne, Bettman, and Johnson 1992).

A Digression on Methodology: Psychology and Economics

Grether's experiments are designed to address many criticisms some economists have of methods used by some psychologists. The Bayesian judgment problem was operationalized using physical devices (bingo cages) rather than a vignette like the cab problem. Subjects made choices rather than simply reporting probabilities; they were paid \$10 for one of their choices, randomly selected, was chosen. (In Grether [1991], a typical error cost 5 to 20¢.) Subjects made repeated choices, with an opportunity to learn; in the psychology studies, subjects often answer each question once because the purpose of the experiment is to study initial intuitions, not learning. The existence of some errors was reasonably robust to all these changes in conditions in Grether's data, but not in Harrison's. Incentives also reduced the number of incoherent and outlying responses (Grether 1981; cf. Smith and Walker 1993).

The difference between psychological and economic experiments should not be overstated. In the 1960s, long before Grether's work, psychologists and others used random devices to study judgment (Edwards 1968) and used the BDM procedure to study valuations (Lichtenstein and Slovic 1971). Even recently, there is substantial overlap across disciplines in methods, and substantial variation within disciplines. However, the typical differences in methods are worth analyzing because they usually follow from different background presumptions about human nature and different target domains investigators hope to generalize to. It is presumptuous to argue that either general method is superior.

For example, many psychologists are curious whether people can recognize and apply statistical rules to everyday situations, like the cab problem in which statistical structure is not transparent. They often use vignettes or problems drawn from natural settings (rather than problems based exclusively on random devices) because (1) they want to learn how people reason about natural events and (2) they think people may reason differently about events and about random devices. Given these interests and presumptions, word problems are well suited to doing their research and bingo cages are not. Economists are interested in different questions (not how people reason but whether people violate Bayes's rule) and are also more inclined to presume that reasoning about bingo cages and taxis is similar. For these purposes, cages and dice are better because they lay bare the statistical structure (making detection of a Bayesian error clear) and are presumed to be good substitutes for word problems.

Another area of typical difference is financial motivation of subjects. Psychologists do not always motivate subjects financially; though many have and a few are adamant about doing so—because incentives usually complicate instructions and psychologists presume subjects are cooperative and intrinsically motivated to perform well (Nawroth stimuli are also thought to keep subjects mentally involved and raise their intrinsic motivation, which substitutes for financial motivation.)

Repetition is another area of typical difference. The psychologists' tasks are often not repeated, with stationary replication, because psychologists are often most curious about initial behavior in a complicated environment. In addition,

many psychologists think stationary replication overstates the frequency, speed, and clarity of feedback the world actually provides. Economists tend to think oppositely: they are mostly curious about equilibrium behavior-the last period, not the first-and they think extensive laboratory feedback is the best time-compressed imitation of the strong learning forces present in natural settings.

To reiterate, there is substantial overlap in the way psychologists and economists do experiments. When their methods do differ, very roughly speaking, psychologists use natural stimuli, do not pay subjects, and do not repeat tasks. Economists pay subjects, prefer blandly labeled random devices as stimuli, and insist on repeating tasks. My view is that these different methods are preferred by different investigators because they effectively produce answers to different questions. Broad-minded students of individual decision making should have a healthy tolerance for variety in methods. (And variation in methods is essential to gathering data, to determine whether different methods do affect behavior substantially.)

It is worth noting that judgment errors, like those revealed in the cab problem, have been a lively topic of research within psychology too (e.g., Cohen 1981). Many of the arguments made in that literature are like those economists have made about methods or interpretations of results. For example, Gigerenzer, Hell, and Blank (1988) used physical devices to operationalize base rates and found some reduction in base rate neglect (though Grether, and others mentioned later, also found substantial base rate neglect using physical devices).

A more interesting argument is that some apparent biases might occur because the specific words used, or linguistic convention subjects assume the experimenter is following, convey more information than the experimenter intends.¹⁰ An example is the famous "Linda problem" (Tversky and Kahneman 1983). Subjects are told the following:

Linda is 31 years old, single, outspoken, and very bright. She majored in philosophy. As a student, she was deeply concerned with issues of discrimination and social justice, and also participated in anti-nuclear demonstrations.

Then they are asked to rank several statements about Linda by their probability:

Linda is a teacher in elementary school.

Linda works in a bookstore and takes Yoga classes.

Linda is active in the feminist movement (F).

Linda is a psychiatric social worker.

Linda is a member of the League of Women Voters.

Linda is a bank teller (T).

Linda is an insurance salesperson.

Linda is a bank teller and is active in the feminist movement (F&T).

Any ranking of probability should satisfy the conjunction law: Linda is less likely to be a feminist bank teller (marked F&T), than to be a bank teller (T) or a feminist (F), since the event F&T is a conjunction of the events F and T. In fact,

about 90 percent of subjects exhibit a conjunction fallacy, ranking the event F&T as more likely than one (or both) of the events F and T, usually T. (In a sample of well-trained Stanford decision sciences doctoral students, 85 percent made the same mistake.) The standard psychological explanation is that the description of Linda is more *representative* of a feminist bank teller than of a bank teller; subjects mistakenly think it is therefore more *likely* that Linda is a feminist bank teller.

The potential linguistic problem is this: in the presence of the statement "Linda is a feminist bank teller," subjects might think that the statement "Linda is a bank teller" tacitly excludes feminists; they might think it actually means "Linda is a bank teller (and is not a feminist)." If subjects interpret the wording this way, none of the statements are conjunctions of others and no probability rankings are wrong.

The linguistic interpretation can be tested in several ways. For example, use a between-subjects design in which some subjects rate the T statement without seeing the F&T statement (and vice versa); or replace "Linda is a bank teller" with the clearly comprehensive "Linda is a bank teller, who may or may not be a feminist," or with the more specific "Linda is a bank teller (and is not a feminist)" and see whether conjunction errors persist.

In fact, the purely linguistic interpretation appears to be wrong. Tversky and Kahneman (1983) tried both the between-subjects and the clearly-comprehensive variations and still found persistent conjunction fallacies. Others manipulated subtle details of wording and found no substantial changes in some conjunction problems (Morier and Borgida 1984) and some error reduction in others (Krosnick, Fan, and Lehman 1990).

2. Underweighting of Likelihood Information (Conservatism)

A second bias is underweighting of likelihood information, or "conservatism." Conservatism has been observed in Bayesian updating tasks like the one Grether studied. Consider two bingo cages, A and B. Bingo cage A contains seven red and three blue balls; B contains three red and seven blues. Suppose each cage is equally likely. Suppose a sample of eight reds and four blues is drawn (with replacement, of course), which clearly favors the A cage. What is $P(A \mid 8 \text{ red, } 4 \text{ blue})$? The typical response is between 7 and 8 but the Bayesian posterior is actually .97. Subjects are far too conservative in drawing conclusions from samples like these. One estimate derived from experimental data suggests that it takes two to five observations to produce a perceived diagnostic impact equal to the Bayesian impact of one observation (Edwards 1968).

McKelvey and Page (1990) ran a study in which subjects observed different parts of a full sample, then reported probability estimates to each other. After hearing the estimates of others, people reported new estimates (taking into account the estimates of others), and so on for several rounds. (This iterative process resembles the aggregation of information through polls and other processes; see McKelvey and Ordeshook [1985]). They observed some conservatism in updat-

ing of probabilities. Eger and Dickhaut (1962) found some conservatism when accounting students simply reported probabilities, but the effects were substantially reduced when stimuli were described an accounting context. The conservatism also disappeared when subjects revealed probabilities by betting against an experimenter in a way that penalized Bayesian errors. By contrast, Sanders (1968) found conservatism in a proper scoring rule with no financial incentives.

At first glance, the connection between evidence of base rate neglect (reviewed in the last section) and conservatism seems to indicate that people use Bayes's rule on average, but sometimes they weigh base rate, too little and sometimes too much. This is a weak justification for adhering to Bayes's rule as a descriptive principle in all circumstances, if one can predict the situations in which the two errors occur. By analogy, it might be the wrong thing to wear on a trip with stops in Alaska and Taipei, even if it is appropriate for the average temperature of the two places). Whether error is predictable across situations is then the crucial empirical question.

There seem to be several reasons why base rates are underweighted in some settings (the taxicab problem) but sample information is underweighted in others (conservatism experiments). Base rates are incorporated when they are salient or interpreted causally, as they are likely to be in the conservatism experiments (partly because judgments are repeated). Also, sample information may be underweighted in the conservatism tasks because it is not highly representative (as witness accuracy is in the taxicab problem); for instance, no sample of draws exactly matches the contents of the bingo cages. Furthermore, Griffin and Tversky (1992) argue that conservatism and base rate neglect are opposite sides of the same coin. They argue that both phenomena result from people overweighting the strength of evidence and underemphasizing its weight. Conservatism is a kind of underconfidence that results when people underemphasize the large size (or weight) of a sample of weak evidence. Base rate neglect occurs because people overemphasize strong evidence.

3. The Law of Small Numbers and Misperceptions of Randomness

In the logic of the representativeness heuristic, there is more probability for sample size: a small sample can represent the population or process that generates it as well as a large sample can (for example, a sample of two coin flips, one head and one tail, represents the Bernoulli trial very nicely). The belief that all samples will closely resemble the processes or populations that generated them is an intuitive extension of the law of large numbers to small samples, facetiously called "the law of small numbers" (Tversky and Kahneman 1971). The law of small numbers predicts that agents will gather too little data and will overgeneralize from small samples to distributions. In economic applications, they will search too little (see the evidence in section III.K) and form opinions quickly, compared to the default of optimal sampling and inference.

The law of small numbers also causes biases like the "gambler's fallacy": when people are asked to generate or identify random sequences their sequences often

show negative autocorrelation (Wagenaar 1972), because the prototypical representative random series repeatedly self-corrects to keep the sample proportion close to the population proportion (see Bar-Hillel and Wagenaar [in press] for a recent review). For example, lottery betting on a given number actually drops off sharply—only nearly half in the several days after that number wins (Osfeldt and Cook 1993).

Mathematically sophisticated subjects are better at generating truly random numbers, but so are children who have not yet learned the law of small numbers (Ross and Ury 1958; Chapman 1953). People can also be taught to choose randomly after several hours of training with excellent feedback (e.g., several measures of the randomness in the previous block of responses, Nuring et al. 1986), and see Edwards (1961b). These training data suggest that experienced agents in some settings might be able to learn to choose randomly. Whether they do in other settings, under natural conditions, is an empirical question.

Truly random sequences will show no negative autocorrelation. Observers who expect negative autocorrelation in a random series will be surprised by the number of times they see and will come to believe the series is positively autocorrelated. This misconception appears to be the origin of the unshakable belief among basketball fans and players that outcomes of shots are positively autocorrelated—players have "hot hands"—even though both field data and experiments show hits and misses are remarkably close to independent (Gilovich, Vallone, and Tversky 1985). Camerer (1989b) found that mistaken belief in winning streaks—teamwide hot hand effects—errors in betting odds on professional basketball games of about one point (The error is small because professional teams score roughly 100 points a game.) And Brown and Sauer (1993) question whether streak beliefs are mistakes, but their tests are inconclusive. Regardless, betting markets are active and patterns like the perceived hot hand are very easy to observe and profit from; such markets might be the worst place to find bias. The modest one point error suggests that larger effects might exist in markets that are less well policed.¹¹

Misperceptions of random sequences are important in game theory because mixed strategy play in repeated games assumes subjects can generate independent random draws (or appeal to independent privately observed hunches that others do not observe). O'Neill (1987) reported that average play corresponded to mixed strategy proportions in a zero-sum game with a unique mixed strategy equilibrium. Skeptical of the strength of O'Neill's conclusions, Brown and Rosenthal (1990) reanalyzed his data. Their reanalysis and subsequent work by others are a good case study illustrating how careful critique of an imaginative experiment can lead to further designs and a cumulation of knowledge. Briffin and Rosenthal first pointed out that some test statistics O'Neill used, like the percentage of times the row player won, had little statistical power to distinguish equilibrium mixed strategy play from various, dis-equilibrium alternatives. O'Neill random choices with equal probabilities for all strategies). Then they showed that despite the lack of power to detect deviations from mixed strategy predictions, about a third of the players did deviate significantly, in different directions.

Furthermore, choices were not independent across plays; choices often depended on one's own previous plays as well as on the opponent's previous play (and sometimes on the interaction).

The work on randomization is now extensive and suggests an area of genuine collaboration and cross-fertilization between economists and psychologists. The conclusions of O'Kell, and Brown and Rosenthal, have been largely replicated by Rapoport and Boebel (1992) and Mookherjee and Sopher (1994). In these games, many players seemed to believe in the law of small numbers: after they won by making a particular choice, they were less likely to try the same choice. Whether players can detect the predictable anomalies; that other players exhibit is an interesting open question.

Rapoport and Budescu (1992) compared the randomness of sequences subjects generated in two conditions: (1) a game condition in which they played a two strategy game with a unique mixed strategy equilibrium (with each strategy equally likely) 150 times and were paid for each choice according to the outcomes; and (2) a choice condition in which the subjects were simply told to choose randomly 150 times (and were not paid for their choices). Formally the two tasks are the same. But in the game condition, subjects created sequences that were more random than those they produced by random choices; for example, subjects reversed their previous choice 59.1 percent of the time in the choice condition (exhibiting "negative recency"), but they only reversed 53.4 percent of the time in the game condition. (A true randomizer would reverse 50 percent of the time.) One explanation for the difference is that subjects played the game more seriously because money was at stake. Rapoport and Budescu suggest a second, psychological explanation: remembering previous choices is essential for choosing nonrandomly; perhaps playing the more complex game inhibited memory of previous choices, making it more difficult not to randomize.

Mookherjee and Sopher (1994) compared behavior in two conditions, in which subjects learned their own payoffs but the choices and payoffs of others were either known or unknown. In many "routine-learning" models, knowing the choices and payoffs of others is inessential because players are assumed to simply choose strategies that yielded high payoffs in the past. These models predict that behavior in the known and unknown conditions should be the same. Behavior was not the same: convergence toward equilibrium mixtures was more rapid when the other players' choices and payoffs were known, suggesting a sophistication that the routine learning models do not capture.

Methodological Digression: Training

Several psychologists have tried to train subjects to avoid judgment errors ("debiasing"). Fischhoff (1982) reports some successful training exercises, mostly using large amounts of well-structured feedback. For example, overconfidence and hindsight bias (discussed in section II.F below) can be reduced by having people generate reasons why their predictions or recollections might be "Wrong." (Groups and organizations might debias individual judgment if several people

generate such reasons, questioning each other's judgment. But the opposite could occur too: groups could inflame bias. If groups generate supporting arguments or if overconfidence is taken as a signal of knowledge.)

Extensive studies suggest reasons to be pessimistic about how well training transfers across time or tasks. When subjects adapt to a setting and optimize in it, it is often the case that they have not learned a general rule they can recognize or apply to a structurally identical task that has different surface features. For example, Kagel and Levin's (1986) subjects learned to avoid the winner's curse in three bidder markets, but overbid when three bidders were added; compare with chapter 7.

Nisbett et al. (1987) and Larrick, Morgan, and Nisbett (1990) trained subjects to use simple statistical rules and ignore sunk costs. For example, 45 percent of their trained subjects gave a correct response on a sunk cost problem, compared to 29 percent of untrained subjects. A month later, the trained subjects reported they had bought and used 1.14 objects in activities (e.g., they returned a rented videotape without watching it—compared to 0.84 by untrained subjects). These are modest victories for training, but the breadth of the rules people have learned to use is subject to debate.

4. Market Level Tests of Representativeness

A central question for economics is whether individual judgment errors aggregate to create errors in market prices, allocations, and efficiencies (and whether they aggregate in groups, firms, and societies). The aggregation question has been addressed theoretically by Haltiwanger and Waldman (1985), Russell and Thaler (1985), Akerlof and Yellen (1985), and many others. Whether individual errors affect market behavior depends on the answers to many deep questions (see Camerer 1992b): do rational agents have more impact than irrational agents? (Do the more rational agents know who they are? Can they get more capital? Is it always optimal to behave rationally when others are not?) Do irrational agents learn from others? Can they buy and sell? Can they go bankrupt? (Will they be replaced by other irrational agents if they do?)

The answers to these questions will undoubtedly vary across markets (cf. Zeckhauser 1986, table 1). They are fundamentally empirical questions. To answer them, some evidence of judgment error in market behavior has been gathered (see Thaler 1992). But the evidence is inevitably controversial because it is easy to construct rationalizations of apparent market anomalies based on risk aversion, transaction costs, unobserved variables, or the current fashion-information asymmetries. To test whether individual errors affect markets, it is therefore helpful to conduct market experiments in which competing explanations can be ruled out. The next section describes some studies of whether errors in Bayesian judgment, caused by representativeness, affect prices and allocations in markets.

Duh and Sunder (1986) published the first market study. They tested whether underweighting of base rates (see section II.D.1 above) affected prices in a market

for 100-period assets. Their design and those used by others (Camerer 1987, 1990) were both inspired by Grether (1980) I will describe my own design in some detail.

Asset values depended on which of two states, X and Y, occurred. States were physically represented by bingo draws. If X had occurred (for example, three balls were drawn with replacement from an X bingo cage, hidden in a box, containing one red ball and two black balls; If Y occurred the cage had two reds and one black.

Subjects were given two shares of an asset and loaned experimental currency (fifty cents each period). Subjects earned a state-dependent dividend for each share they held at the end of the period. (There were two dividend schedules, creating type I and type II traders.) The value of the assets to subjects therefore depended on their subjective probabilities of X and Y, which depended on the sample of balls drawn from the X or Y bingo cage. A sample of one red and two blacks indicates the state is likely to be X. A Bayesian subject would calculate

$$P(X | 1 \text{ red}) = \frac{P(1 \text{ red} | X)P(X)}{P(1 \text{ red} | X)P(X) + P(1 \text{ red} | Y)P(Y)} = .75$$

The judgment literature suggests several alternative hypotheses about how subjects estimate $P(X | 1 \text{ red})$. One could interpret representativeness to imply $P(X | 1 \text{ red}) = .5$ (Duh and Sunder's JLBRI model). An interpretation more faithful to representativeness is $P(X | 1 \text{ red}) = P(Y | 2 \text{ red})$, which follows if the base rates $P(X)$ and $P(Y)$ are ignored. Both of those theories are rejected by the data. An hypothesis that fits better is "exact representativeness," in which $P(X | 1 \text{ red}) > .75$ because a sample of one red and two blacks exactly matches the contents of the X bingo cage. Given any such assumption about $P(X | \text{sample})$, predictions about prices and allocations can be derived by assuming risk neutrality and competitive equilibrium.

Each experiment had thirty to forty trading periods with stationary replication, using different samples each period. Since the equilibrium price varied among four possible prices from period to period (one price for each different sample), prices were volatile and convergence was slow. Camerer (1990) reports the time series of mean prices in all sessions. Camerer (1987) condenses the data in a compact way, shown in Figure 8.3, giving a time series of confidence intervals around the mean price across several sessions with the same parameters. Bayesian predictions for each sample are shown by horizontal lines.

Figure 8.3 shows that mean prices vary across sessions (the confidence intervals are wide). Across periods with the same sample, prices converge roughly to the Bayesian predictions. An "R" denotes the direction of price deviations expected by (exact) representativeness. Prices in zero-red and three-red periods are remarkably close to Bayesian (though the Bayesian probabilities are also close to 0 and 1 in those cases, so there is less room for error). Prices in the one-red and two-red periods begin well below the Bayesian prediction, and converge above it in the direction of representativeness.

Dividend parameters were carefully chosen so that final allo-




Figure 8.3. Confidence intervals for mean price in multiple experimental sessions. "R" denotes the direction of price deviations expected by (exact) representativeness. Camerer (1987).

cations of shares, as well as prices, would distinguish the Bayesian and representativeness explanations. When the sample had one red, the Bayesian expected values were equal for type I and type II traders, so the Bayesian theory predicts an equal allocation of shares among the two types (even if subjects are risk averse or risk preferring). The representativeness theory predicts that type I traders will hold more shares. In fact, about 80 percent of the shares were held by type I traders (90 percent after traders were once again supporting the representativeness theory).

The price biases are modest in probability terms, but small in cost—about a nickel per trade, or a couple of dollars over an experiment. However, paying subjects five times as much magnifies the difference.

Experience reduces price bias but does not eliminate it. Experience also reduces noise in prices and trading, which makes biases more statistically significant and causes allocations to reject the Bayesian assumptions even more strongly.

Duh and Sunder (1986) ran similar experiments using a design with several interesting variations (including a wider range of prior state probabilities and a labeling difference). Their results are similar. Subjects were close to Bayesian, but they erred in the direction of an extreme representativeness theory that predicts that a sample of one draw is taken as perfectly diagnostic. Auctions were strongly supportive of representativeness.

Anderson and Sunder (1989) report further experimental sessions comparing students with professional securities and commodity traders, in a design similar to Duh and Sunder (1986) and Camerer (1987). In three of the 100 sessions, the prices predicted by the Bayesian theory were below the prices predicted by repre-

representativeness theories. Prices were far from Bayesian in the student sessions, and closer to Bayesian in the sessions with traders. (However, in their design representativeness pushes prices below Bayesian levels, just as risk aversion does. It is possible the professionals are simply less risk averse, and therefore appear more Bayesian.) Allocations did not favor either the Bayesian or representativeness theories, though there is some movement toward Bayesian allocations in the professional trader sessions. The professional traders allowed less underweighting of base rates in a word problem similar to the cab problem discussed in section II.C1. While these data are limited, they do suggest professional traders are less prone to neglect base rates, and more likely to trade assets closer to Bayesian prices, than students.

Ganguly, Kagel, and Moser (in press) ran market experiments using an asset whose value depended on the success of a hypothetical firm. A balance of success and likelihood information were given, making the firm's success isomorphic to the cab problem, which elicits large underweighting of base rates. Their study has two innovations: a word problem is used to test for representativeness (the studies above used bingo cages); and both individual judgments and market prices are measured each period, so one can measure whether market prices reduce individual error. In two sessions, judgments and prices were much closer to the base rate neglect prediction than to the Bayesian prediction. There is no apparent convergence to Bayesian price levels across sixteen periods.

Plott and Wilde (1982) studied product markets in which agents give advice to buyers about product quality, based on samples of data that agents can see but buyers cannot. They report some evidence that agents used representativeness, rather than Bayes's rule, in drawing inferences from the samples.

This small collection of work on market level effects of Bayesian judgment errors suggests that errors caused by representativeness are the rule, not the exception, in experimental market prices and allocations. The errors are relatively small when uncertainty is generated by chance devices (bingo cages). Errors are quite large when uncertainty is generated by word problems like the taxicab problem. Market experience appears to reduce error, though not eliminate it (Camerer 1987); professional trading experience does appear, tentatively, to reduce error (Anderson and Sunder 1990). Future research could profitably continue to map out the boundaries of the influence of judgment error. Could individual and their market behavior more carefully, and replicate whether experience and education reduce error (as they appear to, in some studies). Studies should also begin to disentangle the influences of incentives, competition, task repetition, learning from other people's feedback, and learning from actions of others, which are currently confounded in market treatments.

D. Confirmation Bias and Obstacle to Learning

An important source of disagreement between psychologists and economists concerns learning. Psychologists often suspect that the immediate, frequent, clear, exogenous feedback subjects receive in economics experiments overstates how well people learn in natural economic settings. Economists, in contrast, think that

experiments underestimate the rate of natural learning because context, advice, higher incentives, and added time to reflect or calculate are absent from experiments, and probably improve performance in natural settings.

Much of the psychologists' pessimism arises from evidence of heuristic tendencies in judgment that present obstacles to learning (e.g., Einhorn and Hogarth 1978). I mention only two obstacles here. One tendency is called "confirmation bias." A tricky problem due to Wason (1968) can be used to illustrate this:

You are shown four cards, marked E, K, 4, and 7. Each card has a letter on one side and a number on the other. You are given the following rule: Every card with a vowel on one side has an even number on the other side. Which cards must you turn over to test whether the rule is true or false?

In the four-card problem, most subjects answer that E must be turned over, and E and 4. They think you should check the one card with a vowel (E) you should—and perhaps check the even-number card (4) too. The correct answer is E and 7. Few subjects think to turn over the 7, but they should because the rule is falsified if the 7 card has a vowel on the other side. The four-card problem suggests that in testing an hypothesis, people instinctively seek evidence that could confirm the hypothesis: for example, finding a vowel on the other side of 4 would provide support for the rule, but that evidence could actually never test whether the rule is always true (as turning over the 7 can). Confirmation bias is one force that may inhibit learning.

A related problem is the production of treatment effects (or self-fulfilling prophecies). When people believe an hypothesis is true, their actions often produce a biased sample of evidence that reinforces their belief. A busy waiter who thinks poorly-dressed patrons tip badly will give them poor service and receive a bad tip, reinforcing his theory. Treatment effects inhibit learning whether one's underlying belief is false. (The only way to test the belief is by experimenting, giving good service to a poorly-dressed patron.) Expectations of a bank run, or a bubble in asset prices, can be self-fulfilling in a similar way.

E. Expectations Formation

There is a large literature on expectations formation, and a smaller literature about expectations people form about variables generated endogenously by their own collective activity—future prices, for instance.

There are many studies of whether price expectations are rational. Most of the studies use published forecasts by economists, businessmen, or professional economists (see Lovell [1986], and Williams [1987], for reviews). Generally, they find that forecasts are biased: forecast errors (forecasts minus actual results) have a nonzero mean. Forecasts usually violate rationality of expectations: they are correlated with observable variables (typically past forecast errors and current forecast levels), implying that some available information is ignored when forecasts are made. Forecasts also usually follow an adaptive process in which forecast changes are related to past forecast errors (Nerlove 1958, and see below).

Apparent violations of rationality of naturally-occurring forecasts could be due to Bayesian learning in an economy where the statistical process generating out-

com keeps changing (Caskey 1985; Lewis 1989). For example, the surveys show that businessmen were consistently surprised by the persistence of price inflation in the 1970s (their forecast errors had a negative mean), it could be argued that their forecasts were rational. And, ante, but the stochastic inflation process chilled during the period and it took forecasters some time to learn whether the change was temporary or permanent.

To control for changes in the statistical process generating outcomes, several experiments examined forecasts of outcomes of a statistical process that is unknown to subjects but fixed throughout the experiment, and known to be fixed (Schmalensee 1976; Garner 1982; Solle 1983). Their results are generally consistent with rationality of expectations too, but suggest some learning and rationality in special settings (a random walk with no drift, in Dwyer et al. [1993]).

Several researchers have gathered forecasts that subjects in an experimental market make of the future prices that they themselves generate (Carlson 1967; Knez, Smith, and Williams 1985; Williams 1987; Daniels and Plott 1988; Smith, Suchanek, and Williams 1988; Wellford 1989; Caruana and Weigelt 1990; Peterson 1993). I will describe the Williams (1987) study in detail. He conducted experiments with a series of five double auctions for single period goods. Starting after the first period, subjects forecasted the mean price in the next period. The person with the lowest cumulative absolute forecast error earned \$1. The forecasts were generally remarkably accurate, but the small deviations from rationality were statistically significant. Nearly half the forecast errors were zero (the mode), but on average prices were about a penny too high. (Since goods cost about \$5, a penny deviation is tiny.) Forecast errors were modestly autocorrelated ($r = .15$, $p < .01$). Expectations were estimated to be adaptive¹⁷ if the adaptation coefficient was positive in the specification:

$$(1) \quad E(P_t) - E(P_{t-1}) = b[P_{t-1} - E(P_{t-1})] + e_t$$

(where $E(P_t)$ denotes the forecast of prices in period t and P_{t-1} is the actual price in period $t-1$). Williams estimated $b = .86$. Forecasts of experienced subjects, who had participated in other auction experiments, were less biased and less error prone, but not by much. Peterson (in press) found that the estimated bin equation (1) converged to one across periods of an experiment, and changes were judged for the least experienced subjects.

Besides contributing to the debate about rationality of forecasts, the studies by Williams and others allayed methodological fears that simply gathering forecasts might affect market behavior. One concern was that asking subjects to forecast prices before each period might increase or decrease their attention to market behavior and affect convergence. But patterns of convergence looked like those in previous double auction experiments (see chapter 5 for examples), suggesting there was no such effect. Another concern is that subjects who are rewarded for making accurate forecasts may sacrifice trading profit to collect the best-forecaster bonus. There was no evidence of such an effect either.

Williams's methods and results are typical of most OUE studies. Forecasts are usually slightly biased (too low if prices are rising; too high if prices are falling).

Forecast errors are autocorrelated and correlated with some observables (previous price changes or current forecast levels). And forecasts are generally adaptive; estimates of the adaptiveness coefficient are remarkably constant across a wide variety of studies, between .6 and .8.

A notable exception is Daniel and Plott (1988). They studied forecasts in goods markets with price inflation that was induced by shifting supply and demand curves upward by 15 percent each period (until the last few periods). Prices adjusted to the inflation a bit sluggishly. Graphs of average forecasts and prices suggest that forecasts were biased and autocorrelated (they were too low during inflation, and overshot when the inflation stopped). But regressions indicated that subjects' forecasts were rational rather than adaptive. It is not clear why the expectations of their subjects (Cal Tech students) are not adaptive, as they are in most other studies.

Price forecasts are easy to gather. Perhaps experimenters should collect them routinely. So far, the forecasts have not generally been put to much use to inform either psychology or economics, but they could be. A good example is Smith, Suchanek, and Williams (1988) who use evidence of systematic forecast error to explain why price bubbles persist in experimental asset markets (see chapter 11).

Psychological Studies of Expectations

Psychologists have done two kinds of studies germane to understanding rationality of expectations. One kind is studies of "multi-cue probability learning" (MCPL) (e.g., Castellan 1977). In MCPL studies subjects try to predict a dependent variable from given values of predictor variables, in a series of 100 or so trials. The studies indicate that learning is very difficult except in simple, deterministic situations (e.g., when dependent variables are a linear combination of independent variables; Brehmer 1980). Learning stochastic rules is especially difficult.

A second body of literature concerns judgments made repeatedly by people (many of them experts) in natural settings where stochastic outcomes depend on some observable predictors (e.g., test scores) and some unobservables, examples include medical or psychiatric diagnosis (severity of Hodgkins' disease, schizophrenia), predictions of recidivism with parole violation by criminals, ratings of marital happiness, and bankruptcy of firms. About 100 careful studies have been documented so far. The remarkable finding in almost all these studies is that weighted linear combinations of observables predict outcomes better than individual experts can (Meehl 1954; Dawes, Faust, and Meehl 1989). In a typical study (Dawes 1971), it was discovered that academic success of doctoral students could be predicted better by a sum of three measures—GRE scores, a rating of the quality of the students' undergraduate school, and her undergraduate grade—than by ratings of a faculty admissions committee. (Put bluntly, the faculty's deliberation just added noise to the three measures.) The only documented exceptions to the general conclusion that models outpredict experts are a few kind; of esoteric medical diagnosis,

In these studies, experts routinely violate rational expectations by using ob-

servable information inefficiently (worse than simple models do). The violations have two common forms: (1) experts often add error to predictions by using complicated interactions of variables (weighting grades from low-quality schools more heavily, for example), rather than more robust linear combinations of variables; (2) experts pay attention to observable variables that they should ignore because the variables are not highly predictive of outcomes (personal interviews, for example). These psychological tendencies can be traced to some of the judgment biases discussed above (e.g., Camerer and Johnson 1991).

E. Iterated Expectations and the Curse of Knowledge

In many economic settings, agents must guess what others think. These guesses are "iterated expectations," or expectations of expectations. We can express these formally as follows: suppose agent i and j have information sets I_i and I_j and agent i is guessing j 's expectation about a variable X . Then j forms the expectation $E(X | I_j)$ and i forms an iterated expectation, $E(E(X | I_j) | I_i)$.

Most asymmetric-information settings are modeled by assuming one agent knows strictly more than another ($I_i \supset I_j$). These models usually revolve around the less-informed agent's attempt to learn what the more-informed agent knows, perhaps by observing a signal (cf. the asset market examples in chapter 6). A hidden assumption in the models is that the more-informed agent has an accurate mental model of the less-informed. The psychology of memory and imagination suggests that assumption may be wrong: it is hard for the more-informed agent to forget what she knows and imagine what the less-informed agent is thinking, because her extra information is available in memory.

Normatively, $E(E(X | I_{\text{more}}) | I_{\text{less}})$ should equal $E(X | I_{\text{less}})$. If the extra information in I_{more} is hard to forget, empirical estimates of $E(E(X | I_{\text{more}}) | I_{\text{less}})$ will be biased away from $E(X | I_{\text{less}})$ toward $E(X | I_{\text{more}})$. This bias is called the "curse of knowledge" (Camerer, Loewenstein, and Weber 1989). It seems common: teaching is made difficult by knowing too much; after solving a problem it seems obvious that others should see the solution too¹⁹ (Nickerson, Baddeley, and Freeman 1987); and what writer of computer manuals—an expert, usually—has ever written one that novices can understand?²⁰

I. False Consensus and Hindsight Bias

Two brands of curse of knowledge have been studied in some depth. One brand is called "false consensus" (an unfortunate misnomer): people use their own tastes and beliefs as information in guessing what others like and believe (Ross, Greene, and House 1977). In one study, students were asked whether they would walk around a campus for thirty minutes wearing a sign saying "Eat At Joe's." Some did, others refused. The interesting finding is that 60th kinds of subjects thought others were likely to make the same choice they made. Those who wore the sign estimated 62 percent of others would; those who refused thought 67 percent would refuse. Using one's own tastes or beliefs as information is not a mistake

unless that information is overweighted. (It is reasonable to use one's own tastes as a single draw from the population distribution of tastes. Then two Bayesians with different tastes will have different posterior beliefs about the population—as the two groups of sign-wearers and refusers did—but a difference in posterior beliefs is not necessarily an error. See Dawes [1990].) I suspect overweighting one's own tastes might contribute to the high failure rate of small businesses: owners think more consumers share their tastes than actually do and either underinvest in market research or ignore its result.

A second kind of curse of knowledge is "hindsight bias": current recollections of past judgments tend to be biased by what actually happened since then (see Fischhoff 1975; Hawkins and Hastie 1990; Christensen-Szalanski and Willham 1991). Fischhoff and Beyth (1975) asked subjects about the likelihood of various events occurring before Nixon's historic trip to China (Will Nixon meet Mao?). Several months later, after the trip was over, subjects were asked to recall what probabilities they gave before the trip. They remembered having given higher probabilities than they actually had for events that happened, and lower probabilities for events that didn't happen. Subjects were not paid for accurate recollection but I bet the hindsight bias persists even with financial reward (assuming subjects cannot record their initial answers and look them up afterwards). Hindsight bias is often modest in magnitude but robust, and affects events with low ex ante probabilities most strongly. Hindsight bias appears to create second-guessing in firms, courts, and political institutions, which may create added employment risk when good ex ante decisions result in bad ex post outcomes (cf. Baron and Hershey 1988).

2. Market Level Test of Curse of Knowledge

Camerer, Loewenstein, and Weber (1989) tested whether the curse of knowledge affected prices in experimental markets.

Before the markets began, one group of "uninformed" subjects guessed the 1980 earnings-per-share (EPS) of several actual companies, based on accounting data from 1970–1979 and a Value Line profile of the firm's 1980 prospects. Call the uninformed subjects' average estimate $E(\text{EPS} | \text{data})$.

Traders in asset markets then traded a one period asset that paid a dividend equal to the average estimate of uninformed subjects, $E(\text{EPS} | \text{data})$. To value the asset correctly, market traders had to make the best possible guess of what uninformed subjects thought 1980 earnings would be. Market subjects knew the actual 1980 earnings per share. Their guess about the uninformed subjects' average estimate is therefore an iterated expectation, $E[E(\text{EPS} | \text{data}) | \text{data} + \text{EPS}]$. If market subjects suffer from the curse of knowledge, asset prices will be closer to true EPS than they should be. (It will be hard for traders to imagine that subjects could not have guessed the true EPS.) A separate control group of subjects did not trade in any markets, but simply made judgments about what the asset value would be (knowing the 1980 EPS); they were rewarded for accurate forecasts just as market traders were.

Titles setting is similar to underwriting, in which a group of expert buyers must purchase goods that are resold to a group of less-expert consumers. The consumers' opinions establish the value of the goods, which the expert buyers must anticipate. Financial underwriting or buying of clothes, art, or wine for retail sale are examples. The empirical question is whether the experts will let their extra knowledge get in the way when figuring out what nonexperts will buy.

Traders traded assets based on eight different companies, for two trading periods each. (We used two trading periods to measure the change in forecasts and prices between periods. We used only two periods so we would have time to trade several companies because we suspected—correctly, as it turned out—that there might be inherent variation in the degree of curse of knowledge across companies, which we hoped to average out by using eight companies.) The degree of curse of knowledge, or hindsight bias, was estimated by having traders give forecasts of the asset value three times, before each of the two trading periods and after the last period. We compared their forecasts to the actual asset value (0 percent bias) and the true 1980 earnings (100 percent bias). Figure 8.4 illustrates the degree of bias across the eight companies. The mean degree of bias in the forecasts of market subjects after two trading periods is shown by a regular line. The mean bias in the control-group individual judgments is shown in Figure 8.4 with a dotted line.

The market traders' forecasts exhibited roughly 30 percent bias, whereas the individual subjects' forecasts exhibited roughly 60 percent bias. Thus, market forces reduced the curse of knowledge in traders' judgments by about half, compared to the control group subjects, but did not eliminate the curse entirely.

A closer look at individual behavior suggests why market forces had an effect. Less biased traders were slightly more active in the markets (56 percent of bids, offers, and acceptances) than more biased traders (44 percent). Prices generally began between the 0 percent bias and 100 percent bias levels. Of the price changes that moved toward one benchmark and away from the other, 63 percent of the changes were toward the 0 percent bias benchmark. In these experiments, the market is actually a bundle of forces that could be separated in future work. Compared to the individual-subject control group, traders in the market made three forecasts (rather than one); spent more chronological time thinking about each company; and had the opportunity to learn from bids, asks, and acceptances by others. We suspect the third force is most important. The double auction market is a specialized communication mechanism that allows people to express their opinions, and learn from others, in a limited form. It could be usefully compared to other opinion-aggregation schemes (e.g., open group discussion) and other exchange institutions for bias-reducing properties.

The curse of knowledge implies that more information might hurt those who are trying to guess what people without the information think. If a harmful good is freely disposable—toxic waste, or curse producing information—it should have a price of zero; but information is *not* freely disposable if it is hard to forget or ignore. We ran two market experiments, with eight traders in each, to see whether traders would bid zero for curse-producing information (Camerer




Figure 8.4. Degree of hindsight bias (or curse of knowledge) in individuals (dotted line) and market traders (solid line). Source: Camerer, Lowenstein, and Weber 1989.

1992b). We auctioned off the actual 1980 EPS to four traders in a uniform price auction where the price was determined by the fifth highest bidder (cf. chapter 7). Information was auctioned off once for each of the eight companies.

Bids began at very high levels, close to the asset value itself (around \$3), probably because some subjects had no idea what the information was worth. Others bid zero immediately. After two or three company auctions, the market price converged to zero. Some subjects even made small negative bids—they wanted to be *paid* to know something worthless, perhaps because they knew it was not freely disposable and might hurt their judgments.

6. The Illusion of Control

People sometimes act as if tasks that involve only chance have an element of skill too. Psychologists call this belief the "illusion of control." Gamblers throw dice hard to produce high numbers. Lottery bettors buy "dream books" that explain what numbers to bet after a certain dream. People wait longer or pay more for specific numbers in a lottery than for randomly-assigned numbers (Langer 1975).

The illusion of control is one kind of "magical thinking," a misunderstanding of causal relation, akin to rain dances and superstitions!¹ There is some remarkable evidence that control illusion improves mental health. In one experiment, subjects were allowed to bet varying amounts of money, with varying expressions of confidence, on a chance device. Subjects who *did not* suffer from the illusion of control—they bet less money, less confidently—were more likely to be clinically depressed than others. (They are "sadder but wiser," Alloy and Abramson 1979.) Taylor and Brown (1988) review a wide range of evidence that suggests

that unrealistic illusions of optimism and control, rather than realism, are associated with mental health.

Control illusion might be important in agency relationships and compensation. In the standard economic model of agency, output is assumed to be a function $v(e, 0)$ of an agent's effort e and a random variable θ . The illusion of control implies that agents and employers overestimate the effect of effort (valde). In equilibrium, they will tie compensation too closely to output, and reward or punish more than is optimal.

In an experimental study of agency contracts, Berg (1995) found that control appears to matter in an interesting way. In her Mating, agents choose an effort level. Output is correlated with effort and with an observable signal. In the "control" condition agents have control over the signal because their effort is correlated with the signal's value. In the "no-control" condition, their effort is not directly correlated with the signal (so they have no control over it); but the signal is informative to principals because the joint distribution of output and signal is correlated with effort. Optimal contract should use the signal to determine an agent's compensation in both conditions, since the signal is always informative about the agent's effort (when coupled with observed output).

Subjects did use the signal in the control condition, but not in the no-control condition. They acted as if penalizing an agent for outcomes (signal values) beyond her control was pointless or unfair. The result jibes with evidence from natural settings that sharing housing contracts and executive compensation does not depend strongly on variables that are uncontrollable but informative, as it should (Wolfson 1985; Ahtle and Smith 1986).

H. Judgment: Summary and New Directions

The studies reviewed in this section suggest a variety of heuristic rules people use to make complex judgments: they rely on what's available in memory, and similarity, to judge likelihoods and correlations; they are overconfident when they state probabilities of events or forecast numbers; their expectations are adaptive, responding to observed marketwide behavior rather than expressing a rational understanding of the market (but see Lucas 1986); iterated expectations expectations about the expectations of others--are incorrectly influenced by memory; and people overestimate the influence of personal control.

Only a few of these findings have been replicated with the methods of experimental economists. More replications would test robustness of the findings.

Only a few of the replications took place within economic simulations (markets). The market experiments, often in Bayesian updating and iterated expectations, suggest markets reduce simple judgment errors but do not eliminate them. (Experience and expertise seem to reduce errors too.) More tests in which individual errors might be manifested in economic settings, including games or markets, would be useful. Tests could introduce institutional features such as overlapping generations, bankruptcy, access to capital, and advice markets, to carefully dissect precisely how markets reduce error.

m. Choice under Risk and Uncertainty

Most economists are familiar with theories that represent choice by numerical functions. (e.g., a utility function). Sometimes functional forms are simply posited, but utility theorists search for primitive axioms or preferences that imply a specific functional form (e.g., expected utility). A less familiar notion of choice theory is a process model that expresses the procedure a person uses to make choices, in an algorithm. (Expected utility maximization is an example of an algorithm.) Process models will generally not obey the axioms of utility theory, so the preferences they generate cannot be neatly summarized by a utility function.

Within economics there is a vast amount of work on automatic utility representations and a little work on process models. I will try to summarize both.

A. Expected Utility

During the development of statistical reasoning, it was taken for granted that proper choice meant maximizing expected monetary value. Provoked by the St. Petersburg paradox, in 1738, Daniel Bernoulli proposed maximizing some concave function of money (he suggested logarithmic), to reflect diminishing marginal value of dollars. Expected utility was born.

Almost two hundred years later von Neumann and Morgenstern (1944) showed, en route to game theory, that if preferences obeyed a particular set of axioms then those preferences could be represented by the expectation of some utility function.

The discovery of underlying axioms was important because it is easier to judge the intuitive plausibility of specific axioms than to judge the appeal of the utility representation they imply. (Establishing the axioms also laid the groundwork for modern theorists to weaken specific axioms and generate surprising alternative theories.) The utility representation can also discipline preferences by pointing out inconsistencies and violations of appealing properties (Strotz [1953, 1992] gives an example). Expected utility also provided a natural way to establish "measurable utility" (cf. Zeuthen 1937), which was in great demand at the time.

L. Notation and a Diagram

Some notation is useful. Denote lotteries by X, Y, Z , and probabilistic mixtures of lotteries by $pX + (1-p)Y$. (Specific outcomes are just degenerate lotteries with full outcome probability of one.) Denote X preferred to Y by $X \succ Y$, and X indifferent to Y by $X \sim Y$.

Predictions and data can be usefully displayed in the triangle diagram developed by Marschall (1950) and put to good use by MacLennan (1982) and others. Fix three gambles X_L, X_M, X_H (the subscripts represent low, medium, high) such that $X_H \succ X_M \succ X_L$, and $x_{L,M} \succ X_L$. (In most experiments, the gambles are




Figure 8.5. The Marschak-Machina triangle diagram (showing EU indifference curves). Sources: Marschak 1950; Machina 1982.

degenerate gambles with certain outcomes, such as 0, \$5, or \$10). Take the set of compound gambles which are probability mixtures in which each of the three gambles occurs with (objective) probabilities P_L , P_M , P_H . If we assume that two stage compound lotteries are equally preferred to single stage gambles with the two stage probabilities multiplied together (an important assumption we return to later) then this set of gambles can be represented in two dimensions, in $P_L - P_H$ space, as in Figure 8.5. (The third dimension, P_M is implicit in the graph because $P_M = 1 - P_L - P_H$.) Since the sum of the probabilities cannot be greater than one, the set of feasible probabilities is a triangle bounded by the lines $P_L = 0$ (the left edge), $P_M = 0$ (the hypotenuse), and $p_H = 0$ (the lower edge). A utility theory makes a specific prediction about the shape of indifference curves that connect equally preferred gambles in the triangle diagrams.

2. The Axioms

The axiom system von Neumann and Morgenstern used to derive EU got refined by others (e.g., Marschak 1950; Herstein and Milnor 1953). The crucial axioms are as follows:

1. **Ordering.** Preferences are complete (either $X \succ Y$, $Y \succ X$, or $X \sim Y$) and transitive ($X \succ Y$, $Y \succ Z \implies X \succ Z$). Graphically, completeness guarantees that any two points in the triangle are either on the same indifference curve or on two different curves; transitivity guarantees that indifference curves do not cross *within* the triangle (e.g., Fishburn 1984).
2. **Continuity.** For all $X \succ Y \succ Z$, there exists a unique p such that $pX + (1-p)Z \sim Y$. Continuity guarantees that there are no open spaces in the indifference map; uniqueness of p guarantees that indifference curves are not "thick."
3. **Independence.** If $X \succ Y$, then $pX + (1-p)Z \succ pY + (1-p)Z$ for all Z and $p \in (0, 1)$. Independence implies indifference curves are parallel straight lines.

The axioms imply that preferences can be represented by a numerical utility index, and the utility of a gamble is the expected utility of its possible outcomes. For a discrete lottery with several outcomes x_i , each with a p_i chance (denoted I, p, X_i) the functional form for EU is:

$$(2) \quad U(I, p, X_i) = \sum p_i u(x_i)$$

In expected utility, probabilities of outcomes are assumed to be objective and known; choices are made under "risk." But in most natural settings probabilities are not well known or agreed upon; choices are made under "uncertainty." In subjective expected utility (SEU) (Ramsey 1931; Savage 1954), people take acts that yield consequences in uncertain states (see section III.F later). If act preferences obey several axioms like those in EU, preferences can be represented by an expected utility of consequences weighted by beliefs about states (their subjective, or "personal" probabilities). Anscombe and Aumann (1963) fused EU and SEU by allowing outcomes with objective probabilities ("roulette lotteries") and uncertain states with subjective probabilities ("horse lotteries"). Since most of the debate about EU holds for SEU too, and most experiments test only EU, I defer further discussion of SEU until below.

B. Some History and Early Evidence

The publication of von Neumann and Morgenstern's book in 1944 caused quite a stir. Economists had just become satisfied with ordinal utilities, unique only up to monotone transformations, and knew how much analysis could be done using preferences that reveal only ordinal utility (Hicks and Allen 1934²⁴). Then, just as economists became convinced that cardinal utility was unnecessary, von Neumann and Morgenstern discovered a simple way to derive utility cardinally: $u(X) = p$ when X is judged to be indifferent to $pH + (1-p)L$ (and $u(H) = 1$, $u(L) = 0$ arbitrarily).

1. Three Controversies

The first of three immediate controversies was mathematical. In their book, von Neumann and Morgenstern said nothing about an outcome set, indifference, or an independence axiom. In brief symposium papers in *Econometrica*, Samuelson (1952) and Malinvaud (1952) solved these mysteries and showed how the now-familiar independence axiom followed from von Neumann and Morgenstern's axioms. The second controversy was confusion over whether a von Neumann-Morgenstern utility function was a riskless value function too (à la Bernoulli), and could either be derived from preferences over lotteries or by directly comparing differences in lottery outcomes (Ellsberg 1954).²⁵ It's *not* a riskless value function.

The third, and greatest controversy came at a symposium in Paris in 1952, where Maurice Allais presented two papers critical of the descriptive power of the theory of the "American school" (including Friedman, Savage, de Finetti,

Mruuhak, and oth«s died above) and introduce«l llis fnmous paradoxes (Allais 1953). {More about them below.)

After von Neumann and Molgeruitem's book was publihed, empirical lem began to trickle in. Excellent reviews of early v-Ork rue Edwards i 1954c, 1961a} and Luce and Suppes (1965). The collection edited by Thrall, Coombs, and Davis (1954) wile, -tt; the spirit of grop with the new modd& in an hlstorically fasdw nating way.

2 Initial Tests

Preston and Baratta { 1948) did the first tesL They auction«! off chances to vill r points with probability p, for 42 ti, p) pairs.: " Bids were approximalely linear in outcome x and nonlinear in probability p. Low-probability gambles ($p = .01, .05$) were sold for ,eveal times e ted value, high-probability gambles for slightly lt\$S than expected \<alue. (Indeed, the probabHity weight function they estimated, shown in Figure 8.6, looks strikingly like the "decision weight" function hypothesized thirty years later by Kahnman and Tversky f19791, but the interpretation is a bit different?i) The modern reader will suspect the quality of their e-vldence because 100 method& are casual and unorthodox, the sample is small, and hypotheses ant not tested. But similar findings of nonlinear probability weights were reported by Griffith (1949) using muclrlek betting (later McGlothlin 1956), Attneave (1953) using guessing games, and Yaari (1965) using indifference judgments.

Mosteller and Noguee ()951) estimated utility curves by offerin_g subjects complicated bets on three-die outcomes. Subjects played all bets they chose, and earned about \$1 per hour. Since the same bets were offered repeatedly, <etlainty-equivalents were estimated by observing which bets subjects took half the time they were offered.

Student subjects were -Hghtly ri averse: National Guardsmen were rather riik seeking.⁷⁹ Using utilities estimated from ooe sample to predkt freh choices, EU got about 70 per-ent right {compared to 50 percent right for expected value). There was strong evidence of nonlinear probability weighting by National Guardsmen (muclh as in Preston and Baratta's study) but not by students, as shown in Figure 8.6.

There were several complaints about Mosteller and Noguee's design. One was that the bets had complicated probabilities. Edwards began to test EU with an eye toward measuring subjective weight,; of probabilities, using simpler stimuli.⁸⁰ In choices aJlOlg gambles that were played out he discovered coosistent "probability preferences" {overweighting of specific probn.hllities), notably a preference for .5 {Edwards 1953, 1954a) Replkmiou in a military context (hypoHetical choices of attack targets) yielded ,;lightly different result; (Edwards 1954bt Nobody has found quire these probm;hty preference" since. Edwards's data also show that probability weights for potential gains and losses differ, indicating a kind of "wishful thinking" (cl. Irwin 1953³¹).

In the early 1950s the psychologml Clyde Coom began trying to measure subjective probability and ulllity simultaneously. (Edwards showed that th! was




Figure 8.6 Empirical decision weight <iveI from IW<early studies. Source: Preston and Baratta (1948); Mosteller and Noguee (1951).

important to do, since nonlinearities in subjective probability were observed when utility was assumed to be linear, but he did not show precisely how to do it.) Some of Coombs's early results were supportive of expected utility. But Coombs inihally scorned gambles over money (thinking subjects could not avmd responding to numerical money values rather than psychological utilities), so it is impossible to compare the predictive accuracy of EU to the llitlural benchmark expected value. Using money, Coombs and Komorita (1958) found thal utilil ties did satisfy a kind of additivity. Hum and Siegel (1956) found that an "ordered metric" utility model ompredicted expected value, in an experiment where pris.ooser chose bets over cigarettes (a common setting for early p,ychology experiments).

Da\idson, Suppes, and Siegel (1957) untangled probability weights and utility most carefully. Their technique depends on finding an event that is perceived to be exactly as attractive to bet on as to bet against. After rejecting coins and reguW dice, they ch(& a sii:sided die with two nonsense syllables (ZEV and ZOV, shown by others to have few almost oomertial associations) on three sides each. Since svbjcts did not care whether bets pud off on ZEV or WV, their subjective probabilities could be taken to equal .5. Using choices over bets on the special die, bounds on utilities could be determined while holding subjective probability fixed.;

llie estimated utility fuction.s were remarkahfy ronsiitent i<cr>ffm sessions. Twelve of fifteen subjects had nonlinear Clives, typically showing risk preference

for gains and risk aversion for losses. An experiment betting on one side of a four-sided die showed that people gave an event with objective probability .25 a decision weight of about .25.

In later studies Tversky (1967a, 1974b) and others were able to operationalize the independence of probability and utility (the crucial feature of EU) as a kind of additivity that was easy to check experimentally. Tversky experimented with prisoners, playing gambles for money, candy, and cigarettes using the Becker-DeGroot-Mohr (1964) procedure. Their choices generally obeyed additivity... i.e., independence but showed either a utility for gambling (when riskless value functions and risky utility functions were compared) or subadditive weighted probabilities (see also Edwards 1962). Experiments by Wallsten (1971) supported independence too.

In the 1950s people also began exploring stochastic choice models (Debreu 1958; Luce 1958, 1959; Luce and Suppes 1965), which allow subjects to choose differently when facing the same choice several times. In these models a gambler's probability of being chosen out of a pair increases with its utility³¹ (in EU, the probability increases in a step, from 0 to 1). Many experiments (e.g., Mosteller and Nogee 1951) were designed with stochastic choice models in mind, which meant long repeated sessions in which subjects made the same choice many times. (Wallsten 1971 had four subjects, each making choices for thirty hours.) This technique is largely out of fashion now except in some domains of mathematical and experimental psychology.

C. Mounting Evidence of EU Violation (1965-1986)

Except for some evidence that probabilities were weighted nonlinearly, and the simmering impact of Allais's paradoxes, EU emerged relatively unaffected from the first waves of tests. Elicitation yielded reasonable utility functions: expected utility predicted choices better than expected value did; independence of probability and utility was generally satisfied. Then evidence of paradox began to mount.

1. The Allais Paradoxes

Many felt Allais's examples used such extreme sums that they did little damage to everyday application of EU. But the examples were provocative and were repeated repeatedly in the 1960s and later with smaller sums (and paid subjects).

The most famous Allais example illustrates a "common consequence effect." Subjects choose between A1 and A2, where A1 = 0 million francs and A2 = 10 chance of 5 million francs, .89 chance of 1 million francs, and .01 chance of 0, denoted A2 = (.10, 5 million francs; .89, 1 million francs; .01, 0). They also choose between B1 = (.11, 1 million francs; .89, 0) and B2 = (.10, 5 million francs; .90, 0). It is easy to show that the frequent choice pattern, A1 > A2 and B2 > B1 violates expected utility. The choices are shown in a triangle diagram in Figure 8.7. It requires that indifference curves be parallel lines. Since the chords connecting the choices A1 and B2 and the choices B1 and A2 are parallel, subjects with parallel indifference curves must choose either A1 and B1, or A2 and B2.




Figure 8.7. Allais's common consequence paradox (A > A', B > B'): a violation of EU effects.

The Allais paradox attacked EU in a fundamentally different way than the painstaking empirical tests of the 1950s did. The paradox circumvents elicitation of utilities and directly tests consistency required by the axioms, using two pairwise choices. Most recent tests have followed this route too. An important drawback is that the degree of inconsistency and differences among individuals are hard to measure with this method.

The first Allais replication, in a dissertation by MacCrimmon (1965), reported a 40 percent EU violations. Morris (1967) used gambles over aerial grade points, and found 30 percent violation. (Open discussion among subjects improve consistency with EU slightly.) Skwirc and Tversky (1974) found 60 percent violation and presented subjects with written arguments pro and con EU. After reading both arguments, slightly more subjects switched their choices to become inconsistent with EU than became consistent.

MacCrimmon and Larsson (1979) revealed a weak axiom 11 and empirical evidence, and reported new data resting robustness of the paradoxes. They found fairly robust common consequence effects with different parameters (though the effects are strongest for extreme payoffs and probabilities). They also studied a so-called "common ratio" problem due to Allais (see Figure 8.7): choose either C1 = (1 million francs) or C2 = (.80, 5 million francs; .20, 0), and either D1 = (.05, 1 million francs; .95, 0) or D2 = (.04, 5 million francs; .96, 0). Notice that

the payoffs in C1 and C2 have the same ratio of probabilities, i.e., as in D1 and D2 (.05/1/4): hence the term "common ratio effect." People often choose C1 > C2 and $m > D1$, violating EU.

MacCrimmon and Larsson (1979) found the common ratio effect was less robust than the common consequence effect: A majority violated it with large ratios, when the winning probabilities in the C gambles were much different than those in the D gambles, but the rate of violation fell to a third or less as the

and probability differential fell. The best evidence was the strongest indication that the proportion of subjects violating EU might vary dramatically with choice parameters, suggesting a direction for further tests and ripe opportunity for alternative theories.

2. Process Violation

As Allais' paradoxes continued to provoke debate, a second wave of psychological evidence began to rise, even more deeply critical of EOM, a descriptive theory. In experiment after experiment, subjects appeared to use procedures or processes which were much simpler than EU (or even EV). For instance, in one study the value of gambles was better predicted by an additive combination of probability and outcomes ("risk dimensions") than by their product (Slovic and Lichtenstein 1968), even when subjects were told the gambles' expected values (Lichtenstein, Slovic, and Zink 1969).

In some experiments subjects compared the probabilities of winning in two lotteries, and their outcomes, in a way that led to intransitive cycles (Tversky 1969). Loomes, Starmer, and Sugden (1991) give a recent illustration. In their study, many subjects chose $(.6, £8) > (.3, £18)$ and $(1, £4) > (.6, £8)$ but also chose $(.3, £18) > (1, £4)$. (These cycles accounted for about 17 percent of the patterns resulting from the three pairwise choices.) An intuitive explanation is that these subjects chose the gamble with the larger probability as long as the payoffs were close, but chose $(.3, £18) > (1, £4)$ because the payoff £18 is sufficiently greater than £4. In some experiments on multiattribute choice under certainty (finding an apartment), subjects were not told the attribute values for each choice, like the rent on apartment 1 or the size of apartment 3, unless they asked to see those values. (Forcing subjects to ask for information is a practical way to measure their information search and draw inferences about their thinking processes.) Utility-maximizing subjects should ask to see all the information but most subjects did not (Payne 1976). Instead, subjects often chose a single attribute, such as rent, then eliminated apartments with rent above some threshold and never asked to see the other attribute values for those eliminated apartments.

3. Prospect Theory

Strong evidence and an alternative "prospect theory" was offered by Kahneman and Tversky (1979). They replicated Allais's common ratio paradox and introduced others. Prospect theory has four important elements: an editing stage

in which rules either dictate choices or transform gambles before they are evaluated; choice of a reference point (from which gains and losses are measured); a "value function" over gains and losses; and a function that weights probabilities nonlinearly and applies the resulting "decision weights" to outcomes to evaluate gambles. Each element in the theory is derived from experimental evidence.

The idea that people value changes from a reference point, rather than wealth positions, is an old one (e.g., Markowitz 1952). It extends to the financial domain the widespread evidence that in many psychophysical judgments, like brightness and heat, people are more sensitive to changes from adapted levels than to absolute levels (Helson 1964). As many people have noted, there is no axiom in EU implying that wealth positions are valued rather than changes in wealth, but it follows from the integration of assets and could be viewed as a basic principle of rational choice (like "description-invariance" and some other principles described later). Furthermore, it is easy to construct examples in which an EU maximizer will make consecutive choices which are suboptimal (compared to simultaneous choices) if she values changes rather than wealth positions (e.g., Tversky and Kahneman 1986, 255-256).

Kahneman and Tversky presented new data suggesting the value function has two important properties: (1) it is steeper around the reference point for losses than for gains ("loss-aversion")¹⁹; and (2) risk attitudes "reflect" around the reference point—the value function is concave for gains (risk averse) and convex for losses (risk seeking). Fishburn and Kochenberger (1979) also reviewed published studies showing reflection. Like the existence of a reference point, reflection can be interpreted as a psychophysical phenomenon, diminishing marginal sensitivity (marginal gains feel less and less good, marginal losses feel less and less bad).

The decision weight function in prospect theory is akin to earlier measurements of subjective weights of probabilities, by Edwards, Preston, Baratta, Mosteller, and Noges, et al. The data suggest low probabilities are overweighted and high probabilities are underweighted, as shown in Figure 8.6, with a crossover point roughly between .1 and .3. Underweighting of high probabilities implies a "certainty effect," in which special weight is given to certain outcomes compared to slightly uncertain ones (i.e., the decision weight function is convex and steep near one).

4. Elicitation Biases

Many researchers discovered systematic biases in elicitation of utility functions. In the chained certainty-equivalence technique, a value of p is chosen and people are asked for X' such that $X' \sim pH + (1-p)L$ (where H are high and low amounts), or X' such that $K' \sim pH + (1-p)X'$, etc. (Karmarkar 1974) and McCord and de Neufville (1983) found that using higher values of p in a chained procedure yielded more concave utility functions. Hershey and Schoemaker (1980) found that subjects more strongly preferred a loss of \$10 to a gamble (.01, -\$1,000) when it was called an insurance premium than when it was called a loss.

f199t] corroborated this finding in an interesting experimental study ofne<lgng-) Hershey and S.:hoemaker (1985) also found that utility functions elicited using probabmt)' and certainty equivalents were systematically different !cf. Johnson and Schkade 1989), violating the presumption that utility is invariant to the pteOe- dure used to elicit it (see section UIJ later), Hershey, Kunreuthoc, and Schoe- maker (1982) i,ummarited many of these elicitation biases and others.

Wolf and Pohlman (1983) elidted parameters of a specific kind of utility func- tion from a Treasury bin dealer, by di<:liting cenainty equivalents for several hypothetical gamb over hlll wealth. Then they timated the same parameters: using the dealer's actual bids (combined wrh the dealer's forecasts of the resale price of the bids). The utility functions: derived in these two ways were similar in foml (decreasing abwlute, roughly constant proportional risk ave.rasion) but the degree of risk aversion evldem in btds <was much larger than in the hypothetical choice&. (His bids would have been four times as large as lhey actually were if he had bid according to the utility function derived from hypo!het!cal choice.5.) The study is. oot condruive evidence of a hypothetical: real difference, because that difference was confounded with the method by which utility was elicited {cer- tainty•equivalence s. actual bids). However, the differeoce 5l.lggests caution in extrapolating from a utility function measured one way. to its application in an- other domain.

D. Generalizations of Expected Utility and Recent Tests

By the mid-1970 or iio, ooverall developmoots !nad convinced many hers that it was time to take altcmru:tives to EU seriously. Important milestones were grudging acceptance of the power of Allai.s's exarnplei: !lle ubiquity of EU viola- tions in choices, process data, and elicitation procedures: the elegaoce of Katme- man and Tversky's batch of new paradoxes; and Machina's (1982) assimihltion of some of the empirical eviOOnce against EU and imrodoction of sophistk.JJed tools for doing economic theory without the independence axiom:¹⁰

The anomalies motivated theorists to propose genetalitations of expected util- ity in which axioms are weakened or replaced. Most of the theories weaken i- pendence, but theorists have explored generalizations of other axioms too.⁴

Several recent empirical studies test these generalizations of Et.: Tv.ill first describe some of the theories, then describe tesl, of various theories in some delall (see also Camerer 1992a). Table 8.1 Mimmarizes several theories and their predictions about the shape of indifference curves in the Marschak-Machina trian- gle diagram,

Before reviewing !he various theories, a note about the modern influeoce of Allais and his European colleagues is appropriate (see Hagen 19'J1). Allais (1979) himself felt there were two main sources of EU violations: a certainty effect, and the fact !hat EU expressed aversion to risk only indirectly, through curvature of the utility function. He proposed a "neo-Bemoullian" model that presumes a car- dinal utility function ofoutcomes. obeys stochastic dominance, and assumes peo- ple choose gambles according to both the expectation and the variance of the

gamble's utilitiQ. (Higher expectation, and lower variance, are preferred,) Hagen (1%9) proposed a similar model in which positive skewness of utility is preferred as well.

These contributions got relatively linle attention in the United States and En- gland, for both sociological and scientific reasons. Their articles are blunt!v „criti- cal of EU {and of some other alternative theories); I suspect many Arimll readers are pl!t oft by the critical tone. Most of the work & published in book -chapters or in journals like Theory rul Decisicm and Journal of &-onomic Psy- chology, which are more widely read in Europe than ill the United States. Mo<it unponamly, !he Allais and Hagen formulations have free functions thal seem to be especially difficult to measure and test. FOExample, one cannm easily concoct a paradox like Allais's to telt Allais's own theory, by using pairwise chokes thal hold cOO-stant the influence of statiBtiral moments of flu: utility of gambles, with- our knowing the underlying utility function. (And the utility function cannot be easily measured using eertainty-equivalents. as in EU, because choices are u- sumed to depend on expectation and on higher moments.)

I. Predictions of Generalized EV Theories

Wt>ighled wility theory (Chew and MacCrimmon 1979a; Chew 1983) aSSLlmes a weakened form of independence. As above, denote gambles by capital letters (X, Y, Z), compound gambles. by $pX + (1 - p)Z$, and preference for X over Y by $X \succ$. Their axiom is:

Weak Independence. If $X \succ Y$, then for all p in $(0, 1)$ there exists a unique q in $(0, 1)$ such that $pX + 0 - p;Z \succ qY + (1 - q)l$ for all Z .

Weak independence in combination with the other EU iuioms implies a represen- tation of the form

$$(3) \quad U(1_p x_i) = \frac{1P_{,w}(x M t'')}{!p; > v(x, J)}$$

In weighted utility, indifference curves are straight lines that meet at a point out- -J the triangle. (If $w(XM) < 1$, for instance, in the dorruln of three outcome gambles where X_M denotes the middle outcome, then the curves will look like those in Figure 8.8.)

There is a generalization of f weighted utility called "skew-symmetru: bilinear" (SSB! utility theory (Fishburn 1982, 1983, 1988). In SSB utility. preferences are represented by a function of both lotteries, (X, Y) . $X \succ Y$ if and only if $c/(X, Y) > 0$; $4(X, \Pi = 0$ implies $X \sim Y$.

SSB re.wlu frl.)ill replacing EU with a weakened fmm of independence calle<i -symmetry." adding a betweenness axiom {sec below), retaining „umplemess and oontimmy, and abandoning transitivity. When transitivity ls added back in, surprisingly, SSB reduces to weighted utility,

Regret lh&ary (Bell 19S2; Loomes and Sugden 1982, J987a) genccraiu:es SSB further by exiending it lo droi_e„ between lotteries with oorrelated ootcofi:es




figure 1.8. indifference
curves that fan OHL

(i.e., where the outcome from one lottery depends statistically on what the other lottery's outcome was). Some tests are described in section H.2 below.

Implicit EU. A weakened form of SSB or weighted utility, called "Implicit weighted utility" by Chew (1989), or "implicit EU" by Dekel (1986), depends on a weakened form of independence called "betweenness."

Betweenness. If $X \succ Y$, then $X \succ pX + (1-p)Y$ for all p in $(0, 1)$ (or $X \sim Y$ implies $X \succ pX + (1-p)Y \sim Y$ for all p in $(0, 1)$).

Betweenness implies neutrality toward randomization among equally-good outcomes. It yields an implicit utility representation of the form

$$(4) \quad U^* = U(\sum p_i x_i) = \sum p_i u(x_i, U^*)$$

The utility function $u(x_i, U^*)$ denotes the utility of an outcome x_i , but the utility function used to value x_i depends on U^* . (Utility is therefore defined implicitly: U^* is an expected utility that is calculated using a utility function that depends on U^* .) In implicit utility, indifference curves are straight lines (which are positively sloped, and don't cross over), but they are not necessarily parallel as in EU.

Betweenness can be violated in either of two ways. If preferences are strictly "quasi-convex," then $X \sim Y$ implies $X \succ pX + (1-p)Y$ (people are averse to randomization). If preferences are strictly "quasi-concave," then $X \sim Y$ implies $pX + (1-p)Y \succ X$ (people prefer randomization). In the triangle diagram, quasi-convex (quasi-concave) preferences imply concave (convex) indifference curves.

Chew, Epstein, and Segal (1991) propose a weakened form of betweenness, called "mixture symmetry": If $X \sim Y$ then $pX + (1-p)Y \sim (1-p)X + pY$. Together with other axioms, mixture symmetry implies that preferences switch from quasi-convex to quasi-concave, or vice versa, as gambles improve (in the

set of stochastic dominance). (Curves then switch from concave to convex, or vice versa, as one moves northwest in the triangle diagram.)

In weighted utility theory, indifference curves may "fan out," getting steeper as one moves from the lower right-hand corner (or southeast) to the upper left-hand corner for northwest), as in Figure 8.8. (They can also "fan in," getting less steep to the northwest.) Straight indifference curves correspond to more risk averse behavior: the steeper the curve, the more P_u a person demands in order to accept a unit increase in P_L . Therefore, fanning out occurs if people are more risk averse toward gambles that are better (in the sense of stochastic dominance). Machina (1982) showed that several empirical anomalies could be explained by the *fanning out* hypothesis (without requiring the further assumptions of weighted utility theory).

Gul (1991) proposes a theory incorporating *disappointment*. The intuition is that the probabilities of outcomes below and above a gamble's certainty-equivalent are weighted differently to reflect the additional satisfaction from having "beaten the odds," or unhappiness from having lost to them. His theory is a special case of implicit EU that satisfies betweenness and allows curves to fan in (for better gambles in the northwest part of the triangle) and fan out for worse gambles (in the southeast).

In *expected utility with rank-dependent probability weights* cumulative probabilities are weighted, and the utilities of outcomes are weighted by the differential in the weighted cumulative probability (Quiggin 1982; Segal 1987b, 1989; Yaari 1987; Green and Juven 1988). (This procedure ensures that stochastic dominance is never violated, which can happen if probabilities are weighted separately.) The weight of an outcome depends on its probability and its rank order in the set of possible outcomes. Suppose we rank outcomes from high to low, $x_1 \succ x_2 \succ \dots \succ x_n$. Then the functional form for rank-dependent utility is

$$(5) \quad U(p, x) = \sum_{i=1}^n w(p_i) [g(p_i + p_{i+1} + \dots + p_n) - g(p_i + p_{i+1} + \dots + p_{i-1})]$$

Note that if $g(p) = p$, the bracketed expression reduces to p_i and equation (5) reduces to EU. In rank-dependent theory, indifference curves will not be straight lines unless the probability transformation function $g(p) = p$; otherwise, they are curved in a way which depends on $g(p)$ (Roell 1987; Camerer 1989a). A convex $g(p)$ expresses risk aversion in a novel way, by underweighting the probabilities of the highest-ranked outcomes and (since weights sum to one) overweighting the lowest-ranked outcomes; similarly, concave $g(p)$ expresses risk preference.

Lottery-dependent utility theory (Becker and Steiner 1987) assumes only stochastic dominance, ordering, and continuity. The theory is quite general:

$$(6) \quad U(p, x) = \sum p_i u(x_i, c)$$

where $c = h(x_i/p_i)$. Becker and Sarin (1987) suggest an exponential form for the utility function $u(x_i, c)$ which makes the theory more precise and useful. Indifference curves fan out in the exponential form and lottery-dependent preferences are quasi-convex (curves are concave) if $h(x)$ is concave (Camerer 1989a, 73),

In prospect theory (Kahneman and Tversky 1979), indifference curves may vary with the choice of reference point, but for testing purposes it is useful to consider a variant of prospect theory which assumes that the reference point is current wealth. Then the "value of a gamble is simply

$$(7) \quad V(p_1x_1 + p_2x_2 + P, ?) = \pi(p_1)\pi(x_1) + \pi(p_2)\pi(x_2)$$

In general, $p_1 + p_2 \leq 1$ (and $1 - p_1 - p_2$ is the probability of getting nothing). If $p_1 + p_2 = 1$ and both x_1 and x_2 have the same sign, then the value is

$$(7') \quad V(p_1x_1 + p_2x_2) = \pi(x_1) + \pi(p_1)\pi(x_1) - \pi(x_2) = (1 - \pi(p_1))\pi(x_1) + \pi(p_1)\pi(x_2)$$

The shape of the indifference curves determined by equations (7) and (7') depends on $\pi(p)$, but they will certainly be nonlinear unless $\pi(p) = p$. If $\pi(p)$ is most nonlinear near 0 and 1 (as originally proposed by Kahneman and Tversky) the curves will change slope and shape dramatically at the edges. As a result, choices between gambles inside the triangle will violate EU less than choices involving gambles on the edges.

Prospect theory was originally restricted to gambles with one or two nonzero outcomes. In "cumulative prospect theory," Tversky and Kahneman (1992) extend the theory to gambles with many outcomes (including continuous distributions) using a rank dependent form like equation (5) (cf. Starmer and Sugden 1989a, 99-100). The difference is that, in their formulation, probabilities of gains and losses can be weighted differently. They also suggest a parsimonious, one parameter probability weighting function⁴³:

$$(8) \quad g(p) = p^\alpha \ln p Y + (1 - p)^\beta Y^{1-\beta}$$

Other Theories

Handa (1977) proposed a precursor to prospect theory in which probabilities were weighted nonlinearly (see, much earlier, Edwards 1954:3). Kannarkar (1978) and Viscusi (1989) propose specific functional forms for probability weights $\pi(p)$. Leland (1991) suggests an "approximate" EU in which people lump outcomes (and possibly probabilities) into discrete categories, making their utility functions discontinuous (cf. Rubinstein 1988). One special feature of this approach is that it allows convergence to EU with experience, since experience permits finer grained categorization of outcomes (cf. Friedman 1989).

2. Empirical Studies Using Pairwise Choices

There are many recent studies using pairs of choices to test EU and its generalizations. Each of the theories described above predicts different indifference curve shapes in some part of the triangle. The predictions are summarized in Table S1. By choosing pairs carefully from throughout the triangle, each theory can be tested against the others. The first study of this kind was done by Thaler and Waller (1986). I will describe their design because it is typical and raises basic methodological questions.

Table S1. Predictions of Competing Theories about Properties of Indifference Curves

Functional Form for U^*				Properties of Curves			
Theory	Continuous, $U^*(F(x))$	Discrete, $U^*(\sum p_i x_i)$	Restrictions	Straight Lines?	Fanning Out?	Fanning In?	Miscellaneous
Expected utility	$\int u(x) dF(x)$	$\sum p_i u(x_i)$		Yes	No	No	Curves parallel
Weighted utility	$\frac{\int u(x) \pi(x) dF(x)}{\int \pi(x) dF(x)}$	$\frac{\sum p_i \pi(x_i) u(x_i)}{\sum p_i \pi(x_i)}$	$w(X_{n+1}) < 1$ $w(X_{n+1}) > 1$	Yes Yes	Yes No	No Yes	Curves meet in a point
Implicit expected utility	$\int u(x, U^*) dF(x)$	$\sum p_i u(x_i, U^*)$		Yes	Maybe	Maybe	Only testable property is betweenness
Fanning-out hypothesis	$\frac{-U''(x, F)}{U'(x, F)} \geq \frac{-U''(x, G)}{U'(x, G)}$ if $F(x) \leq G(x)$ for all x			Maybe	Yes	No	Movements to northwest cause steeper slopes
Lottery-dependent utility	$\int u(x, c) dF(x)$ $c_F = \int u(x) dF(x)$	$\sum p_i u(x_i, c_F)$ $c_F = \sum p_i u(x_i, p_i)$	h concave h convex	No No	Yes Yes	No No	Curves concave Curves convex
Prospect		$\pi(p_1)u(x_1) + \pi(p_2)u(x_2)$ $(1 - \pi(p_1))u(x_1) + \pi(p_2)u(x_2)$ $p_1 + p_2 = 1$ and $0 < x < y$ or $y < x < 0$		No	Lower edge	Left edge, hypotenuse	Parallel along $P_F = (1 - P_F)/2$
Rank-dependent utility	$\int u(x) dG(F(x))$	$\sum_{i=1}^n u(x_i) [g(\sum_{j=1}^i p_j) - g(\sum_{j=1}^{i-1} p_j)]$	g concave g convex	No No	Lower edge Left edge	Left edge Lower edge	Parallel along hypotenuse




Figure 8.9. Chew and Waller's HLO structure and results. O, I, and H denote the four gamble pairs, and L is the common ratio test. The number of subjects (total 56) choosing each gamble in each pair is shown in parentheses. Dollar amounts are in thousands. Source: Chew and Waller 1986.

Chew and Waller used an ingenious, compact set of choices (originally developed by Chew and MacCrimmon 1979b) called the "HLO structure." Figure 8.9 shows one set of four HLO pairs they used (their context 2b), drawn in a triangular diagram. The three outcomes in the gambles, in thousands, are $X_H = \$20$, $X_I = \$10$, and $X_L = 0$. Every subject chose one gamble from each of the four pairs, four choices in all.

The pairs were picked to efficiently detect several different effects, with only four choices: the common ratio effect (pairs O-L and O-H), the common consequence effect (I-L and I-H), and a test of the betweenness axiom (I-O). Figure 8.9 shows the number of subjects (out of fifty-six) who chose the gambles in each pair. Subjects were not allowed to express indifference and didn't play any gambles.

The common ratio effect occurred because thirty-five subjects chose risk-aversely in the O pair (thirty-five picked (0, \$10) and twenty-one picked (\$20, \$10)) but only seventeen did so in the L pair. A weak common consequence effect occurs because twenty-three subjects chose risk-aversely in the I pair and fifteen did in the L pair.

The framing hypothesis predicts that people become more risk-averse, and indifference curves fan out, with movement toward the northwest corner. The prediction appears false in the Chew and Waller data: fewer subjects chose risk-

aversely in the R pair (twenty-five) than in the O pair (thirty-five),⁴⁴ and about the same number chose risk-aversely in the I pair (twenty-three).

Pairs I and O test whether the betweenness axiom holds (since the inner gamble {\$.75, \$10; .20, \$20} in pair I is a probability mixture of the outer gambles (0, \$10) and (\$20, \$20)). Betweenness is violated because 35 subjects chose (0, \$10) in the O pair but only 23 chose it in the I pair.

*A Methodological Transformation:
Between-Subjects vs. Within-Subjects Analysis*

Psychologists call the sort of analysis expressed by Figure 8.9 "between-subjects": the fractions of subjects behaving in a particular way in two different settings are compared (but the subjects may be different in the two settings). In a "within-subjects" analysis, the fact that a single subject made choices in two or more settings is exploited. Within-subjects tests are always more statistically powerful, but they run a certain risk: presenting several stimuli to a single subject could conceivably induce her to behave more consistently than she would if she saw the stimuli one at a time. For example, in tests of whether two different descriptions of an equivalent problem elicit the same choices ("framing effects"), presenting the two problems one after the other might cause a subject who recognizes the equivalence to guess that the experiment tests for consistency, and respond accordingly. (The same issue arises in some of the judgment research reported in section U above.)

Many people think within-subjects analysis is the only proper analysis in choice experiments, because BU requires consistency of individual preferences. But, of course, between-subjects tests are equally legitimate (though less powerful) if the subjects in different groups can be presumed to have the same distribution of tastes, up to sampling error, because they were drawn from a single population.

Chew and Waller (1986) did a within-subjects analysis by counting how many subjects exhibited a particular pattern across choices. For instance, in the O and L pairs (the common ratio test) twelve subjects chose the less risky gamble in both pairs and sixteen subjects chose the more risky gamble in both pairs, so twenty-eight of the fifty-six subjects satisfied it. Of the twenty-eight who chose inconsistently, twenty-three (82 percent) chose risk-aversely in O and risk-preferentially in L, manifesting the standard common-ratio pattern.

*A Further Methodological Offshoot:
Judging Violation Rates*

Is a 50 percent rate of EU violation (twenty-eight of fifty-six chose inconsistently) large or small? If we allow random error in expression of preferences, then the appropriate benchmark for violation rates should be the fraction of people who switch their choices when making the same choice twice (called "reliability" in psychometrics). Chew and Waller did not measure the fraction of random switching but other studies suggest that percentage is about 25 to 35 percent for

-hokes like these (Starmer and Sugden 1987b; Camerer 1989a). Using that ndlmark, a χ^2 -test shows that the fraction of inconsistent subjects in the 0-L pairs (28/56) is much too high to be a chance deviation from random switching. Another way to test whether violations are systematic is to compare the fraction of inconsistencies in both directions. In the 0-L pairs, twenty-three of twenty-eight switched in one direction and five in another, an asymmetry unlikely to occur by chance.

Of course, the violation rate is likely to be sensitive to the gamble pair chosen. Two familiar gambles will have a violation rate close to 50 percent. That simply means that if your goal is to reject theories, using such a pair of gambles will require a large sample to detect true systematic violation statistically.

A Third Methodological Observation: Incentives

In Chew and Waller's study, choices were entirely hypothetical. In most other studies subjects played one of the gambles, they chose. The common procedure of randomly choosing one of several gambles to play has been tested by Camerer (1989a, 82) and Starmer and Sugden (1991a): it elicits roughly the same preferences as when subjects make only one choice and play the gamble they picked.

Several studies have compared hypothetical choices with real choices (in which one choice was played). They found either no effect or a slight tendency for playing gambles to yield more risk aversion (Edwards 1953; Becker, DeGroot, and Marschak 1963; Camerer 1989a; Battalio, Kagel, and Jiranyakul 1990; Hogarth and Einhorn 1990; Schoemaker 1990; cf. Slovic 1969).

Harrison (1990) reported the only evidence that actually playing a gamble substantially reduced the rate of EU violation (Battalio et al. [1990] and Camerer [1989a] report no such effects). In his common-ratio experiment, seven of twenty subjects violated EU when choices were hypothetical and three of twelve (different) subjects violated EC when choices were real. The difference is not significant at conventional levels ($p = .15$).

In a contrasting study, Kachelmeier and Shehata (1992) elicited certainty-equivalents from Chinese students (using the BDM procedure), for gambles with high or low outcomes. Subjects played every gamble, and payoffs ranged from 1 to 100 yuan, which are substantial amounts for students who spend about 60 yuan per month. Certainly, Chinese students exhibited dramatic overweighting of low probabilities at both high and low payoff levels (e.g., the certainty equivalent of (p, X) was about twice pX for $p = .05$ or $.10$); similar patterns were observed with Canadian students playing hypothetical gambles. Risk aversion among the Chinese students was also greater for high payoffs than for low payoffs. Thus, the overweighting of low probability observed in so many experiments (which leads to many EU violations) is present when very large payoffs are used, but large payoffs also induce some differences in risk aversion.

In experiments where gambles are played, the amount of expected value (EV) subjectively lost by violating EU is usually small. Unfortunately, there is no simple way to design experiments that have large expected value penalties for violating

EC but that can also sharply distinguish EU from competing theories (including EV). Since the effect of dominant payments is difficult to determine empirically, a pragmatic way to approach the problem is to spell out and test whatever theory or intuition underlies the claim that higher payoffs would reduce EU violations (such as Smith and Walker's [1993] analysis of "decision cost"). The intuition seems to be that subjects will calculate more, think harder, or somehow see the appeal of axioms when they are faced with larger stakes. But in experiments where subjects were told expected values (Lichtenstein, Slovic, and Zink 1969) or given written arguments explaining the independence axiom, it is easy for them to think harder—EU violations were not reduced (Slovic and Tversky 1974). Indeed, some studies suggest that the main effect of paying subjects is a reduction in variance of their responses, which increases the statistical significance of EU violations (Harless and Camerer 1994).

The effect of paying subjects is likely to depend on the work they perform. In many domains, paid subjects probably exert extra mental effort, which improves their performance, but in my view choice over money gambles is not likely to be a domain in which effort will improve adherence to rational axioms. Subjects with well-formed preferences are likely to express them truthfully, whether they are paid or not. If their preferences are not well formed, it seems unlikely that subjects would be both sophisticated and lazy enough to make an expected utility calculation when they are paid, but not when choices are hypothetical. (Furthermore, if payment does induce more formal reasoning, it is likely to be expected value maximization, not EU maximization.)

Other Studies of Common Consequence and Common Ratio Effects

Chew and Waller's data replicate the common ratio and common consequence effects in the southeast corner, but they show little fanning out in the northwest corner. Camerer (1989a) found the same general pattern. In a replication using the Allais payoffs, Conlisk (1989) found fanning in in the northwest corner and Prelec (1990) found dramatic fanning in (55 percent of subjects) close to the fowling edge and southeast corner. (Their subjects did not play gambles.)

Starmer and Sugden (1989a, 1989b) studied a wide variety of common ratio problems. The rate of EU violations was only significant in three of fifteen pairs. They found strong fanning out along the lower edge, and some fanning in along the left edge. Starmer (in press) found common consequence effects, showing fanning in rather than the more typical fanning out, especially along the lower edge.

Battalio, Kagel, and Jiranyakul (1990) observed fanning in along the lower edge, in common ratio problems where gambles had strictly positive payoffs. Fanning in also appeared in the northwest corner with gains and the southeast corner with losses.

The patterns in these studies are complicated. Most of the differences in results can probably be traced to small differences in the particular gambles being stud-

ied. It would be useful to have a composite picture of the indifference curves revealed by all the studies, but nobody knows how to create such a composite (but see Harless and Cainera 1994, discussed below). It does seem that any composite picture is likely to cast doubt on the generality of fanning out. For example, the fanning out that is observed in the Allais problems along the lower edge of the triangle is generally reduced or reversed-sometimes fanning in on the left edge. The fanning out hypothesis was suggested by MacLennan (1982) to explain evidence that had come primarily from the lower right corner of the triangle (e.g., the Allais paradoxes illustrated in Figure 8.7). His inference from the data is the entire triangle was ingenious, but apparently not quite right.

Two classes of theories can account for most of the mixed fanning evidence from common ratio and common consequence studies. One class of theories posits mixtures of fanning in and fanning out (Neilson in press) or derives them from axioms (Gui [1991], which adds only one parameter to EC). Another class weights probabilities nonlinearly, which generates indifference curves that fan out or in depending on the weighting function.

Evidence of Betweenness Violation

Camerer and Ho (1994) reviewed nine studies in which the betweenness axiom was tested. (Tests of betweenness efficiently test several theories that assume if w -weighted and implicit EU, disappointment, SSB, and regret-in one fell swoop.) In these studies, betweenness is frequently violated but the pattern of violation is complicated. Gambles with gains generally show quasi-convexity (concave indifference curves) except close to the lower and left edges (El Bernasconi 1994); loss gambles show the opposite. The fact that convexity reflects for gains and loss implies people seem to weight gain and loss probabilities differently (See Tversky and Kahneman 1992). Evans (1992) found comparable degrees of betweenness violation when gamble valuations were elicited with the BDM procedure, and with second- and fifth-price sealed-bid auctions (see chapter 7).

We fitted data from all the studies to a variant of EU with nonlinear rank-dependent probability weights (using the weighting function (8)). The maximum-likelihood coefficient estimates γ are around $\gamma = .60$, which implies a nonlinear weighting roughly like that pictured in Figure 8.6, with probabilities under .30 overweighted and probabilities above .30 underweighted. The data therefore reject betweenness (which requires $\gamma = 1$) and corroborate both guesses from the earliest studies from the 1940s and sharper estimates from studies almost fifty years later (e.g., Tversky and Kahneman [1992] estimate $\gamma = .61$ for gains and $\gamma = .69$ for losses). But the very latest evidence (Wu 1994) shows some violations of the rank-dependent approach which should renew interest in preference editing rules.

The most prominent theory that can account for the mixed fanning reported in the last section, Gui's (1991) disappointment theory, assumes betweenness. The fact that betweenness is often violated casts doubt on Gui's theory and leaves the nonlinear-weighting theories as the best available account of both complex fanning out patterns and betweenness violation.

Evidence of Better Conformity to EU Inside the Triangle

One of the most interesting and robust effects in the new wave of tests is that EV violations are much smaller (though still statistically significant) when subjects choose between gambles that all lie inside the triangle. Conlisk (1989), Gigliotti and Sopher (1990), Camerer (1992b), and Harless (1992b) all discovered this phenomenon independently.

The shrinkage of EU violations inside the triangle does not vindicate EU. Inside gambles have probabilities p_L , p_M , and p_H which are all nonzero; edge gambles have at least one of the three probabilities equal to zero. The fact that violations disappear when moving from the edge of the triangle inside suggests they are probably due to nonlinear weighting of probabilities near zero (as the rank-dependent weighting theories and prospect theory predict).⁴¹ Much as Newtonian mechanics is an adequate approximation at low velocities, but relativistic mechanics is accurate at all velocities, the linear weighting of probabilities imposed by EU may be an adequate approximation when outcome probabilities are not too low or high. Morgenstern (1979) drew the analogy with mechanics, and anticipated such a conclusion: "(T)he domain of our axioms on utility theory is also restricted. . . . For example, the probabilities used must be within certain plausible ranges and not go to .01 or even less to .001, then be compared to other equally tiny numbers such as .02, etc." (178).

Whether nonlinear weighting of low probabilities affect choices among natural gambles with many outcomes is a fundamental empirical question. In any case, theories that assume nonlinear probability weights are here to stay: they are the only theories that can explain evidence of mixed fanning, violation of betweenness, and approximate EU maximization inside the triangle.

Informal and Formal Summaries of Evidence

There are at least two ways to summarize this evidence on experimental testing of various utility theories. Camerer (1992a) summarizes the evidence informally in a list of "stylized facts." Many of the facts can be expressed as shapes of indifference curves in the Marschak-Machina triangle. The crucial facts appear to be that indifference curves vary in (local) slope from risk averse to risk seeking; indifference curves fan in and out in a systematic, complex pattern (strict fanning out can be rejected); betweenness violations imply that indifference curves are not straight; indifference curves are more nearly parallel inside the triangle than on the interior and curves for gains and losses reflect around a 45 degree line.

The right kind of nonlinear probability weighting function can have all these properties. Figure 8, 10 shows curves plotted by Tversky and Kahneman (1992), for gambles over gains and losses, for cumulative prospect theory using their one parameter weighting function. The dotted line shows pairs with equal expected value. The plotted indifference curves have all the properties mentioned in the previous paragraph: They are sometimes flatter (more risk preferring) and sometimes steeper (more risk averse); they both fan in and out; they are not straight; they are more curved near the triangle boundaries; and they reflect from gains (panel a of Figure 8JO) to losses (panel b),




Figure 8. Empirically derived indifference curves of cumulative prospect theory for EU. (a) gains and (b) losses. Source: Tversky and Kahneman (1991).

A Statistical Summary

Harless and Camerer (1994) summarize the evidence more formally. In studies with several pairwise choices, like those discussed above, theories are usually judged by what fraction of the choice patterns they predict correctly. For example, in the Chew and Waller (1986) study people made one choice from each of forty pairs. There are $2^{40} = 16$ possible patterns of choices. EU predicts only two of those patterns, and excludes fourteen of them. Most statistical tests compare the percentage of patterns that are predicted—two of sixteen in the EU example, or 12.5 percent—with the percentage of subjects actually choosing those patterns (25 percent in their data) and perform some hypothesis test. Harless and I show that this method ignores useful information. If one adds the possibility that people choose erroneously with some probability, then EU predicts that some of the fourteen excluded patterns are more likely than others. Considering all fourteen patterns equally likely mistakes ignores information in the relative frequency of the fourteen patterns that could be used to judge the theory.

We developed a method that uses all the data and generates a chi-squared statistic testing each of several alternative theories. Since chi-squared statistics can be added across independent experiments, the results of several studies can be easily aggregated using our method. The result is an "efficient frontier" of theories that are most accurate (best-fitting), given the number of patterns they allow (most parsimonious). A compilation of twenty-three data sets, from a total of 2,000 subjects making 8,000 choices, shows that the following theories are on the efficient frontier: mixed fanning, prospect theory, EU, and EV.

The difference between boundary gambles and interior gambles emerges strikingly from the compiled data. For interior gambles, there is a broad range of

parsimony-accuracy tradeoffs for which EU is the best theory. But for boundary gambles, EU is preferred. (If one is willing to trade off enormous predictive accuracy for parsimony, preferring EU to a generalization, then one should go even further and choose EV over EU.)

A Field Observation and Experiment

Marshall, Rikhard, and Zarkin (1992) studied an interesting field situation. They begin with a fact: in 1983 in the United States, twice as many accidents (per trip) occurred during evening rush hours (4–7 P.M.) than occurred during morning rush hours (7–10 A.M.). The choice of whether to wear a seat belt is a choice between two risky accident distributions: wearing a seat belt is, roughly speaking, a risk-averse choice: it reduces the probability of having a serious accident but, curiously, slightly raises the probability of having some kind of accident. Since accident rates are lower in the morning, the safer morning choice is a better pair of gambles than the evening choice (in the sense of stochastic dominance), so the fanning-out hypothesis predicts that people will behave more risk-aversely—wearing their seat belts—in the morning than in the evening. They do: more people wear seat belts in the morning than in the afternoon (15.4 percent vs. 13.9 percent). Marshall et al. show that observations of increased seat belt use in the morning, when accident rates are lower, reject the EU hypothesis and support fanning out.

Their conclusion is debatable because strong, if realistic, assumptions must be made to conduct the statistical tests and an experiment with students, in which choices and accident rates were prefiltered abstractly, did not corroborate the field observation. (Students obviously made the choice corresponding to wearing seat belts in both time periods.) Nonetheless, their method is a unique illustration of how field data could test alternative choice theories.

Cicchetti and Dubin (1994) did a similar study (but with no experimental component): they fit an EU model to consumer decisions to buy insurance against the cost of repairing wires in their phones. The EU model fits reasonably well. They cannot reject the hypothesis that consumers estimate repair probability accurately, but the data also show some modest support for nonlinear weighting of probabilities.

3. Empirical Studies: Measuring Indifference Curves

The most direct approach to testing generalized utility theories through their predictions about the shape of indifference curves is by directly assessing indifference curves. Hey and Strazzer (1989) estimated curves by eliciting a set of equally good gambles, using a series of "lottery-equivalent" (McCord and De Neufville 1986) indifference judgments. Some curves crossed (reflecting intransitivity) or had negative slopes (reflecting stochastic dominance violations). The fit to the elicited curves generally obeyed EU (parallelism could not be rejected), but the statistical power to reject EU was limited by the number of curves and the number of points on each curve.

Abdellaoui and Munier (1994) did a similar curve-fitting exercise. They found evidence of fanning on using common-ratio type choices, and some fanning out along the lower and left edges of the triangle. Curves were roughly parallel in the triangle interior.

Hey and Di Cagno (1990) used choices scattered throughout the triangle, rather than indifference judgments, then fitted probit models to the choices. They strongly reject a special form of regret theory, but accept a more general form.

4. Empirical Studies Fitting Functions to Individuals

A few researchers have estimated functional forms or parameter values for individual subjects (Currim and Sarin [1989] for prospect theory, and Currim and Sarin [1990] and Daniels and Keller (1990) for lottery-dependent utility). In a typical experiment, a subject makes a series of choices or states certainty-equivalents for several gambles. To fit EU, for example, a particular form of utility function is assumed, such as $u(x) = 1 - e^{-cx}$, and a best-fitting value of c is estimated (by minimizing squared deviations from the stated certainty-equivalents, for example). Alternative theories have more degrees of freedom and require estimation of more parameters. Then the fitted parameter values are used to predict choices in a hold-out sample of fresh choices. The method in Hey and Orme (1994) is similar except goodness-of-fit tests are used (EU is tested as a restriction on non-EU theories) instead of making predictions for a hold-out sample.

The results of these studies are somewhat discouraging for new theories. Alternative theories always fit initial choices better than EU, since they have extra degrees of freedom, but they do equally as poorly as EU predicting new choices, getting only about 60 percent right. (In the Hey and Orme study, the EU restrictions cannot be rejected for about half the subjects; for the other half, prospective reference theory or the one-parameter version of cumulative prospect theory appear to fit best.) These early results serve as a reminder that many subjects obey EU, and the lean functional form in EU is more statistically robust to estimation error than more complex functional forms are (e.g., Carbone and Hey [1994]). Of course, these are initial efforts. Refined techniques and larger samples might work better and enable more precise estimation of non-EU functions and parameters, and better predictions of hold-out sample choices.

Lattimore, Baker, and Witte (1992) also fit variants of EU with nonlinear probability weights to choices by students and prisoners (involving gambles over crimes and sentences). They assumed a power utility function $u(x) = x^\alpha$, and a two-parameter probability weighting function of the form:

$$(9) \quad w(p) = \frac{p^\alpha}{(p^\alpha + \beta p^\gamma)}$$

The parameter α expresses additivity of weights (weights add to one if $\alpha = 1$) and β expresses the degree of overweighting of certain probabilities.

Lattimore et al. found reflection effects between gain and loss gambles, and nonlinear weighting of probabilities (roughly similar to Figure 8.6; Camerer and

Ho [1994]; and Tversky and Kahneman [1992]). Their data indicate weighting of losses and gains are slightly different: loss weights are more likely to be subadditive. They did *not* fit parameters to one set of choices then predict new choices.

5. Cross-Species Robustness: Experiments with Animals

Besides the many animal experiments studying risk aversion, "consumer" choice, optimal foraging, etc., there are a few experiments testing whether animals obey expected utility. Rats exhibit the common ratio effect like people do (Battalio, Kagel, and MacDonald 1985). (The cross-species replication is especially notable because the rats respond only sluggishly to incentives--e.g., they only choose a stochastically dominant lever 55 to 90 percent of the time.) Rats also fan out and fan in, just as people do, in the northwest and southeast corners of the triangle, when they choose between levers that give "losses" (delays in dispensing food) (Kagel, MacDonald, Battalio 1990). Rats also fan in over gains (cups of water) in some parts of the triangle, and violate the betweenness axiom (MacDonald, Kagel, and Battalio 1991), though no attempt has been made to check whether people exhibit those patterns for comparable gambles.

So *all* the available evidence to date indicates rats exhibit the same EU violations people do.⁵¹ This cross-species generalizability is profound; it encourages the search for theories that can explain both human and non-human behavior in one fell swoop. My own view (which may be biologically naive), is that rats and humans are more likely to share misperceptions of probabilities than to share feelings of regret and disappointment. So if one prefers a common theory across species, the animal data shifts a bit more support to theories with nonlinear probability weights.

6. Some Conclusions from Recent Studies

The evidence collected in the last five years is as voluminous as the evidence gathered at any stage of testing EU. There are important lessons in the data.

Common ratio and common consequence effects are easy to replicate, but their strength varies across probabilities and outcomes. Fanning out is systematically violated for several kinds of gambles. Risk attitudes, and weighting of probabilities, appear to reflect around a reference point. There is an interaction between the degree of EU violations and the size and sign of outcomes (cf. Edwards 1954a)--i.e., violations are more frequent when outcomes are larger. Some models can account for the interaction by abandoning separability of outcomes and probability weights (e.g., Hogarth and Einhorn 1990; Becker and Sarin 1987), but the loss in parsimony may be too high a price to pay for better fit.

Many studies (beginning with Preston and Baratta 1948) are consistent with overweighting of low probabilities (below .2 or so) and underweighting of higher probabilities. Mixed fanning, betweenness violations, disappearance of EU violations inside the triangle, and replication of human results with animal subjects, might all be accounted for by nonlinear weighting. Some other studies, not mentioned above, also found evidence that probabilities between .3 and .8 are under-

weighted (Cohen, Jaffray, and Said 1985; Cohen and Jaffray 1988; de Neufville and Delquie 1988).

The few attempts to fit models for individual decisions suggest that more general theories fit better than EU (since they have more degrees of freedom) but are no better in predicting new choices. More studies of this sort are crucial for establishing whether the new theories can actually make better predictions than EU.

Among classes of theories, we can declare a few winners and losers in the empirical sense. Theories that incorporate nonlinear utility in probability—such as the rank-dependent approaches, particularly, and cumulative prospect theory—to have the necessary empirical properties (Figure 8.10) and, I think, prove relatively easy to work with formally. Betweenness-based theories (including implicit and weighted EU, and disappointment theory) have elegant formal properties but cannot accommodate apparent nonlinearity in probability adequately.

Finally, the continued use of EU can only be justified in two fairly narrow ways: first, EU is not so badly violated in choosing over gambles with the same set of possible outcomes (triangle interior), or with probabilities well above and below 1/2 though it is still statistically rejectable. Second, EU might be preferred in an application where parsimony is very highly valued compared to predictive accuracy (but even then, EU is often just as good).

7. Investments in Risky Assets

Loomes (1991a) conducted a novel experiment. He gave subjects lotteries over three events, denoted A, B, and O (a zero payoff), which occurred with probabilities p_A , p_B , and $1 - p_A - p_B$ (A was always more likely than B). Subjects could allocate £20 between the events A and B, in any proportion they liked. If an event occurred, they earned the amount of money they allocated to that event. (Subjects played one lottery.) For example, suppose $p_A = .6$ and $p_B = .4$. A person who is risk averse might allocate A = £10 and B = £10 (then A is certain to earn £10); a risk seeker might allocate A = £20, creating a .6 chance of winning £20. The money splitting task resembles allocation of wealth in a portfolio of risky assets; making A and B close to £10 is like buying bonds; making A close to £20 is like buying stocks.

An EU maximizer will divide the money to solve $t(A)u(£20 - A) = p_A/P_B$. A risk neutral or risk preferring person will put all £20 in the more likely outcome, A. A sufficiently risk averse person will put less than £20 in A.

This simple problem provides a remarkably powerful test of several alternative choice theories. Under EU, as p_A and p_B fall, the amount A should stay the same. If the ratio p_A/p_B is held constant. (Regret should not affect A either, in the display that was used.) Fanning out predicts the amount A will rise, and a restricted form of rank-dependent EU predicts it will fall. When $p_A = .6$ and $p_B = .4$, subjects put £13.15 in A on average. Half the subjects chose amounts proportional to the probabilities (A = £12, B = £8). As p_A and p_B fell, to .2 and .3, a third chose A = £10, and the average A fell to £11.58. The drop in A violates EU and fanning out.

No current theory explains all the observed behavior. Current theories assume the choices of A and B are induced from underlying preferences between gambles—i.e., a person chooses A = £12 and B = £8 if and only if $(p_A/12; p_B/8) \succ (p_A/A; p_B/20 - A)$ for all other values of A. Maybe subjects do not induce divisions from preferences; instead, they regard money splitting as akin to problem-solving and use a simple heuristic (such as $A/B = p_A/p_B$) that generates allocations that are inconsistent with complete pairwise preferences.

The money-splitting task vaguely resembles investment in risky assets. Kroll, Levy, and Rapoport (1988a, 1988b) studied asset portfolio problems more directly. Their experiment design operationalizes portfolio theory (Tobin 1958; Markowitz 1959): Subjects can invest in two risky stocks with normally-distributed returns (one has a higher mean and variance than the other). They can also invest in risk-free bonds or borrow, up to a limit, by issuing bonds. Portfolio theory makes a remarkably counterintuitive prediction, called "portfolio separation"; differences in risk tastes determine how much is invested in riskless bonds (or is borrowed) and how much is invested in stocks, but everyone will hold the same proportions of the two stocks in their stock portfolio, regardless of risk attitudes. (For example, an optimal portfolio might place 60 percent of the stock investment in stock 1 and 40 percent in stock 2; those proportions should be the same for all investors, regardless of the overall amount they have invested in stocks instead of bonds.)

Relative to the theoretical prediction, subjects invested too heavily in the high-return, high-variance stock and did not issue enough bonds. (About a quarter of their portfolio choices were stochastically dominated by portfolios in which more funds are borrowed and invested in the low-return stock.) Since different subjects chose different stock mixtures, portfolio separation was badly violated.

Kroll et al. paid some subjects ten times as much money, to see how behavior would change. Highly paid subjects invested more heavily in the low-return, low-variance stock, bringing their portfolios closer to the optimal mixture. (They also searched more for information about stock returns in previous periods, even though they were told returns were independent each period.) However, my guess is that the increased incentive brought their portfolios closer to the optimum simply by increasing risk aversion and reducing investment in the high-variance stock, which coincidentally moved allocations toward the optimum predicted by portfolio theory.

A third study by Kroll and Levy (1992) varied several features of the earlier studies: MBA students participated in weekly sessions over a semester, competed for grade points in a tournament structure, and could see the publicly posted decisions and performance of others. The students' portfolios were much closer to those predicted by portfolio theory than in the two earlier studies; however, they did not reallocate portfolios in response to changes in between-stock intercorrelations as sharply as the theory prescribes (but subjects in the earlier studies did not respond at all). Students also tended to mimic the decisions of high-performing students, and the tournament payment structure appeared to create a borrowing frenzy at the end of the experiment.

£ Subjective Expected Utility

In "subjective expected utility" (SEU), probabilities are not objectively known as they are assumed to be in EU. The events over which people have subjective (or "personal") probabilities are called "states." Decision makers choose acts X which yield consequences $x(r)$ that depend on which of several states ($s \in S$) occurs. The SEU axioms show when preferences can be represented by subjective expected utility, with utilities $u(x(r))$ and (subjective) state probabilities $p(s)$ both derived from preferences, as follows:

$$(IO) \quad SEU(X) = \sum_s p(s)u(x(s))$$

SEU was inspired by Ramsey (1931) but made dear by Savage (1954), who combined the von Neumann and Morgenstern (1944) EC approach with de Finetti's (1937) calculus of subjective probabilities. Since probabilities in SEU are derived from preferences, rather than assumed (as in EU), SEU applies more widely. See Fishburn (1988, 1989) and Kami and Schmeidler (1990) for technical review, or Camerer and Weber (1992).

Most of the empirical evidence specifically critical of SEU concerns precisely the distinction between whether probability is known or unknown. This basic distinction goes by many names: risk versus uncertainty (Knight 1921); unambiguous versus ambiguous probability (Ellsberg 1961); precise or sharp versus vague or fuzzy probabilities. In SEU the distinction between known and unknown probability is pointless because subjective probabilities are never unknown—they are always known to decision makers (and inferable from choices). But empirical evidence suggests people do make such a distinction.

1. The Ellsberg Paradox

The first serious challenge to SEU was posed by the paradoxes of Ellsberg (1961). Two similar problems were posed in his remarkable paper. Here is one of them:

A decision maker chooses from an urn that contains thirty red balls and sixty balls in some combination of black and yellow. There are two pairs of outcomes, X and Y , and Z and W . If a red ball is drawn, X is won; if a black ball is drawn, Y is won; if a yellow ball is drawn, Z is won; if a yellow ball is drawn, W is won. The outcomes are defined in Table 8.2. For example, the act X pays W if a red ball is drawn; Y pays W if a black ball is drawn.

Many people choose $X > Y$ and $Z > W$. The number of black balls that yield a win if act Y is chosen is unknown (or ambiguous); people prefer the less ambiguous act X . The same principle, applied to the second choice, favors Y because exactly sixty balls yield W (The number of black balls is true for losses, $W < 0$.) In this example, people prefer acts with a known probability of winning. That is, they take confidence in estimates of subjective probability into account when making choices.

Such a pattern is inconsistent with the sure-thing principle of SEU. Suppose $p(r)$, $p(b)$, and $p(y)$ are the subjective probabilities of drawing a red, black, or yellow ball. Under SEU, $X > Y$ if and only if $p(r)u(W) > p(b)u(W)$ (selling $u(0) = 0$), or $p(r) > p(b)$. Similarly, $Z > W$ implies $p(b)u(y) > p(y)u(y)$, or $p(b) > p(y)$. If probabilities are additive, then $p(b)u(y) = p(b) + p(y)$ (since $p(b)u(y) = 0$). Then $Y > X$ implies $p(b) > p(r)$, which conflicts with the earlier inequality.

Table 8.2. The Ellsberg Paradox

Act	30 balls	60 balls	
		Black	Yellow
X	W	0	0
Y	0	W	0
Z	W	0	w
W	0	W	w

Source: Ellsberg 1961.

There are several reactions to the paradox. One is to deny it (e.g., de Finetti 1977; Howard 1992)—whether there are thirty balls, or some number between zero and sixty, shouldn't matter—but this does not explain the evidence. Another reaction is reductionist (Marschak 1975): even if probability is not sharply known, there may be a sharp "second order" distribution of probability (SOP), or probability of various probabilities, that restores the usefulness of SEU. For example, in the Ellsberg case a person might not know the number of black balls, but might think that each possible number of balls from zero to sixty is equally likely (i.e., $p(k \text{ black balls}) = 1/61$ for $0 \leq k \leq 60$). But there is no guarantee that a sharp second-order probability exists, or that it captures subjects' intuitions about ambiguity. A more constructive reaction is to study the paradox empirically.

2. Conceptions of Ambiguity

Defining ambiguity is a popular pastime in decision theory. Ellsberg's (1961, 657) definition is typical, if messy: ambiguity is the "qualitative (depending on the amount, type, reliability, and 'unanimity' of information), giving rise to one's degree of 'confidence' in an estimate of relative likelihoods." I favor a slightly pithier definition: ambiguity is knowledge of missing information, or not knowing relevant information that could be known (Frisch and Baron 1988; cf. Heath and Tversky 1991).

The missing information definition includes other kinds of ambiguity as special cases. The composition of the ambiguous Ellsberg urn is missing information, which is relevant and could be known but is not. Doubts about the credibility of sources and disagreements among experts create missing information (namely, whether a source or expert can be believed). Keynes (1921) proposed that the weight of evidence be taken into account along with its implications. For instance, in Scottish law there are three verdicts: guilty, innocent, and unproven. While evidence might imply guilt, if there is too little evidence it has low weight and the verdict will be "unproven." If the weight of evidence is defined as the fraction of available information, then missing information lowers evidential weight,

3. Empirical Tests

Ellsberg did not run any formal experiments,⁵¹ but his thought experiments frequently replicated and extended. Becker and Brownson (1964) did the first careful study. Their subjects chose between urns containing 100 red and black balls. Drawing a red ball⁴ paid \$1. The number of red balls fell within a different range for each urn. For example, one urn had exactly fifty red balls; another had between fifteen and eighty-five balls. The urns with unknown contents were covered; they did not report how many balls were actually used in each. Subjects chose between pairs of urns, differing in the range of red balls, and said how much they would pay to draw from their preferred urn.

Subjects always picked the less ambiguous urn and paid high amounts to avoid ambiguity (which increased with the degree of ambiguity). For example, they paid an average of \$.36 to choose from an urn with 50 red balls instead of an ambiguous urn with 100 red balls. (The expected value of a draw was \$.50!)

There were several other studies in the 1970s and 1980s, all using similar paradigms. In most of the studies a choice was picked randomly and played for money. Some stylized facts emerged from these studies.

Ambiguity aversion is found consistently in variants of the Ellsberg problems (many of them using small actual payoffs). Ambiguity averters are typically immune to written arguments against their paradoxical choices (e.g., Stovic and Tversky 1974), and pay substantial premiums to avoid ambiguity—around 10 to 20 percent of expected value or probability (MacCrimmon and Larsson 1979; Curley and Yates 1989; Bernasconi and Loomes, 1992). Risk attitudes and ambiguity attitudes are uncorrelated (Cohen, Jaffray, and Said 1985; Hogarth and Einhorn 1990). Subjects would rather bet on known probabilities p than on known probability distributions of probability (compound lotteries) with a mean of p (Yates and Zukowski 1979; Ullson 1980). Increasing the range of possible probabilities increases ambiguity aversion (Curley and Yates 1985). There is some evidence of ambiguity preference for betting on gains with low ambiguous probability, or betting on losses with high probability (e.g., Kahn and Sarin 1988). (This may be due to perceived skewness in ambiguous distributions of low and high probability, which makes their means higher and lower than the actuals.)

Curley, Yates, and Abrams (1986) tested several psychological explanations for ambiguity aversion. Subjectively said the urn could not be biased against them were ambiguity averse too, suggesting ambiguity aversion is not due to an expressed belief in "hostile" generation of outcomes. Subjects were ambiguity averse when indifference was allowed (cf. Roberts 1963). Subjects were more ambiguity averse when they knew the contents of the ambiguous urn would be revealed afterward to others.

Competence

Ambiguity aversion implies there may be a gap between subjects' beliefs about an event's likelihood and their willingness to bet on the event. (In SET) there can be no such gap, since beliefs are derived from betting preferences.) Heath and

Tversky (1991) suggest that competence-knowledge, skill, comprehension is what causes the gap. They ran one set of experiments in which subjects gave probability assessments for natural events; they knew a little or a lot about them. (They were rewarded with a scoring rule.) The subjects were then asked whether they would like to bet on the event, or on a chance device constructed to have the same probability as the event. If people ate ambiguity aversion they should always prefer to bet on chance devices since events are inherently ambiguous. But subjects who knew a lot about a domain of events preferred betting on events: those who knew little preferred betting on chance. People preferred betting on events if they knew a lot about holding beliefs consistently.

The competence hypothesis broadens the study of choice anomalies in SEU by suggesting that ambiguity about probability is just one of many forces that makes people reluctant to bet, by undermining competence. Heath and Tversky suggest that competence influences betting because personal and social assignments of credit and blame are asymmetric; competent people can take credit for winning in a way that incompetent people cannot, and incompetent people may suffer more blame.

Ambiguity in Markets

Camerer and Kunreuther (1988) studied ambiguity in an experimental market for hazards that incurred a loss; subjects could pay "insurance" to get rid of them. A second-order probability distribution was used to induce ambiguity in loss probability (i.e., the probability had three possible values, which were equally likely). Ambiguity had little effect on market prices or volume, but it did increase the variance in the distribution of sales across sellers (i.e., some sellers sold more insurance policies and others sold none).

Sarin and Weber (1993) studied ambiguity in market settings. They auctioned off ambiguous and unambiguous lotteries in double-oracle and sealed bid auctions with German subjects. An unambiguous lottery was a draw of a ball from an open urn with live winning and five losing balls (a winning draw paid 10 marks). An ambiguous lottery was a draw from a hidden urn with an unknown composition of balls. They found persistent ambiguity aversion around $p = .5$ (but not around $p = .05$), even under stationary replication, and even when the ambiguous and unambiguous lottery markets operated simultaneously.

While these two studies explicitly tested the effect of ambiguity in an experimental market, other studies may provide indirect evidence of ambiguity effects. In most market experiments subjects are not fully informed about the market's structure (they know only their own valuations), so they face some ambiguity due to missing information. In games they usually know the entire structure, but their opponent's rationality is ambiguous. Anomalous behavior in markets and games might therefore be explained by ambiguity aversion (Camerer and Karjane, 1994). For example, subjects in market experiments often behave very conservatively early in the experiment (e.g., failing to trade when it is optimal); their behavior is usually labelled "risk aversion" or "confusion." The apparent confusion may be a manifestation of ambiguity aversion: subjects would rather do

nothing, foregoing profits, than take action in an ambiguous environment where information is missing and they feel incompetent. As the experiment progresses, subjects learn and the amount of missing information shrinks, reducing their ambiguity aversion and conservative behavior. Thus, notions like competence, missing information, and ambiguity aversion might help us make sense of disequilibrium behavior early in experiments, which we currently call "confusion" and largely ignore.

4. Formal Models

Several kinds of formal models have been proposed to accommodate ambiguity effects. Some of the models are axiomatic generalizations of SEU. Others invoke psychological principles or propose ad hoc decision rules. I mention just a few (see Camerer and Weber 1992, for details).

There are several ways of modifying SEU without abandoning the expectations principle. Utilities can depend directly on ambiguity (Smith 1969; Sarin and Winkler 1989). Expected probabilities can be undetweighted (Fellner 1961; Einhorn and Hogarth 1985), or possible probabilities in a second order distribution can be weighted nonlinearly before taking their expectation (Segal 1987a; Kahn and Sarin 1988; Becker and Sarin 1990). The nonlinear weights might depend on outcomes (Hazen 1987).

Another way to modify SEU is to replace the expectation principle with more general decision rules that combine expected SEU with the minimum SEU over a set of possible probabilities (this was Ellsberg's proposal; cf. Hodges and Lehmann 1952, and Giiridenfors and Sahlin 1982). A special case of the combined model is maximizing the minimum SEU over some set of possible probabilities (Gilboa and Schmeidler 1989), perhaps weighting possible probabilities by a person's willingness to bet on them (Nau 1989). (Allowing preferences over ambiguous acts to be incomplete can yield a constrained maximin representation too; Bewley 1986.)

More radically, subjective probabilities may be precise but nonadditive: that is, $p(A \cup B) \neq p(A) + p(B) - p(A \cap B)$ (Luce and Narens 1985; Gilboa 1987; Schmeidler 1989; Wakker 1989; Tversky and Kahneman 1992). The idea is that unwillingness to bet on ambiguous states can be expressed by attaching lower subjective probability to those states. Nonadditive SEU must be calculated with a special summation or integral, first discovered by Choquet (1955). If states are ranked by the utilities of their consequences $f(s_i)$ for a particular act f , from $u(f(s_1)) > \dots > u(f(s_n))$, then the finite Choquet integral is

$$(11) \quad u(f(s_1))p(s_1) + \sum_{i=2}^n u(f(s_i)) [p(\bigcup_{j=1}^i s_j) - p(\bigcup_{j=1}^{i-1} s_j)]$$

Note that if $p(\cdot)$ is additive and the states s_i are mutually exclusive, then the term in brackets reduces to $p(s_i)$ and equation (11) reduces to SEU. There is obviously a close kinship between nonadditive probability in SEU and rank-dependent probability weights in EU (compare equations (11) and (5)).

5. Applications to Economics

Theoretical and applied work on variants of SEU is several years behind the work on EU reviewed in section 11LD. For example, there are virtually no empirical tests pitting the formal models against each other.

There are several applications to economics. Dow and Werlang (1992) apply SEU with nonadditive probability to financial markets. They show that in theory increased ambiguity creates wider bid-ask spreads. Blank (1991) reports a large *American Economic Review* experiment comparing single-blind refereeing, when a paper's author is known to the referee, with double-blind refereeing, when the author is unknown. (Not knowing the author of a paper, and knowing that you could, creates ambiguity.) Incoming papers were randomly allocated to the two reviewing conditions. Single-blind papers are rated more highly by referees (3.47 vs. 3.37, on a five-point scale), and accepted more frequently (14.1 percent vs. 10.6 percent) than double blind papers; one could ascribe the difference to ambiguity aversion. French and Poterba (1991) document a global preference for home-country investments, which costs investors the equivalent of about 3 percent in annual returns in foregone diversification. The preference for home-country investment and the revealed preference for publishing known-author papers are consistent with the Heath and Tversky (1991) finding that people prefer to bet on events they know more about (holding likelihood constant).

In surveys using hypothetical vignettes (based on naturally occurring risks), Hogarth and Kunreuther (1985, 1989) found that pricing decisions by professional actuaries and insurance underwriters reflect ambiguity aversion. (Indeed, ambiguity premiums are *required* by many of the pricing rules used in insurance companies; see Hogarth and Kunreuther 1990.) Knight (1921) and Bewley (1986) suggest entrepreneurship can be understood as ambiguity neutrality. Bewley also sketched some applications of his theory to labor contracting.

The idea that missing information creates uncomfortable ambiguity implies that agents will demand information simply to reduce ambiguity (and discomfort), even if the information does not help them make better decisions. (Demand for information could then be modeled as a primitive, instead of deriving it from preferences over outcomes that result from informed decisions.) Demand for ambiguity-reducing information may help explain potential anomalies like uninformative advertising and alleged medical overtesting.

F. Choice over time

In traditional models of choice over time, simple axioms imply that preferences over consumption streams $X = (x(1), x(2), \dots, x(n), \dots)$ can be represented by an additive discounted utility representation (e.g., Fishburn and Rubinstein 1982):

$$(12) \quad X \succ Y \iff \int d(t)x(t) > \int d(t)y(t)$$

where $d(t)$ represents the present value of consumption at time t . Further axioms imply the familiar exponential forms $d(t) = \delta^t$ and $d(t) = e^{-\delta t}$.

Considering the importance of intertemporal choice in economic theory, there have been relatively few experimental tests of choices over time. The tests conducted so far indicate several systematic violations that are remarkably similar to violations of EU. (1) Implicit discount rates, derived from choices over income or consumption streams, decline with time horizon: people are much more impatient about immediate delays than about future delays of the same length. Put formally, discount rates seem to be hyperbolic, $d(t) = (1 + \alpha t)^{-\beta}$, rather than exponential ($d(t) = e^{-\alpha t}$). (The exponential form is the limit of the hyperbolic form as α goes to zero.) (2) Discount rates are larger for smaller amounts of income or consumption, and are larger for gains than for losses of equal magnitude. (3) People demand more to delay consumption than they will pay to speed it up. Loewenstein and Prelec (1992) give a generalization of discounted utility, using a value function akin to the one in prospect theory, which can explain all these anomalies.

Some of these findings, especially hyperbolic discounting, have been replicated by Thaler (1981), Loewenstein (1988), Benzion, Rapoport, and Yagil (1989), and Shelley (in press). Hyperbolic discount rates are also commonly observed in experiments with animals (e.g., Ainslie 1975).

Human subjects were not financially rewarded in many of these experiments. Indeed, economic experiments on choices over time are extremely difficult to run because subjects cannot be paid immediately. Even if experimenters intend to pay, subjects may not trust them. Making subjects return to collect their money also imposes a cost that must be accounted for. A few experimenters have used monetary payments and observed the same anomalies observed in hypothetical choices (Loewenstein 1988; Holcomb and Nelson 1989; Horowitz 1991; Carlson and Johnson 1992).

Time affects choices in some other interesting ways (see Loewenstein and Elster 1992). Temptation arises if immediate discount rates are too high, creating a problem of self-control (examples include smoking, overeating, procrastination, and using credit cards; see Ausubel 1991). Many people have modeled self-control as a neo-Freudian conflict between multiple selves: a myopic self greedily consumes immediately while a foresightful self tries to restrict the myopic self's binges (Thaler and Shefrin 1981; Schelling 1984; Laibson 1994). The neo-Freudian view helps explain the existence of institutions that make self-control easier—Christmas clubs, "forced saving" by overwithholding income taxes from wages, voluntary diet plans that ration food, and so on.

Ex ante anticipation of consumption, or cherishing of ex post memories of consumption, may affect timing of choices too. Loewenstein (1987) asked college students when they would prefer to kiss their favorite movie star. (Kisses were hypothetical.) The median subject preferred to wait three days, exhibiting a negative discount rate, to savor the consumption before it occurred. Savoring and its negative counterpart, dread, are probably important parts of the explanation for why people buy insurance and gamble.

Varey and Kahneman (1992) explore judgments of the overall quality of temporally extended episodes (such as watching a movie or taking a week-long vacation). A natural model is that people judge overall utility by cumulating or integrating momentary utilities experienced across the episode. This model appears

to be wrong in interesting ways. For example, the cumulation model predicts that episode length will be correlated with overall quality, but in experiments using short and long films, length is only a weak correlate. People also appear to dramatically overweigh the momentary utilities at the peak and end of the episode (Fredrickson and Kahneman in press). It is easy to generate dominance violations as a result of this overweighting: in one study, subjects preferred long, mildly painful colonoscopies that ended with a gradual easing of pain to shorter ones that had less overall discomfort but ended painfully (Redelmeier and Kahneman 1993).

G. Process Theories and Tests

There is a large body of work, mostly by psychologists, spelling out procedures people use to make choices. I will mention only a few germane contributions.

When the number of alternatives is large, people often use conjunctive and disjunctive rules to reject alternatives, especially in complicated settings with many alternatives having several dimensions or attributes (e.g., Einhorn 1970). These rules specify cutoff levels for all attributes. With a conjunctive rule, if an alternative falls below any cutoff it is rejected. In a disjunctive rule, an alternative that satisfies any cutoff is not rejected. Tversky (1972) proposed and tested a related rule called "elimination by aspects" (EBA): subjects pick one attribute with a probability that depends on its importance, eliminate alternatives without the attribute (or below a cutoff level), then choose another attribute and repeat the procedure.

Grether and Wilde (1984) studied conjunctive rules experimentally. Subjects were forced to use conjunctive rules to make choices. "Inspecting" an attribute had a cost; subjects could choose cutoffs and the order of inspection. Subjects tended to inspect the lowest-inspection-cost attributes. They often violated normative predictions of optimal cutoffs and inspection order (Wilde 1982) because they ignored the influence that a drop in cost of inspecting one attribute should have on inspection order and the level of other attributes' cutoffs. (Roughly speaking, subjects used a partial-equilibrium mental model and ignored important general equilibrium effects.)

Using simulations, Johnson and Payne (1985) showed that in some choice settings simple decision rules, like conjunctive rules, are almost as accurate as more complex rules, like EV-maximization, and take much less effort. (Estimates of effort come from psychological studies of how long it takes to perform various mental operations such as adding, remembering, multiplying.) Johnson and Payne's simulations give an underpinning to the "labor theory of cognition" proposed by Smith and Walker (1993) (cf. Marschak 1968; Wilcox, in press). In the labor theory of cognition, mental effort is like manual labor: People dislike thinking hard; more effort is supplied in response to higher incentives (except when constrained at computational limits); and greater effort reduces the variance of responses around an optimum.

The labor theory is a natural way for economists to comprehend thinking, but in some ways it represents a psychological step backward. Some history is needed

to explain my point. From about 1930 to 1960, psychology was dominated by "behaviorism," the study of the response of animals (including people) to stimulus. Behaviorism ignored the details of cognitive process and treated the brain as a black box. In the 1960s a better way to model cognition came around—the brain is like a computer—and the information processing paradigm was born. Behaviorism was largely abandoned (except for some domains of animal learning) because the computer metaphor was so appealing and because of mounting evidence, such as transfer of learning across tasks, which was anomalous for behaviorism and cried out for a cognitive explanation.

The labor theory of cognition is a partial return to behaviorism because it concentrates on the relationship of stimulus (incentives) to response (choices), compressing the details of cognitive processing into the catch-all category, "effort." The empirical question is whether research that incorporates more detail of thinking processes, like that discussed in section II and in parts of this section, generates better predictions than the labor theory. I think it does in many cases. However, the simpler labor theory may still be useful in some economic applications (especially in formal theorizing) because of its parsimony and rough accuracy.

H. Description invariance

Utility theories makes several invisible background assumptions that are usually considered too innocuous to spell out. Two crucial assumptions are "description invariance" and "procedure invariance"; different representations of the same choice problem, and different elicitation procedures, should yield the same preference (Tversky and Kahneman 1986). Money illusion is an example: doubling all wages and prices, or denominating them in Irish pounds rather than dollars, shouldn't make anyone feel richer.

Both principles are sometimes violated. Invariance violations are especially troublesome for utility theories (including generalizations of EU and SEU) and provide the strongest indications that preferences are constructed from procedural rules.

Luce and von Winterfeldt (1994) offer a similar decomposition of axioms. They distinguish between axioms of "structural rationality" (or "accounting equivalences"), which prescribe indifference between formally equivalent descriptions of a gamble, axioms of "preference rationality" (like independence), and axioms of "quasi-rationality," which prescribe how consequences are coded as gains and losses.

I. Framing Effects

The most famous violations of description invariance are "framing effects," reversals of preference induced by changes in reference points. For example, McNeil, Pauker, and Tversky (1988) gave some American doctors and Israeli medical students data on survival rates after treatment for lung cancer by radiation therapy and surgery, as shown in Table 8.3. (Survival rates are the percentage of

Table 8.3. Survival and Modality Framing of Lung Treatments

	Survival Frame (% alive)		Mortality Frame (% dead)		Both Frames Presented	
	Radiation	Surgery	Radiation	Surgery	Radiation	Surgery
After treatment	100	90	0	10		
After one year	77	68	23	32		
After five years	22	34	78	66		
Percentage choosing each:						
American doctors and medical students	16 (117)	84	50 (80)	50	44 (1213)	50
Israeli medical and science student,	20 (126)	80	40 (132)	56	34 (144)	66

Source: McNeil et al. 1988.

Note: Numbers in parentheses are *n*.

patients surviving a given length of time.) Others were given the same data, phrased as mortality rates (percentage of patients who died before a given length of time). A third group got both frames. In the survival frame, 16 to 20 percent favored radiation therapy, since surgery only reduces immediate survival from 100 to 90 percent and keeps more patients alive in the long-run. But in the mortality frame the 10 percent death rate after surgery looms large; then nearly half favored radiation. Given both frames, 40 percent favored radiation (suggesting the mortality frame is more natural or potent).

The most pressing question is whether framing effects are systematic and predictable. The evidence is mixed. Fischhoff (1983) used a simple pairwise choice that could be framed several ways. Subjects' ratings of which frames seemed most natural were unrelated with the frames they appeared to use in making choices. Van Schie and Van Ommen (1990) found similar results: expected frame preferences were only weakly correlated with risk aversion in choices (replicating Fischhoff's discouraging result), but the initial description of the choice did affect risk aversion. (If the problem was initially posed in terms of losses, rather than neutrally, frames emphasizing losses were preferred and choices were more risk seeking.)

Gertner (1993) reports an interesting framing finding using data from actual bets (averaging \$3,200) on a television game show. He reports that when the cash stake available for betting increased by \$1, bets increased by about \$.60 (cf. the "house money" effect in Thaler and Johnson [1990]). But when the amount of cash winnings that couldn't be bet increased (or when a contestant had won a car) by \$1, bets increased by only a penny. The data are clearly inconsistent with the theory that contestants integrate assets (bettable cash, unbettable cash, and car) then bet based on their integrated assets.

Thaler (1985) proposed a "hedonic editing" rule for choosing reference points. For example, people should segregate two \$50 gains (resetting their reference point to absorb the first gain before valuing the second one) because $v(\$0) + v(\$0) > v(\$100)$ if $v(\cdot)$ is concave for gains. But Thaler and Johnson (1990) found that hedonic editing did not explain choice in two-age senin where prior gain or loss was followed by a choice. People use a variety of rules instead.

Consider a prior loss of -\$7.50 followed by a second-stage choice between \$0 and (.5, \$2.25; .5, -\$2.25). If subjects integrate the loss and the choice (and obey prospect theory), they choose between -\$7.50 and (.5, -\$5.25; .5, -\$9.75), and take the gamble, because they are risk seeking over losses. But if subjects segregate the loss they will take \$0 over the gamble (which is unappealing because of loss aversion). Sixty percent of subjects rejected the gamble, which suggests they are segregating the prior loss.

Now consider a prior loss of -\$7.50 followed by a choice between \$2.50 and (.33, \$7.50; .67, \$0). Under integration, the choice is between -\$5.00 and (.33, 0; .67, -\$7.50); subjects should gamble. Under segregation, subjects should take the sure \$2.50. In fact, 71 percent of the subjects preferred the gamble, which suggests they are integrating the prior loss. Integrating the loss is appealing because the integrated gamble gives a .33 chance of breaking even by winning \$7.50 after recouping the prior loss.

In the firm situation, the prior loss appears to be segregated from the gamble choice. In the second situation the loss appears to be integrated (when breaking even is possible). Framing principles appear to be a long list of rules like these dictating whether people integrate or segregate depending on contextual details (cf. Luce and von Winterfeldt's, (1994) discussion of "joint receipt").

2. Lottery GifTelation, Regret, and Display Effects

The correlation between outcomes of two lotteries is an element of choice description which, according to EU, should not affect preferences. Consider the top two choices in Table 8.4 (see Loomes 1988a). The probability distributions of A and B outcomes are the same in both choice 1-A is (.3, 10) and B is (.75, 5)—but their correlation is different. In choice 1, the payoffs are negatively correlated and in choice 2 they are positively correlated. Under any utility theory that assigns a number to outcome distributions, $u(A)$ and $u(B)$ should be the same in both choices. But regret theories assign a number to $u(A; B)$, so the correlation between A and B outcomes can matter.

Loomes and Sugden (1987a) suggest a "convexity" (or "regret aversion") hypothesis to make the theory testable. Convexity implies that the comparative utility from getting 0 and foregoing 10 is worse than the sum of comparative utilities from getting 0 instead of 5, and getting 5 instead of 10 (That is $u(0, 10) < u(0, 5) + u(5, 10)$.) Then people might switch preference from $B \succ A$ in choice 1 to $A \succ B$ in choice 2 (but they should not switch the opposite way).

In experiments with choice displayed graphically such as the top two choices in Table 8.4 (when one gamble actually played), about a third of subjects switch

Table 8.4. Choices in Regffl E

		Probability of State	
Choice 1	A1	.3	.7
	B1	10	0
Choice 2	A2	.3	A
	B2	10	0
Choice 3 (collapsed table)	A3	.3	.7
	B3	10	0

preferences between the two displays, 80 percent of them in the direction predicted by convexity (Loomes 1988a, 1988b, 1989a; Loomes and Sugden 1987b; Starmer and Sugden 1989c). Regret-aversion also predicts that if payoffs are juxtaposed properly people will choose a MOChastically dominated lottery (e.g., (.38, \$3; .2, \$10) will be preferred to (.42, \$3; .2, \$10)). About half of subjects did violate dominance in this way in studies by Tversky and Kahneman (1986) and Loomes, Starmer, and Sugden (1992).

Battalio, Kagel, and Jiranyakul (1990) found much less switching when the A and B payoffs were shown horizontally, rather than vertically as in Table 8.4. In tests with several different displays, Harless (1992a) only found systematic regret effects-switching in the direction predicted by regret aversion-in the vertically aligned Table 8.4 display. He also experimented with displays in which states yielding the same payoff were collapsed for each act, as shown in choice 3 in Table 8.4. Regret effects presumably arise from comparing two acts' outcomes in the same state (i.e., comparing two row entries within a column). The choice 3 display might weaken regret effects by making it more difficult to compare outcomes within a column. Indeed, Harless found no regret effects when choice 3 displays were used.

Starmer and Sugden (1993b) found no regret effects using choice 3 displays ("strip displays") either. They discovered something else even more subtle, and astonishing: Compare choice B2, (.3, 5; A, 5; .3, 0), with choice B3, (.7, 5; .3, 0). Both describe the same lottery (a 70 percent chance of winning 5), but the B2 choice has two states that yield the prize of 5, one with .3 probability and one with .4 probability, while choice B3 has a single winning state, with probability, .7. Yet they found that about 10 percent more subjects preferred the two state choice B2 over a third gamble than preferred the one state choice over the third gamble.

Simply splitting one state into two states with the same outcome, and the same total probability, increased preference for an act noticeably. (A similar preference-for-splitting can occur in theories with nonlinear weights, since $w(.1) + w(.1)$ can be greater than $w(.2)$. Indeed, Tversky and Kahneman [1986] exploited this property to construct violations of stochastic dominance in "opaque" choices.)

The regret studies show the interplay of experimental studies, and the cumulation of discoveries, at its best. Several studies by Loomes, Starmer, and Sugden established an important anomaly, the effect of lottery correlation on choices, which violated description invariance but could be neatly explained by regret aversion. Then a small test by Battalio et al. hinted that regret aversion could be sensitive to display effects. Harless landed a second, harder blow. Then Stanner and Sugden confirmed that regret aversion largely disappeared when displays were altered. In so doing, they (re)discovered the remarkable fact that splitting states made an act substantially more attractive. This is a story of successful detective work. The effects are small in magnitude. The studies used large samples of subjects (40–200 or so) for statistical power and a wide variety of gambles for robustness. Less methodical investigators might have falsely accepted the null hypothesis that there were no regret effects in the first place or underestimated their sensitivity to display.

Earlier studies tested the effect of using different lottery displays—trees, matrices, verbal descriptions—on the rate of Allais paradox choices (Moskowitz 1974; Keller 1985a). Display effects were generally too small to eliminate violations, but large enough that more work on display effects could prove fruitful.

3. Compound Lottery Reduction

Most utility theories contain an explicit or implicit reduction axiom stating that whether a lottery is described as a compound gamble with several probabilistic stages or as a single-stage gamble should not affect preference.

The reduction axiom has great normative appeal but is often violated in direct comparisons between one-stage choices and multiple-stage choices (Bar-Hillel 1973; Ronen 1973; Kahneman and Tversky 1979; Keller 1985b; Bernasconi and Loomes, 1992; Conlisk 1989). Most violations seem to be caused by the tendency to "isolate" the uncommon elements of two gambles, by cancelling common first-stage probabilities or common final-stage payoffs. For example (Kahneman and Tversky 1979), gamble 1 is a two-stage gamble in which there is a .75 chance of winning nothing at the first stage and a .25 chance of moving to a second stage gamble that pays a certain 3000. Denote gamble 1 by $(.75, 0; .25[1, 3000])$. Gamble 2 is a .75 chance of winning nothing at the first stage and a .25 chance of moving to a second-stage gamble that pays 4000 with probability .8, or nothing with probability .2. Gamble 2 is denoted $(.75, 0; .25[.8, 4000; .2, 0])$. In reduced form, these compound lotteries are equivalent to $(.25, 3000)$ and $(.2, 4000)$; most people pick the riskier gamble 2, $(.2, 4000)$. The isolation effect refers to the tendency to ignore the common $(.75, 0)$ stage and choose gamble 1 or gamble 2 by isolating

their uncommon second-stage elements. Then people usually choose gamble 1 because they prefer $(1, 3000)$ to $(.8, 4000; .2, 0)$.

This cloud of violation has a silver lining, however. Since isolation is simply an application of the independence axiom, the independence axiom (and the betweenness axiom) is less often violated when gambles are presented in compound form (see Conlisk 1989; Segal 1990; Luce 1990; Brothers 1990; Camerer and Ho 1994; Bernasconi 1992; cf. von Winterfeldt, Chung, Luce, and Cho 1992). In the example above, if people prefer $(1, 3000)$ to $(.8, 4000; .2, 0)$, then they will prefer the compound gamble $(.75, 0; .25[1, 3000])$ to $(.75, 0; .25[.8, 4000; .2, 0])$.

Reduction is important for some procedures that are widely used in experimental economics. An example is the "random lottery" procedure, in which subjects make many choices but only one choice is picked, at random, and played out for money. If subjects violate independence and the random lottery procedure is used, then if they obey reduction they will *always* choose as if they made only one choice. (Whether they will do so depends on the set of choices and the nature of their independence violations.) But suppose subjects isolate each pair while they are choosing, rather than multiplying its probabilities by the chance of that pair being randomly picked.⁵⁷ Then they will choose as if they were making only one choice, even if they violate independence. In fact, people *do* seem to isolate each pair (Stanner and Sugden 1991a; cf. Camerer 1989a).

Reduction plays a similar role in the Becker, DeGroot and Marschak (1964) procedure for revealing preferences and in the Vickrey auction for gambles (see chapter 1). If subjects obey reduction and violate independence, these mechanisms will not be preference-revealing for gambles (Chew 1985; Kami and Safra 1987). But if subjects violate reduction by isolating, both techniques are preference-revealing. Ironically, the violation of reduction that is implied by the isolation effect works to the experimenter's advantage, *ensuring* the usefulness of the random-lottery procedure and the BDM procedure.

Reduction is also assumed in attempts to induce risk neutrality by paying subjects in lottery tickets (Smith 1961; Roth and Malouf 1979; Berg et al. 1986). Evidence on the effect of the ticket procedure is mixed (see chapter 1 and chapter 7). In the ticket mechanism, reduction violations work against the experimenter.

1. Procedure Invariance

Elicited preferences should be invariant to the procedures used to elicit them, but they seem not to be. The elicitation biases mentioned above in section III.C.4, such as differences in utility functions derived by probability- and certainty-equivalence, violate procedure invariance.

Shafir (1991) discovered a violation of procedure invariance that is both illustrative and startling. Consider two lotteries, $(.5, \$50)$ and $(.8, \$150; .2, -\$10)$. Asked which lottery they would choose if they had neither, 75 percent of subjects ($n = 279$) picked the riskier one, $(.8, \$150; .2, -\$10)$. Asked which they would give up if they had both, 50 percent preferred to give up $(.8, \$150; .2, -\$10)$. The pattern is surprising because complete preferences, and invariance of expressed

preference to whether the choice is an acceptance or a rejection, imply that the two percentages should add to 100 percent. Shafir's explanation is that the second gamble is "enriched" there is a good reason to pick it (it has a higher "winning payoff") and a good reason to reject it (it has a chance of losing). If people choose gambles with more positive features, and reject gambles with more negative features, then enriched gambles with both kinds of features will both be chosen and rejected more often, a dramatic violation of the principle that the procedure for eliciting choice should not affect choices. (Notice that this violation, while not substantially rational, seems to have a clear procedural explanation that is amenable to formal modeling.)

Another violation of procedure invariance with a long history is preference reversal (e.g., Tversky and Thaler 1990; Loomes forthcoming). The reversal literature began when Slovic and Lichtenstein (1968) noticed that the prices subjects gave for bets were highly correlated with bet payoffs, but choices were more highly correlated with probabilities. They conjectured that if subjects were offered two bets, one with a high probability and low payoff (a "P-bet") and another with a low probability and high payoff (a "\$-bet"), they might choose the high-probability P-bet but price the high-payoff \$-bet higher.

They were right. Lichtenstein and Slovic (1971, 1973) observed systematic, widespread preference reversals (see also Lindman 1971). The reversal attracted relatively little attention in psychology; perhaps there were plenty of other demonstrations that how a question is asked influences its answer (cf. opinion polls), and psychologists were busy discovering other anomalies. Then Grether and Plott (1979) replicated the earlier findings, using the Becker, DeGroot, and Marschak (BDM) (1964) procedure to elicit utility-compatible selling prices. Their replication attracted much attention within economics. The early debate is described in chapter 1.

1. New Evidence of Preference Reversal

Recent evidence has established some new facts. Reversals disappear when choosing over portfolios of gambles played repeatedly (Wedell and Bodrenbohm forthcoming); this is unsurprising because repeated play reduces the difference in probabilities and payoffs of the two bets that generate reversals. Irwin et al. (1993) elicited large reversals in choosing vs. pricing consumer goods and environmental values (such as cleaner air). In gambles over real losses from a cash endowment, subjects exhibited reversals in the opposite direction when they chose \$-bets they usually priced P-bet higher (less negatively) (McDonald, Huth, and Taube 1991). Casey (1991) also found oppositereversals using buying prices and high stakes gambles (replicated with real payoffs in Casey [1994]).

Most other recent evidence addresses three general explanations for reversals (laid out by Tversky, Slovic, and Kahneman 1990). Reversals are either (1) an artifact of the method used to elicit bet prices; (2) violations of transitivity; or (3) violations of procedure invariance.

The artifact explanation (1) has attracted by far the most attention from economists.

The idea is that the Becker, DeGroot, and Marschak (BDM) (1964) procedure does not elicit truthful telling if the independence axiom is violated and reduction is obeyed (see Holt 1986; Kami and Sarra 1987; Segal 1988). Apparent reversals might then be due to systematic misreports of true selling prices.

Cox and Epstein (1989) avoided problems with the BDM procedure by using a different method. They asked subjects to value a P-bet and \$-bet concurrently, then compared the ranking of the valuations with subjects' pairwise choices. Subjects were somewhat motivated to give accurately ranked valuations because the high-ranked gamble in one randomly chosen pair was played for money (but they were not penalized for inaccurate valuations so long as the preferred gamble had a higher valuation). Since the BDM procedure was not used, any reversals in preference could not be due to BDM distortions. They observed fewer asymmetric reversals (P-bet > \$-bet, \$-bet priced higher), but many symmetric reversals. The rate of symmetric reversals is roughly similar to the 15 to 25 percent rate of reversal observed in some studies cited earlier, in section III.D.2, but is substantially lower than would be expected from purely random switching.

There are three theoretical, philosophical, and empirical counterarguments to the artifactual explanation, which blames the BDM procedure for apparent reversals. First, the BDM procedure only fails if independence is violated and reduction is obeyed. If subjects exhibit an isolation effect and violate reduction (as they appear to do; see section III.H.3) then the BDM procedure works properly.

Second, virtually identical patterns of reversals are observed when prices were elicited without the BDM procedure (in earlier studies); why can the BDM procedure be blamed for those reversals? (Are they a coincidence?)

Third, Safra, Segal, and Spivak (1990a, 1990b) showed that the artifactual explanation has two testable implications: (1) reversing preference goes hand in hand with violating independence by fanning out; and (2) a gamble's selling price SP (derived using BDM) and certainty-equivalent CE (derived from introspection by subjects) need not be equal, but SP and CE should lie on the same side of the gamble's expected value. New experiments suggest both implications are false. McDonald, Huth, and Tuohy (1991) discovered that implication (1) is false—subjects who exhibited fanning out were no more likely to reverse preferences than others who did not fan out. Keller, Segal, and Wang (1993) discovered that implication (2) appears to be false—although SP and CE lie on the same side of EV nearly two-thirds of the time, a violation of this same-side property is strikingly asymmetric: 22 percent of subjects show $SP > EV > CE$ while only 9 percent show $CE > EV > SP$. (If SP-CE differences are due to independence violations in the BDM procedure, the two patterns are random errors and should be roughly equal in number.)

Given the strength of these counterarguments—the second one of which has been known for twenty years—it appears that the artifactual explanation may have received too much attention from wretched researchers with better things to do.

Tversky, Slovic, and Kahneman (1990) conducted an experiment to separate the transitivity (2) and procedure invariance (3) of reversals. Con-

side a P-bet (35/36, \$4) and a \$-bet (11/36, \$16). If a subject chose the P-bet, then stated prices of \$3.50 for the P-bet and \$5.00 for the \$-bet. she reversed preference in the usual way. Subjects then chose between the P-bet and a predetermined certain amount (\$3.85, in this case) and between the \$-bet and \$3.50. Choosing \$3.85 instead of the \$-bet (for which the subject stated a price of \$5) would indicate the \$-bet was overpriced in the pricing task compared to its value in choice, indicating that choice-based preference and value-based preference are different (i.e., procedure invariance is violated). If procedure invariance holds but transitivity is violated, then $P\text{-bet} > \$\text{-bet}$ but $P\text{-bet} < \$3.50$ and $\$3.85 < \-bet . For their data, most reversals (66 percent) were due solely to overpricing of \$-bets. Only 10 percent reflected intransitivity. Cox and Grether (1991) closely replicated this result (10.5 percent intransitivity).

Loomes, Starmer, and Sugden (1989, 1991) disputed these results (see also Loomes 1991b). They think the method understates the degree of intransitivity and overstates the importance of mispricing. They used a similar design and found more evidence of intransitivity. In one experiment, they compared the frequency of reversal, measured two ways. First, subjects stated valuations and then chose between certain amounts and gambles (as in Tversky et al.). For example, a subject who said the P-bet was worth \$3.50 would later be asked to choose between the P-bet and \$4.00. Second, a subject's preferences in the choice between the P-bet and \$4.00 were automatically determined ("imputed") from their earlier valuation. In the example, the subject would be forced to choose the \$4.00, since he said the P-bet was worth \$3.50. The second method (imputed-choice) forces valuations and choices to be consistent: the first method (actual-choice) does not. If choice valuation discrepancies generate reversals there should be substantially more reversals with the actual-choice method (because the imputed-choice method does not allow discrepancies). In fact, there were slightly more reversals using actual choices 11/75 vs. 14. In another experiment using only choice Loomes, Starmer, and Sugden (1991) got about 17 percent intransitive cycles. Loomes and Taylor (1992) got 25 percent cycles using gambles over tosses. These data show about twice as many intransitivities (in a pure choice setting) as were reported by Tversky et al. and Cox and Grether. Regardless of the precise "market share" of intransitivity, it seems clear that both intransitivity and procedural variance play a role in explaining reversals. Both phenomena are well worth studying further.

Bostic, Hermsstein, and Luce (1990) and Loomes (1991e) elicited prices using an iterated choice procedure in which subjects made choices between a bet and a series of certain amounts that varied up and down until indifference was reached, establishing a price. Replacing pricing (a judgment task) with iterated choice should reduce the number of reversals if they are caused by procedure variance. Reversals were reduced.

Recent process evidence is informative too. Schkade and Johnson (1989) used a computer display, with probability and payoff information hidden in boxes that opened when the subject moved a cursor into them, to trace what information a subject was looking at (and for how long). Subjects looked at payoffs a larger

fraction of the time when setting prices (55 percent) than when choosing (45 percent).

Johns, Payne, and Bettman (1988) made choices more difficult to process simply by multiplying them (e.g., 9/10 became 513/570). The change made elicitation-based strategies more computationally difficult and doubled the number of reversals.

Much of the new evidence corroborates the original Slovic and Lichtenstein interpretation of the cause of reversals. Their interpretation is now known as "contingent weighting": the weight attached to a dimension increases when the dimension is psychologically "compatible" with the response (payoffs and price-setting, for instance). When a dimension is psychologically compatible, the most prominent dimension is weighted more highly (e.g., probability is weighted highly in choosing). Tversky, Sattath, and Slovic (1988) showed contingent weighting effects in several settings, of which preference reversals are just one example. Contingent weighting can also explain the opposite patterns of reversals for gambles over losses (loss amount is weighted more highly in forming valuations than in choosing, leading people to prefer \$-bets but "lose" more negatively). However, explaining opposite reversals using buyings with high stakes gambles (Casey 1991) requires a more complicated theory with framing features.

2. Arbitrage and Incentives

Several experimental economists have studied the effects of arbitrage and incentives on preference reversals.

Chu and Chiu (1990) money-pumped Chinese students who exhibited reversals (as did Berg et al. 1985). Their subjects stated prices for two gambles and chose one of the gambles. If they exhibited a reversal, choosing the P-bet but pricing the \$-bet higher (denoted $c(S) > c(P)$), an experimenter would let the subject the \$-bet (collecting $c(S)$), make her switch the \$-bet for the P-bet (in accord with her choice), then buy back the P-bet for the price $c(P)$. The subject ended where she began, with no bets, but was $c(P) - c(S)$ poorer.

For most subjects who expressed reversals, it took roughly two arbitrage cycles to eliminate the reversals. Money-pumping a subject on the first of three gamble pairs reduced reversals in the second and third pairs (Berg et al. 1985) and that the magnitude of reversals, but not their frequency, was reduced by a similar money-pump procedure. These results suggest that in an environment where preference reversal is a recognizable, costly mistake that outsiders can spot and exploit, then people can learn to switch their expressed preferences (reduce the size of any discrepancy). But there is no evidence of whether subjects who are disciplined this way then learn to express preferences more consistently in the future, or whether reversals actually persist in natural settings.

Bohm (1990) conducted a highly original experiment, recruiting twenty Swedish students to choose between and bid for two used cars, a Volvo and an Opel. (The cars were actually sold!) His goal was to test for reversals of prefer-

ence in a high-stakes, natural setting with eager (self-selected) consumers. Most subjects chose the Volvo, and bid higher for it too. No subjects reversed preference (although four reversed weakly, choosing one car but making equal bids). Böhm concluded that reversals of the usual kind may be uncommon in natural settings. The problem is that there is no strong a priori reason to expect reversals in a choice between a Volvo and Opel. So it is difficult to tell whether reversals disappeared because of the extraordinary incentive in a natural setting with eager consumers, or because of the poor correspondence between the car choice and the P-bet/\$-bet paradigm.⁶⁰

Harrison (1990) criticized preference reversal experiments for providing inadequate incentive, because small misreports of prices cost subjects very little when the BDM mechanism is used. He showed that making the scale of price reports more coarse (e.g., forcing prices to be rounded to the nearest \$.25 or \$.50), or increasing the difference in expected values of the P- and \$-bet, reduced the number of reversals substantially. The effect of coarseness is not surprising (in the *reductio ad absurdum* case, if subjects could report only one price then no reversals would occur). Increasing the expected value, like the effect of money-pumping, makes expressed reversals more costly and does reduce them. The results point out that large differences in expression of preference are less frequent than small ones, but they do little to answer the pressing question of how large such differences are likely to be in natural settings.

Berg and Dickhaut (1990) studied incentives too. Their work uses the two-error-rate model (Lichtenstein and Slovic 1971). That model assumes a fraction q of subjects are truly risk averse and prefer the P-bet. Subjects make errors in expressing choices with probability r , and rank the prices of the two bets backwards with probability s . Certain values of q , r , and s lead to certain fractions of subjects choosing P- or \$-bets and setting higher prices for P-bets or \$-bets. The observed fraction of reversals can therefore be used to estimate q , r , and s . The error rate model provides a clearer view of how subjects respond to incentives than overall reversal rates, because incentives can reduce error but *increase* the number of reversals.⁶¹

Berg and Dickhaut showed that experiments that used hypothetical choices or interdependent incentives were best fit by a three-error-rate model in which pricing errors by risk averters (who prefer P-bets) were more common than pricing errors by risk seekers. Experiments with incentives for truthful revelation of prices (using the BDM procedure) were best fit by the two-error-rate model. Experiments with arbitrage (Berg et al. 1985) reduced error rates substantially. Curiously, the error rates *inferred* from the data were mostly greater than .5. This casts doubt on their interpretation as errors in expression of preference, and poses an important puzzle for work using the error rate models.

Berg and Dickhaut also ran experiments in which subjects earned payoffs in points that were convertible into lottery tickets, rather than in dollars. Risk aversion and preference were induced by converting points into tickets with concave and convex functions (à la Berg et al. 1986). Since prices were stated in points, just as payoffs were, the contingent weighting explanation of Tversky et al.

(1990) predicts the same rate of reversals as in experiments with dollars (unless compatibility of a familiar dimension like money is different than an unfamiliar dimension like points). But reversal rates (and error rates) were actually much lower using points, and most reversals occurred when risk-seeking subjects chose the \$-bet but priced the P-bet higher. Their paper suggests error rates are sensitive to incentives and raises a new puzzle, since opposite reversals were observed with point payoffs (as in Casey's reverse reversals with buying prices for high-stakes gambles).

3. Reversals and Markets

Knez and Smith (1987) studied market trading and preference reversal. In their experiments, subjects hypothetically chose between a P-bet and \$-bet and hypothetically valued both bets (giving selling or buying prices) four times. Between each choice-pricing iteration, they traded the bets in a separate trading period for each bet. The market trading experience reduced the incidence of preference reversal across the four hypothetical choice-pricing iterations, from about 60 to 40 percent. Subjects also sold or offered below their stated minimum selling price, or bid or bought above their maximum buying price, about a third of the time. (The average violation was substantial, about \$1 in bets with expected values of \$3.85.)

Cox and Grether (1991) studied the influence of markets more thoroughly. They compared selling prices elicited three ways: with the BDM procedure, with sealed bid second-price auctions (see chapter 7), and with an English clock auction in which prices fell steadily until only one person remained willing to sell. Notice that the sealed bid auction requires subjects to state prices directly, while the English clock auction requires a series of choices (i.e., whether to sell at the current price or not). The two markets therefore compare the choice and pricing modes of expressing preference.

Cox and Grether observed a typical rate of predicted reversals (around 60 percent) in the BDM and sealed bid conditions. Reversals in the English auction were much more symmetric; pricing the P-bet higher while choosing the \$-bet was more common than the opposite, familiar reversal. Since the English auction is like a series of choices, disappearance of asymmetric reversals is strong evidence that choice-pricing discrepancy underlies reversals. When all three tasks were repeated the rate of predicted reversals declined slightly using BDM (though here they have few data) and both declined substantially and became more symmetric in the two markets. In the markets, bids were highly correlated with the last market price, which suggests that markets reduce reversals by giving traders an observable price to anchor on when generating bids.

Cox and Grether varied incentives too. Some subjects earned a fixed payment of \$10 while others played all their choices and earned large amounts (an average of \$59) or small amounts (\$36). Incentives made little difference in the BDM procedure. In the market tasks, however, fixed-payment subjects behaved in the opposite way to the others, exhibiting no systematic reversals in sealed bid auc-

tions and strong reversals in English auctions. Their data provide a striking example of how incentive effects can vary in predictable ways across domains. In the English auction, watching the price fall is dull. (The fall is larger for the high-payoff \$-bet, since the clock starts at each bet's maximum payoff.) Fixed-payment subjects often dropped out of the auction quickly so they could ignore the computer screen and read newspapers or daydream, thereby establishing high selling prices-especially for \$-bets, because the fall is slower and more boring-and a high reversal rate.

4. Social Comparison and Reversals

Loewenstein, Blount, and Bazerman (in press) report a novel type of reversal that springs from the tendency of people to compare their outcomes with others'. They went to a class to recruit subjects for experiments. They recruited students in one of three conditions. In one condition, subjects could earn \$7 for forty minutes of work; 72 percent of the students ($n = 39$) agreed to participate. In another condition, subjects could earn either \$8 or \$10, depending randomly on the last digit of their social security number. (No convincing explanation was given for the disparity in wages, a weak point in the study.) Of the students who would have earned \$8, only 54 percent ($n = 44$) agreed to participate. In a third condition, students could choose to participate in either experiment. Of those who chose to participate at all, 22 percent chose the \$7 experiment and 78 percent chose the \$8 to \$10 experiment (when they would earn \$8). A preference reversal occurs because *fewer* students participated for \$8 than participated for \$7 when they considered the experiments separately (54 percent versus 72 percent), but many more students elected the better-paying \$8 experiment over the \$7 experiment (78 percent versus 22 percent) in a direct choice between the two. Students appear to compare their wages with others when rating any one activity (see also Bazerman, Loewenstein, and White 1992), but weight their own wage more highly when choosing among different jobs. This intriguing finding (which is only marginally significant) deserves further exploration.

5. Some Conclusions about Preference Reversals

The discovery of systematic preference reversals is now about twenty-five years old. Economists have concentrated on the role of incentives and incentive mechanisms used to elicit preference. Increased incentive appears to lower implicit error rates, though reversal rates are not always lowered (Berg and Dickhaut 1990; Harrison 1990). Arbitrage reduces the magnitude of reversals, and sometimes their frequency (Berg et al. 1985; Chu and Chu 1990). Experience trading bets in markets appears to reduce reversals, perhaps by giving subjects a way to establish a gamble's worth (Knez and Smith 1987; Cox and Grether 1991). No studies in new domains-repeated gambles (Wedell and Bockenholt forthcoming) and used car auctions (Bohrn 1990)-did not find reversals, which simply shows that not all pairs of choice objects are prone to systematic reversals. Perhaps the clearest

result from the last five years of research by economists is that theories that trace reversals solely to problems using the Becker, DeGroot and Marschak mechanism are wrong (see Cox and Epstein 1988; McDonald, Huth, and Taube 1989; Berg and Dickhaut 1990).

Psychologists have been less interested in the roles of incentive and experience, since the replication with casino gamblers by Lichtenstein and Slovic (1973) suggested neither variable was important. Instead, they suspect reversals are caused by a difference in revealed preference that results from different procedures used to elicit preference (Tversky, Slovic and Kahneman 1990; Mellers, Ordóñez, and Birnbaum 1992). Intransitivity appears to play some role too (Loomes, Starmer, and Sugden 1989, 1991).

There are at least three obvious directions for further research. First, since errors underlying reversals, and the dollar size of reversals, can be pounded down by enough incentive and arbitrage, an open question is how much discipline economic settings actually provide. (Knez and Smith [1987] start in this direction.)

It would also be useful to take the psychologists' explanation for reversals-that preferences are procedure-dependent-more seriously. One direction aims at the individual level: for example, Luce, Mellers, and Chang (1993) give a theory for how certainty-equivalents could be constructed for a set of gambles, and yield a preference order systematically different than would be observed in pairwise choices.

Another direction is the market level. (A start down this path is long overdue.) Two examples spring to mind. First, in choice among commodity bundles, if one commodity is the numeraire then the marginal rates of substitution will value that commodity more highly, *ceteris paribus*. Similarly, in negotiations over alternative multi-attribute settlements, the attributes being adjusted to make a settlement acceptable will seem most valuable (cf. Tversky, Sattath, and Slovic 1990). Second, different exchange institutions correspond to different response modes or procedures for eliciting preference (Machina 1987, 140-41). Buyers in posted-offer markets make choices; bidders set prices. If preferences depend on response modes, prices and allocations should differ systematically across institutions (familiar territory for experimental economists; see Cox and Grether 1991). Chapter 7 offers a related interpretation of the discrepancy between sealed bid and open outcry auctions.

J. Endowment Effects and Buying-Selling Price Gaps

Economic theory predicts that the prices a person will pay to buy and sell an object should be about the same. But a wide variety of studies indicate a large gap between buying prices (measuring "willingness to pay," or WfP) and selling prices ("willingness to accept," or WTA). See Kahneman, Knetsch, and Thaler (1991) and Hoffman and Spitzer (1993) for reviews.

The buying-selling price gap, or WTA-WTP gap, was discovered by environmental economists in the 1970s. For example, Hammack and Brown (1974) found that duck hunters would pay \$247 each to maintain a wetland suitable for

ducks, but liked \$1,044 to give up the wetland. Many other studies reported similar large gaps (see Cummings, Brookshire and Schulze 1986), ratios of median VTA-WTP around two or more. "Contingent valuations" like these⁶ are useful for doing cost-benefit analyses to make governmental allocations of non-market goods, and the gap between buying and selling prices raises the difficult question of which price is more appropriate. Knetsch and Sinden (1984) studied price gaps for lottery tickets (with cash or a gift certificate as the prize). In a typical experiment, nineteen of thirty-eight subjects would pay \$2 for a ticket, but only nine of thirty-eight would sell at that price.

I. Market Experiments

An immediate concern was whether these gaps would persist in repeated market settings. Courney, Hovis, and Schulze (1987) studied an unusual "bad," the obligation to hold a harmless bitter-tasting liquid called SOA in one's mouth for twenty seconds. The buying (selling) price was the amount one would pay (accept) to get rid of (assume) the obligation. Prices were bids in a uniform price Vickrey auction in which the four high bidders paid the fifth-highest price. There were large gaps in hypothetical valuations made before the series of auctions, but repeated auctions reduced the gap substantially, to a ratio of 1.5/2.5. However, their conclusions have been disputed (Knetsch and Sinden 1987; Gregory and Forby 1987) because they are especially sensitive to outliers and skewness in WTA values.¹

Boyce et al. (1992) auctioned off houseplants that resemble pine trees (called Norfolk Island pine²). They used the BDM procedure to elicit prices. The prices were then used as bids in an auction among subjects. (Their procedure adds an intermediate step to the usual procedure of having subjects bid directly.) Mean buying and selling prices were \$4.81 and \$10.00. Prices were substantially higher, \$7.81 and \$18.43, when subjects knew that any trees they did not keep or buy would be destroyed by the experimenters. (One subject was drafted as a witness, for credibility; some squeamish ones refused.) They suggest the increase in prices captures the "existence value" people place on mere existence of the trees.

Kahneman, Knetsch, and Thaler (1990) ran several market experiments. They first conducted a choice-based sealed-off-bid auction with tokens of known value, to test whether subjects bid their true values. (They did.) The token market established confidence in the auction mechanism elicited good approximations to true values. Then they conducted auctions with coffee mugs, pens, and other commodities. The median selling price for a mug was \$7.12; the median buying price was \$2.87. Prices did not converge much across four trials (one of which was closed and played afterward). Since mugs were randomly allocated to begin with, if buying and selling prices were the same roughly half the mugs should be traded, but only about a quarter were.

Franciosi et al. (1993) replicated these experiments by changing the instructions to remove terms such as "buyer" and "seller" which, they thought, might be overstating the buying-selling price gap. They observed significantly lower sell-

ing-price values (\$5.36 versus a mean of \$6.89 in the KKT data) but the gap between buying and selling price was still large. They also replicated the finding of undertrailing using a uniform-price double auction (in contrast to the sealed bid mechanism used by Kahneman et al.).

Several experiments studied markets for lottery tickets with money prizes. P. Knez, Smith, and Williams (1985), M. Knez and Smith (1987) tested the hypothesis: buying selling prices with actual trading in markets. Subjects routinely paid more in the test than their stated buying price, or sold for less than their stated minimum selling price. Trading volume was only slightly lower than expected if VTA = WTP.

McClelland, Schulze, and Courney (1991) ran Vickrey auctions for lottery tickets over gains and losses. Elicited buying and selling prices were close together for gain tickets, but selling prices for insurance on loss tickets were roughly bimodal (either zero or several times expected value) and were larger than buying prices.

Harless (1989) measured buying and selling prices using a uniform price Vickrey auction with a within-subject design. Each subject gave a buying price, immediately followed by a selling price (or vice versa), and paid or received money immediately after each lottery. The within-subjects design enables a buying-selling price ratio to be calculated for each subject. The modal ratio was one, but several ratios were very large. The median and mean of the ratios were 1.3 and 2.7. The low ratios Harless observed show the greater consistency that can sometimes result from a within-subjects design, especially when two tasks follow immediately in time. Kachelmeier and Shehata (1992) also measured buying and selling prices within-subjects. The ratio of median prices in their study was about two, and the difference in prices was highly significant.

Overall, the data suggest that competition or learning in markets reduces the buying-selling price gap somewhat, in some settings, but does not eliminate it. The gap is large for environmental and consumer goods (like wetlands and mugs) and small for lottery tickets. Perceptions of the investigators may also play a role, perhaps through the differences in the designs or domains of application they choose.

2. Explanations Based on Experimental Artifacts

The experiments have been careful to eliminate several artifactual explanations for the buying-selling price gap. First, it is possible that subjects do not make hypothetical valuations of duck wetlands or pine trees very carefully, or strategically state high selling prices and low buying prices. Direct comparisons indicate there is some misrepresentation, especially in hypothetical selling prices (e.g., P. Knez, Smith, and Williams 1985; M. Knez and Smith 1987). But in most of the experiments subjects were paid, or got to keep goods they bought or didn't sell, and price gaps persisted.

Second, sellers are wealthier than buyers because they own the object being sold. Wealth effects cause a legitimate buying-selling price gap that can be large

under very special conditions.⁶¹ Studies show that wealth effects explain essentially none of the observed price gap. Coursey et al. (1987) controlled for wealth effects by endowing buyers with \$10 (but see Knetsch and Sinden 1987). Franciosi et al. (1992) regressed prices given by a subject against the wealth accumulated by that subject earlier in the experiment and found no apparent wealth effect. Kahneman et al. (1990) controlled for wealth effects by allowing some subjects to choose between a mug and a sum of money, for several possible sums of money. These "choosers" are in exactly the same wealth position as sellers endowed with mugs, but their median valuation was only \$3.12, close to the median buying price (\$2.87) and much lower than the median selling price of \$7.12.

3. Endowment Effects: Some Psychology and Implications

The leading psychological explanation for the buying-selling price gap is called the "endowment effect": people prefer the things they own, *ceteris paribus*. Endowment effects are thought to arise from the "loss aversion" assumption in prospect theory—losses are more painful than equally sized gains are pleasurable (see Tversky and Kahneman 1991). The stylized fact that buying-selling price gaps are larger for environmental and consumer goods than for lottery tickets (in most experiments) suggests that the gap is larger for goods bought for use, and smaller for goods, gambles, or securities that are routinely sold or easily valued.

Endowment effects are related to at least five other psychological effects. These phenomena are conceptually distinct, in principle, but are empirically entangled ("confounded") in some experiments. All of them follow from two principles: valuation relative to a reference point, and loss aversion (Tversky and Kahneman 1991).

1. "Status quo bias" is an endowment effect in which having a current choice, or default option, enhances preference for it (see Samuelson and Zeckhauser 1988; Knetsch 1989; Hartman, Doane, and Woo 1991). For example, New Jersey drivers now get cheaper insurance, restricting their right to sue, unless they pay extra. Only 17 percent of drivers paid extra in 1988. Pennsylvania drivers make the same choice, but their default option is the more expensive, unrestricted insurance. More of them chose the default option, paying extra for the right to sue, than in New Jersey (Johnson et al. 1993). (Part of the difference could be due to transactions cost, of course—the cost of filling out a form and sending it in—but experiments in which subjects must choose, forcing them to pay the transaction cost, show a comparable status quo bias.)

2. Buying-selling price gaps can result if people are more sensitive to overpaying (which incurs an out-of-pocket cost) than to selling too cheaply (an opportunity cost), as they appear to be in other domains (Thaler 1980).

3. Ritov and Baron (1991) show that people treat errors of commission, or action, as more blameworthy than errors of omission, or inaction. The reluctance to pay too much is an action error; passing up opportunities to sell is an inaction error. Greater fear of action errors will make buying prices too low; ignorance of inaction errors will keep selling prices too high. Schweitzer (in press) showed that

status quo bias is largely due to a bias in favor of inaction. (Subjects preferred a default option, which would be chosen if no action was taken, even if it differed from the current status quo option.)

4. There is some evidence that the purchase price of assets matters in financial decisions, creating a "disposition effect," because people are reluctant to take actions that create an irreversible loss and are eager to take actions that create gains. For example, trading volume is lower for stocks that have fallen in price (Ferris, Haugen, and Makhija 1988; Weber and Camerer 1992). Casual observation suggests the volume of houses sold falls when housing prices fall. In experimental asset markets, volume appears to thin when bubbles burst (Smith, Suchanek and Williams 1988, see chapter 6).

5. Marshall, Knetsch, and Sinden (1986) asked people whether they would buy or sell objects at certain prices, and how they would advise others. The advice people gave others revealed no buying-selling price gap—the gap disappeared because they urged buyers to pay more—which suggests endowment effects are not recognized or encouraged in giving advice.

An important question is how endowment effects change economic predictions. Tversky and Kahneman (1991) show how theories of choice and exchange can be altered to accommodate endowment effects (which they call "reference-dependence"). Bowman, Minehart, and Rabin (1993) showed how a prospect-theoretic kind of loss aversion could explain empirical anomalies (and predict some new surprises) in economic theories of optimal consumption, savings, and asset pricing. Hardie, Johnson, and Fader (1993) found that an expanded logit choice model, incorporating reference-dependence, fit within-household time series data on orange juice purchases better than a conventional model.⁶⁷ Kahneman, Knetsch, and Thaler (1990) showed that trading volume in goods markets was substantially reduced by endowment effects. Hoffman and Spitzer (1993) point out that if marginal buyers and sellers in a market have no endowment effects—if sellers are firms, for instance—then market prices and volumes may not be affected. Endowment effects are therefore likely to have their largest impact when individuals buy and sell on both sides of the market—residential housing, for instance, or some kinds of labor.

Rietz (1991) conducted market experiments that exhibit a surprising endowment effect^M (and deserves further study). His subjects trade state-dependent contingent claims which pay 1,000 francs (= \$.50). A Blue ticket pays off if the Blue state ($p = .3$) occurs. A Green ticket pays off if the Green state ($p = .7$) occurs. The states are mutually exclusive and exhaustive, so a portfolio of one Blue ticket and one Green ticket pays 1,000 francs with certainty. Subjects are initially endowed with at least two portfolios of tickets (i.e., two Blues and two Greens) in each period. In addition, each subject was endowed, in alternating periods, with either four extra Blue tickets or four extra Green tickets.

Behavior in the double auction markets for tickets tests various theories of prices and allocations. After sixteen trading periods, mean prices were around 350 for Blues, and 850 for Greens, well above the expected values of 300 and 700. The sum of the two prices was often 1,200 or more. Prices this high are bizarre

because subjects always began with at least two pairs of Blue and Green tickets. They could sell a pair of tickets, which are worth exactly 1,000 together, for more than 1,000, but they did not do so frequently enough to drive prices down. Furthermore, rational subjects should end a trading period with an allocation of tickets that depends on their risk tastes, but does not depend on their initial endowment. They did not: in periods when they start with four extra Blues they ended the period with more Blues than when they started with four extra Greens and no extra Blues. (Carlson and Johnson [1992] observe a similar endowment effect in experimental bond auctions.)

The puzzling prices and allocations Rietz observed can both be explained by endowment effects: suppose people value whichever tickets they start with more highly. Then Blue-holders will ask a high price for Blues and Green-holders will ask a high price for Greens, which pushes prices above expected value. Because people cannot sell many tickets at the high prices they bid, their final allocations depend on their initial endowments.

Endowment effects have a natural application to law (Hoffman and Spitzer 1993). The Coase theorem presumes that the valuation of a property right is independent of who owns the right, an assumption that is questioned by observed buying-selling price gaps. Cohen and Knetsch (1990) suggest that the law frequently recognizes the special losses that result from reducing one's endowment, and assigns property rights so as to minimize those losses (though see Hoffman and Spitzer 1993).

K. Search

1. Search for Wages and Prices.

One setting in which individual decisions have direct market consequences is search. There is a large theoretical literature on search (e.g., Lippman and McCall 1976). In a typical model, a person looks for work each period. With some probability, she turns up a job offer drawn from a distribution. She can accept the offer or search more. If she decides to take the old offer after searching more, it is available with some probability that depends on how much time has passed. The models can apply to consumer shopping and other settings too.

The optimal strategy is usually to set a reservation wage and accept any job that pays more. The optimal reservation wage comes from a difficult recursion on a lot of series of recursions, if the time horizon is finite. There is little empirical evidence about whether people search as the models predict. So experimental evidence is useful.

Schotter and Bratnućin (1981) and Braunstein and Schotter (1982) studied several variants of the Lippman model with an "infinite" horizon (i.e., subjects could search as long as they liked). Subjects stated reservation wages but were not bound by them. (they could accept lower offers). In a baseline condition mean reservation wages were amazingly close to optimal (134.5 vs. 133) but subjects

sought too little (3.7 actual periods when 4.5 was optimal). When parameters were changed—whether rejected offers could be recalled, the cost of searching, the dispersion in the offer distribution—stated reservation wages and the actual acceptance of offers changed in the correct direction but differences were not always significant. When risk aversion was induced by paying subjects a concave function of points they earned (anticipating the method of Berg et al. 1986), their reservation wage was much too low (100 when 130 was optimal).

Kogut (1991) studied search in infinite-horizon settings too. Subjects paid a search cost each period (typically \$0.08) and drew price offers from a known uniform distribution. In each trial they paid the price they accepted, and search costs, and earned a known value. Their reservation price should be the same each period, which implies that optimal searchers should never reject an offer, then go back and accept it later. They did accept old offers, about a third of the time. They often stopped searching too early (even assuming a large degree of risk aversion). Most of the early stops occurred just before the total search cost was large enough to make their profit negative, suggesting some subjects are sensitive to sunk costs and not to marginal costs, and benefits of search.

Cox and Oaxaca (1989) studied search in finite-horizon experiments with twenty periods of search. (They chose finite horizons to establish more control than was exerted in the infinite-horizon setting of Braunstein and Schotter.) Assuming they were risk neutral, subjects quit searching at the optimal point 80 percent of the time. Searches ended an average of half a trial too early. As in the Braunstein and Schotter experiments, search duration responded in the right direction to several parameter changes, including search cost and offer dispersion, but the size and significance of the changes was not optimal.

Cox and Oaxaca (1990) replicated their earlier results but asked subjects to precommit by stating reservation wages that were used to automatically accept or reject offers. (Precommitment made little difference except for a slight increase in risk aversion and some learning effects that disappeared after several trials; Cox and Oaxaca find, presumably, that subjects followed the optimal path reasonably well to optimality under risk aversion could not be rejected), except they were too low at the start and too high at the end; overall, subjects searched too little, compared to a risk neutral benchmark.

Hey (1982) studied search in a shopping setting where the price distribution was unknown. (Subjects were not paid.) Recorded statements of subjects during the experiment ("verbal protocols," to psychologists) suggested simple rules of thumb. One rule was a reservation wage strategy. Other rules prescribed stopping points depending on the previous sequences of price offers (e.g., stop if the current price is above the previous price). These rules might be manifestations of judgment biases discussed in section 4, since many of them try to capitalize on apparent nonrandomness in the price series. Hey found some subjects using each of the six rules, or mixtures of them.

In Hey (1987) subjects were paid and half of them knew the price distribution. Knowing the distribution increased subjects' use of the optimal reservation price

strategy, but financial incentives did not. The ability to reject offers actually hurt slightly (reducing overall profit), which is puzzling. And they searched too little (perhaps due to risk aversion).

Moon and Martin (1990) extended Hey's work. They spelled out several more alternative heuristics subjects might use. In their data, cutoff rules such as "wait for 11 price k standard deviations above the mean" explain decisions roughly as well as the optimal theory. Simulations show that heuristic rules can be very close to optimal (only 1 percent worse).

Harrison and Morgan (1990) studied several search problems. In variable-sample trials, subjects could buy a sample of n offers in each period k . In sequential trials, subjects could only sample one price at a time ($n_k = 1$). In fixed-sample trials, subjects could only sample for one period ($k = 1$). In theory, the experimental design in the variable-sample method should enable subjects to earn about 10 percent more profits than in the sequential or fixed-sample trials.

Subjects did exploit the freedom in variable-sample trials. They chose bigger samples than in the $n_k = 1$ sequential trials and searched longer than in the $k = 1$ fixed-sample trials. They earned substantially higher profits too, but the increases were not significant by nonparametric tests. The direction and raw size of deviations from optimal sampling are not reported (a result of the authors' obsession with the cost of deviations), but the fraction of subjects who made errors of a certain cost are reported. Subjects are apparently good at deciding whether to keep searching, but not as good at choosing the number of offers to sample each period (cf. judgment errors in section 4 above, especially 4.D.3).

Since the optimal strategies in these search problems are difficult to derive, the subjects' approximation of them in many of the experiments is generally impressive. But there are some anomalies, especially in responses to parameter changes. And in general, people search too little, compared to the amount of search recommended under risk neutrality. It would be useful to induce risk neutrality, to ensure the degree of risk aversion to test whether the observed undersearch can be rationalized by risk aversion. It would also be helpful to know whether heuristic rules could produce such an impressive approximation to optimality (situations by Moon and Martin [1990] suggest they can). Heuristic rules might also explain the persistent tendency to undermatch and the relative inability to choose optimal sample sizes.

A natural extension is to experimental markets where sellers choose prices while buyers shop around. Sellers must understand how buyers search to set prices optimally. A variety of theoretical models predict endogenous price dispersion that depends heavily on shopping habits. Grether, Schwartz, and Wilde (1988) report experimental evidence supportive of some models.

Another interesting direction is to reproduce apparent search anomalies from natural settings. For instance, Pratt, Wise, and Zwickhauser (1979) found that across categories of consumer goods, price dispersion (standard deviation) was a roughly linear function of mean price. This finding is rational if search costs are higher for more expensive goods. A competing behavioral explanation is that

people calculate the marginal benefit from shopping as a percentage, not a dollar amount; they search longer for a \$5 saving on a \$20 calculator than for a \$50 saving on a \$400 washing machine.

2. Search for Information

There are many psychological studies on the purchase of information that is used for making decisions. The results are much like those for search over wages and prices: people are insufficiently sensitive to factors such as accuracy of information and cost, which should affect search, and overly sensitive to factors that should be irrelevant, such as the source of information or the total information available (e.g., Connolly and Gilani [1982] and references they cite). One study found that providing subjects with a decision aid, which converted information into decisions optimally, reduced mistakes in information purchase by about half (Connolly and Thom 1987).

It would be useful to extend the psychologists' results to economic domains in which the value of information is derived from its use in making decisions, such as information markets coupled with asset markets (e.g., Sunder 1991; Copeland and Friedman 1992; and see chapter 6).

L Choice: Summary and New Directions

The studies reviewed in this section suggest a variety of broad classes of anomalies of the standard utility theories under risk and uncertainty. Many of the anomalies can be traced to the ideas that values are judged relative to a reference point, probabilities are not weighted linearly, and decision weights are not the same as beliefs. Preferences also seem to depend on the way choice objects are described (creating framing effects), the procedure by which they are elicited (creating preference reversals), and on one's current endowment (creating a buying-selling price gap). These phenomena and anomalies in portfolio choice and the purchase of information suggest people use simple procedures to make choices, constructing their preferences from procedural rules rather than maximizing over well-formed preferences.

At the same time, studies of search and market trading of risky assets (cf. chapter 6) suggest that models based on maximization are not badly violated (Plott 1986; Smith 1991). Future research should concentrate on three classes of explanation (and exploration) for the disagreement across studies: (1) experience, incentive, and discipline in markets combine to create stable preferences which are well approximated by normative models (contrary to the individual choice results); (2) anomalies in the models, which loom large in a large sample of individual choices, are too small to see in markets; or (3) studies of markets and search have not looked at the settings in which anomalies are likely to be large and common. (For example, buying-selling price gaps are largest with consumer goods and smallest with money gambles.)

IV. Conclusions, and Research Directions

My perspective in this chapter is unapologetically behavioral. I think the search for systematic deviations from normative models of individual decision making has been extremely fruitful.

Economists have had two reactions to data on individual decision making and have made two kinds of contributions, which might be called "destructive" and "constructive." Destructive tests, often motivated by skepticism, are designed to check whether apparent anomalies are replicable, robust across settings, or might be due to flaws in experimental design. My opinion is that some occasional tests of this sort are essential, but too much energy has been devoted to destructive testing with very little payoff. Not a single major recent (post 1970) anomaly has been "destroyed" by hostile replication of this sort.

Constructive reactions of economists to decision research have taken at least two forms. One reaction is the construction of alternative theories to explain anomalies. For example, Kanodia, Bushman, and Dickhaut (1989) show that the failure to ignore sunk costs—managers sticking with projects after learning they are bad investments—can be privately rational, for the managers, if there is information asymmetry about their talent. Theories of this sort are easy to construct (probably too easy); most of the theories posit information asymmetry, then show that an apparently irrational action—sticking with bad projects, following the herd (Scharfstein and Stein 1990), sticking with an inefficient status quo (Petrandez and Rodrik 1991)—is actually rational because the action conveys information. The main problem with this class of theories is that most posit a highly stylized economic setting much different than the setting created in the original experiments demonstrating the anomaly: at best they are sufficient explanations for an anomaly, but they are hardly necessary. Of course, in principle these explanations can be pitted against behavioral accounts and tested (e.g., Berg, Dickhaut, and Kanodia, in press). (I fear many readers do the opposite, exhibiting a "sufficiency bias" by taking sufficient explanations as implying the need to explore a behavioral phenomena further.) Experiments can play a special role because one can test theories of individual behavior directly and simultaneously test their implications in markets, rather than testing only market implications.

Another constructive reaction is expressed by Plott (1986) and Smith (1991). They frame the basic issue as a puzzle of aggregation: why do models that assume individuals behave rationally perform so well describing behavior in market experiments, if individuals behave irrationally in psychology experiments? There are three possible answers: (1) It does not take much rationality to behave nearly optimally in an experimental market; (2) traders learn in market experiments; and (3) market experiments overstate the degree of rationality in naturally occurring markets.

Gode and Sunder (1993) explore the first answer. They show, using simulations, that double auctions can be highly efficient even when simulated traders have very limited rationality.

Learning is a second answer, and has been insufficiently explored. For example, in Cox and Grether's preference reversal experiments, subjects' bids in a period are highly correlated with the previous period's winning bid. That correlation suggests markets are helping traders construct (or "discover") a preference by wringing others, rather than simply revealing their well-founded preferences.

The third answer is that experimental markets overstate the ability of natural markets to erase individual irrationality (The best answer along these lines can only come from further studies of behavioral phenomena in naturally occurring markets, which lie outside the scope of this handbook.) Two experimental approaches have been taken in exploring the boundaries of market magic. One approach begins with individual errors and constructs an experimental market in which they might persist, to search for domains in which experimental markets might fail. The other approach begins with market level anomalies and searches for explanations based on individual errors.

Several studies reviewed above took the first approach, creating experimental settings in which individual errors were tracked, and studying whether errors were reduced by market forces; (Duh and Sunder 1986; Camerer 1987, 1990, 1992b; Anderson and Sunder 1989; Camerer, L-Oewenstein, and Weber 1990; Ganguly, Kagel, and Moser, in press). These studies mostly show a tendency for prices to converge toward, but not to, Bayesian predictions. Psychologically predictable deviations persist. The data suggest the market glass is both half-full of deviations and half-empty because some deviations were drained away by learning.

The second approach tests whether anomalies originally observed in aggregate experimental data can be explained by behavioral models of individual choice. I mention three examples. Lind and Plott's (1991) alternative specifications of nonrational bidding behavior in low-bid common cost auctions with a seller's curse is one. Cox, Smith, and Walker (1983) is another (see chapter 7 for a longer discussion): they give two behavioral models to explain the observed difference between Dutch auction and first price auction prices, and run experiments to test the models. A model in which bidders violate Bayes' rule, by underestimating the risk of losing the auction if time passes, appears to be the better of the two. Guler, Plott, and Vuong (1987) is an especially sophisticated example. They ran experiments based on "zero-out" auctions for airport landing slots, in which a government authority rebates all bidding revenue according to a known formula. Because of the rebates, airlines could bid much more than their reservation prices. Indeed, bids increased explosively over repeated auctions. In one set of experiments and converged in another. The data present a puzzle for traditional analysis: competitive equilibrium predicted badly (because the rebate formulas meant that slot valuations depended on the bid of others) and Nash equilibrium predicted badly too. They then considered two classes of alternative models of bidder decision making: game-theoretic models in which decision rules are derived from system equilibrium conditions (e.g., winning bidders bid slightly more than losing bidders); and decision-theoretic models in which bidders form beliefs over important parameters (e.g., what the losing bid will be) and update their beliefs. Models of the latter sort presume no game-theoretic sophistication and

violate rationality because information is not used efficiently, but they predicted the path of actual bids better than the game-theoretic models.

Theorists, listen up: If a fundamental question for economics is whether individual decision making deviations persist at various levels of economic aggregation (households, groups, firms, markets, societies), then the aggregation question should certainly be studied theoretically as well as experimentally (and in the field too, of course). Theoretical progress is not impossible. In fact, much progress has been made. We now have a rather good understanding of some of the procedures people use instead of Bayesian judgment and utility maximization.¹¹ In the realm of choice, Machina (1982, 1989) has shown that some basic results in economics can be derived without expected utility maximization; many applications have followed (Epstein [in press] reviews some). Gilboa and Schmeidler (in press) describe choice models grounded in the vast psychological literature on analogical "case-based" reasoning. Models of heuristic probability judgment, adaptive expectations, limited memory and costly cognition, preferences that are reference-, context-, endowment-, and procedure-dependent, time preference with hyperbolic discounting, overweighting of low probabilities, and utility maximization with nonadditive probability are now well developed enough to prove tractable for many kinds of economic analysis. For example, variants of subjective expected utility could prove fruitful in explaining demand for information (to reduce ambiguity) and timid behavior that looks like risk aversion but disappears over time (it is probably ambiguity aversion instead).

Naturally, there is a tradeoff between analytical tractability of assumptions—the observed procedures will be less tractable—and their predictive accuracy. But it is hard to know how much tractability is lost in the new generation of behavioral models without trying them out for a decade or two.

Notes

¹ Thanks to Robyn Dawes, Dave Harless, Teck-Hua Ho, Charles Pion, participants at the Pittsburgh conference in June 1990, and especially Amos Tversky and the editors for help and extensive comments. Thanks also to the Russell Sage Foundation, where I finished writing this chapter during 1991-1992.

1. For example, Barro and Fischer (1976) wrote: "A fundamental difficulty with theories of expectations that are not based on the predictions of the relevant economic model [i.e., rational expectations] ... is that they require a theory of systematic mistakes" (163). (They go on to suggest such a theory is impossible.) In fact, the "heuristics and biases" paradigm in decision making is a theory of systematic mistakes.
2. The original work, much of it in meteorology, is by Brier (1950). Murphy (1973), Yates (1982), and many others in between. Yates (1990, chaps. 2-3) is a recent review.
3. Call r_i the reported probability of event $i = 1 \dots 11$. Then if event i occurs, the quadratic scoring rule pays $(1 + 2r_i - 1r_i/Y_i)$, the logarithmic rule pays $\log(r_i)$, and the spherical rule pays $r_i/(1r_i, 1)^{1/2}$. All three rules are "proper" (incentive-compatible). The quadratic scoring rule was first discovered by Brier (1950). The quadratic and logarithmic proper scoring rules were independently discovered by Toda in the early 1950s and reported in an unpublished 1963 paper. Van Naerssen (1962) and de Finetti (1962) rediscovered them. Roby (1965) discovered

the spherical rule. None of the latter seemed to know of Brier's earlier paper (in a meteorology Journal).

4. The spherical scoring rule punishes misreports of probabilities slightly more than the quadratic rule, and both are substantially more punishing than the logarithmic rule (except for misreports near one) (Murphy and Winkler 1970). However, the logarithmic rule requires an infinite penalty if an event that is called impossible occurs anyway. Fischer (1982) found that using a large penalty instead worked well.
5. That is, the standard error bars which are reported in early journal articles announcing estimates of physical constants, like the speed of light, tend to not include very recent estimates.
6. Lightman and Gingerich (1992) suggest a revised version of Kuhn's hypothesis about "paradigm shifts" in science, along these lines. Kuhn's idea was that the weight of anomaly would eventually topple any theory, when it could be replaced by a new theory that could explain the anomalies. Lightman and Gingerich suggest, instead, that many observations are not even recognized as important anomalies—because of misperceptions like the red-heads phenomenon—until, in a kind of "retrorecognition," a new theory arises that can explain them and that makes clear how anomalous they are for the old theory.
7. In economic terms, media information provides entertainment and creates an externality by affecting one's stock of knowledge about the world. If consumers do not recognize the knowledge externality, distortions will result. For example, people may overinvest in preventing accidents-avoiding airplanes—while underinvesting in nutrition and exercise to prevent disease.

$$\begin{aligned}
 8. \quad P(\text{Blue} | \text{Blue}) &= \frac{P(\text{Blue} | \text{Blue})P(\text{Blue})}{P(\text{Blue} | \text{Blue})P(\text{Blue}) + P(\text{Blue} | \text{Green})P(\text{Green})} \\
 &= \frac{(.8)(.15)}{(.8)(.15) + (.2)(.85)} = .41.
 \end{aligned}$$

9. For example, in memory studies subjects are often asked to remember nonsense syllables (such as ZEV and ZOV) to deliberately eliminate the influence of prior memory.
10. The informational-wording argument draws on the work of linguists, who show that listeners expect a kind of cooperation from speakers (e.g., Grice 1975). Therefore, subjects may read between the lines of a word problem and respond appropriately to a perceived problem different than the one actually posed.
11. For example, hot hand effects might be important in labor markets for managers who do separable projects where luck plays a role—like movie producers or studio heads.
12. Argote et al. (1986, 1990) and Snizek and Henry (1989) studied group judgment in Bayesian tasks like those described in this section. They found that aggregation of opinions in groups did not reduce errors much, but there are many open questions about group aggregation. Furthermore, group judgment forms a natural comparison for market-level tests and bridges the gap between traditional domains of economics and psychology.
13. Take all the periods with the same sample (say, O-black) from a particular experimental session, and arrange them in chronological order. Do the same for all the experimental sessions. Now form a 90 percent confidence interval using the mean prices in each of the first O-red periods from all sessions. Form a second interval using the second O-red periods from all sessions, and so forth. Stop at N when the number of sessions with N O-red periods falls to three or less. (Note that the sample of mean prices shrinks as one goes along, because some sessions had very few O-red periods, raising the standard error across sessions and widening the confidence intervals.) Do the same for each of the four possible samples. The procedure is conservative. It treats each session period's mean price as a single data point and draws inferences across sessions (which we can presume to be statistically independent).
14. For instance, inexperienced subjects regard a sample of 1 red (Bayesian $p = .75$) as being almost as indicative of the X state as a 0-red sample (estimated $P = .85$ for 1 red v. $P = .923$ for 0 reds). If the mistakes were any bigger they would violate a kind of monotonicity condition that requires $P(X | k \text{ reds})$ to decrease in k .

15. Robyn Daw, et al. >HOrder example: d11mp1st who believe that t1uld Jbuser> nev er Rnp without thenipeuc treatment h1vt thmr beliefs reinforced every time 1h WJLr v1rts their dink. The abusers who white the theory never CQC to the clime!
16. The f-mXOM, f, f1 raixoru walk, collected by Dwyer et al. (1993) MC the nlo&t Indicativ of rationality of expectations. But even their data am m1rnd: when fore,;asv; are pooled th1ir variance is greater than the variance from a naive f=a s t (which is m optimal f=s t for a Wldom walk); and (w(r,tep-l1he d f0R'CS\$,ould be exactl, equal to one- ahead forecasts if subjects know the proce-s is a random walk, but tM) are not.
17. Adapovene, couflicl, with raJonalit of expe;tbtion bec*w,e it implies 1hal ?WP use information about a previw; forecast error lo alter their n-it foreca;1. Bm if forocast5 att rational w begm with andmnly diotrlbmed around the corre;t price—1hen the previl)ls foreru1 error contains no information and should bf ignored (leading to b"" zero in 11J). Adaptiveness and bitt- in forecast errors are dlerefore Ukly to g5 hand in h11ld, ut ii is pc,sible that forecasts could 1X unbiw;ed :md-oot r liring that-torncastr exhibit 1lap-tiveness,.
18. It rmneo little whether weights are derived by regre;ins; pw(outioomes on ob<ervlddn; {"acuarial" model6), by randardlcng oos<Vablea for vari anl wmghting them eqwdly, Of by regres ing expert judpuient, on ob;.;1-'@lo ("boomrapping" modcis); they inipmv.i accuracy by dimuding (oowy teg:noolQn residuuls).
19. After discovering their M1om of ex utili,y, 'Morgenstern (1970. 809) repons that von Neumann "called out in awmislment 811 didn't: anyoru, . u mat?,"
20. Norman (19U) mllke> a sintlar poiul ab<ut pmdct design. M tedmolsy i:n;rove.. it h easy to de1;gn produm with more function. TM tricky plrt 1s X11V<ying the fncfionu. to conMlmeti, o they can be quickly WJertood 100 1&ad y father bauglu. a VCR d1ac can record up to eighl p10gr1v1& a ytw: in arlvance. He cannot figure oot how ro<3TU>ven ow: program one day in advance.
21. Magical t1inking is akin 10 1WQ false "1uv.8" andropo;ollis\$ tudy: the law of snrahmy (sintihw object, &hare properitt) and the law of contagm (OO Jan transfer by touching). For iMt1Ultt, some primnive poople, even civitiid ehi n. are afraid of photographs or dotl resembling tigen; because trunk th-Om>1>j&t 1F< a dan nt hgw; 1F< Inruite behefs, like ilunc 1F< important even in mndern uullnes. Many species of are endangtmi-the black rhinocem; Merwan sea turtles, some se ml- nous people desire w eat the aui, or weai; pml, of them. to 1mp1O'< t1r1r heath nr ie:ural potency (Gilovkh 1991,cilap, 1J,
22. A> example of m 1U1Controllable, inf<J17M1* uigrutl is the pn{omtlrnr lif other tiftlil in 1111 industry, Th11 effort of oomplllly X', 1W11ale1S doo' oot aft&t othef firm; p;rf 1iy much), bul the performnce of other firms is mfotmative; about csm1MO industry,wideshocb that oompany X faC<d, and hence 1; in j1JdgLag x's pe-rfunnance.
23. E.g., MOrhina (SS112, 1987); Sugdea (1986), and Weber and Ciutwm {1987),
24. As Ed ards 0954c, 3S6> ;oinwd ow; the Hieb and Aile> paper "W1; fur oonoraics something 11ac the bch111orist revol.Kloo in psychwlogy." 1t wwed the seem; fur modem erooo-misi' dim1Vt W1U1ley evidence or introspect1oa, and alm(x;t etthmve relian;e oo choices a dato for C>ru11cting theorica of ilchavtor,
25. That J, crmld theeeffillntyequivilletuX' that w1ve5X' p1! • (1 ~ p;L alrobe got by w1vi1S for X: in the dif'ference-cowparnKmu<H- u{X'1 e:'4X')-1t(LJ1 Ifso. then thevoo Nf11111111r Mrugenstem utility funi:tkn j; a rillldeSS value funcuon (h la 8emouili) loo. While ,nme piwage, 1h VOn N1runt1OU and M<qellSttffi's hook an equ1v;leare, i is shrilltgly dwtd in othtt pla;es, See Fmlhum (in pres;) for w1 ;;;-egesk.
26. AJJus felt that a nc(1-Bemo1J1liatt' value ftiar:tioa should uadvrly the thtt:ey of clv;iet1 indk'' nllk and rhk avmioa C1J1d he captlrod by 1NWS1011 to 1uiance tn value o. di11O1ion of probability. The J;t two hum1md of Allrut and Hag,1 (1979) ail 1n; views.
27. They induced v..Ue xy the player with d1c moo. pomu: 11 the end 11toosc ;aru1y. cigmt:tteS, or cigar, as a prized). It 1S W1j1 11J11le 1ha! wV<rai 11J11oc1 were BCW1y ir:tmbl,

- ("in many cases 1hey were 00ave;1 making active nre of 1his i 1 1he&y" 11116); fucir respos.es were no 1iffetenl
- ii_ The Prestoo-Saraua curve, and thore dmved by many othen, C1W01be Y1E11 io re&ci: pm, probability we.gilt., becaJie did not control for tOn1111111ty of utility in deriving i: Fw1hemKm., they interpreted 1heir weights; ,u psydrok,gkal U11M1ations with 1he funn1l properties of pmhibi1ries (satisfy111g additivity. etcJ which Kaime111111 and T1H1Uy'a d<i-ion ght3 need not satisfy.
29. The Guardsmen were mud! closnr to risk neutral when insuuttirtg -ge11tt how to bet for t1lem. which wggests wmc wilyly fur participation as f11k pre{trt:noo'4-bl1tt they madebels .
30. Edwa1U1\ 1954c) pointed out that bo<h oatcomc; and probability could be weightro linearly or nonlinearly and oomh1ned, resulting in four poS11blr. models, fu'poc'led -1UIC 1inearity of both, e11peel.00 utility &f1111ed linearity of pri;bm1ity only, 11B modtl USUmOO linearity of G1tcomes ooly {until Hinda 1977 and Yaari 1987}. and Edwards proposed an e1;temron of expoc1Cd utility in which both were 0011111Ut He used 1he tetm "utjective expected utility to refer 1Ohis 11J11111111, bur J will refa to "nonlntarptob.1:bility W1111111" instead since Savage (1954) defined subjl1ctive d utility diff1c1fill:ly thw1 EdWards.
31. m their inflation experiment. Daniels and Ploo {198S) tested whether 00Y111 fon;c:med i m-- prices than sellen, reflecting 1had of wishful thinkmg (or richn: of al belief drat forecasting low prices would make prices low), lo forty-eight of sevtmy-t1f1h periods :vera buyer- and £E11<f W11 differem (p"" .02, 000-Wledi. Thi 1 a good example of the kind of behaviorally proro<Wve finding thu emMget by flather111g rich& kind; of data, and asking more unorthodox quemtrns. tha! experimenral erooortu;1S olW1 1b
32. That is..Supposea,b,c, anddare11 meli<edutilitm; a >handc>dshontd imply a + c > b + J. It did, in twenty-nine ofthirty d1SeE.
33. Tire exerciating care in 1heu'tectwiques is remarkable, D11tJg: sl1bjn;:tive probanbil- iity and utility wa<cmwdered me crucial empmcal question in decision making 1il the 19:501. (aft-er Ramsey showed how to do1t in principle 1i 1931) 1111dno effoo w1s ;pnred 10 Jo if properly. All choi.- - operati.maliud with random de10e&; Ubjets1 played oome gam- hles, f1t' f1111Mmtial stakes; instructions and 1111r of 111" data 11C pubu11td 111 1heir mono- graph; subjects were run in up to thret:repeated r1ssioru; ubjects u:nu1y played gambles a: me end, etc. Their insttuctrom relld (52): "Tiu:scJic:ehav-' b u n made especially for 1b, and they an, as fair as dice can be. la fact, they h been ground 10 i1:pecifications 11CU11m to 110,000th of an incl.1n
34. This design another=mp1aint about M'11St1der and Ngee. 1hat utility of money W11 confounded w1th utilit; of gwnbling b1au<1e subjecti: compared trrtain 1ums (no 1.111lty of gamM111g) with g:unbkt. (utility of gaubling).. In 1he Davidstoneul. design 11ley oornpared two bets, rather than stming-certaitty-aiuivalents, S< utility Qf gl11111bfog was held 1ixed.
- 35- This approach is furu1amel1ally different than that taken by M1china (t985) 11rd Crawford (1988) who 800w that with qu11f1-oonca'ie preferences peopleprefr'r to 111td<mit1e m choosing between a 111t; gamble;., generating die appeamoce of 11oc1&stic choice.
36. The dat. are also crmsiste11t with their regret-b11Sed theory (se., rection 11.H2 later) which 1ssues5 a-different lo gic than the altn1nte-compwison process described here. However, the original Tversky (1969) data are 1rm:omistenl with regret aversion. Furthermore, those dobi and the Loomes 111ld Sugden cycleS can both be eJJplained by an additive-difference model di cu.ed by Tversky.
37. 11t1S behavior is consistent with a nw1tiattribute uti1i model in which eath 111tnhute has a separable utility, tu1d the attribute utilities are added to derermine the alkrnalve's overall utility. But cliruinat11g high-rem apartments enlitley implies that the d15y1J1ty from real above 50111111 threshold is negatively infinite,
- J5. fut YOU1id: 11md in bot water and your right hand in oold water for a minnte. Bolh will adapt ro the tempta;tures of the water in which they sit. Then put both hands in the sarm 11h of wum J1111r. Siner, your bands are sensitive to dumges from the temperature 1;-vol to whk:h.

39. For example, subject15 are roughly $md:iff:ts:it$ between getting nothing_z $lild$ accepting a coin flip - \$JO md \$Xwhen X $lild$ around 25 (Tvm:k) and Kalmemfl 1991,1.

40. The breadtrtr of Mar:hina's technical cotribution rln:ith be questioned <n empiri:,al gtoood... His t m - f-J/emple, shJwed Ural if pOOkle *inegrae* asseL Md $lild$ the "globally" $lild$ a,;erse-duitcach of the "local $lild$ illy furu::Mini" use< w vlllue different gamble exhlbrt risk n-Ui:m many tions of EU $lild$ ddes as.urning a single m k an:rw $lild$ lly function would lwld. But the u.c/l:ulness of his rem:lOOlbk proof is underan by empirical evideffla: that (C_S_ unlty functions an: nm unifomly m<-l:averse O risk seeking ts:ce: epe<:ially see: li00S IDD and $lild$.H later); m tlv; u- pan of his "if,tlit(- proof appear; q1lervon&ble.

41. The implication of weakrning other p<:suru; have bffln lthoroughly worked oot 100. Au-UMP (1962) $lild$ owtd that weakening oompicteneis to a;yclkity of preferen= ($X > Y$, $Y > Z$ $lild$ $lie5Z$ $;$ X- /: $;$ not ferrtrrd to X) yickls a unidirectional p.uuat oeder in which u(X) > ;:(l') if $X > Y$!büt ii(X) > .(Y) does not alw imply $X > Y$; because prtfrences $;$ an be inoompklit) Hausner (1954) and Chiprtlan (19601 showffl that weakening coolinnity yields a ve:ror ntihry t:represent:l/n over whi<h preferences att l!i:O)aphic, mther inlm a wigle real"v"l1md flmct:oo.

42. Muchina /1982) $lild$ etined fanning Odl formally as followil: S u w - a gwnble X tocbasticaHy domma $;$ (Then Xlile lu the northweffl of Y-bigher P_n)om:rp ,,- in a triangle diagflin l lndifferen,-e curve& fan out if the slope of their tangent hll- at X u gtt:lW Ulan me slope at /' Fanning in the opp:e: too tUpe.u X is smaller itraa III /;

43. In general, welgbtng funcooni h&:fl two lqurtant femes, the location of the $;$ u>:>Ver point where $M1''' p$, and the gr< of CW"ature lbe <me p,lJLUCter form ln equaoon (S) <K&S not allow t,om teatureli to mmie indepew.iePty; the N<:>MOU falls as c w v - berotne! greater. Add!lS a woond parametrd dewuvs thlic t-o graphical propenies.

44. The R-O comparivn;ices nQt quilte tsfl farming out, because all We gambln lymg along the O-pair chord are not lltocla3tically dominated b; the R-chooi gambleffl; Thum, onr :an con< &fl't lt m llllm:l:al pktem of JOCH unlty functions that flm iNt and gMefflt, the observed data. This is a good emple of how a small dffleredWd in the de/sgo of Pllirs (pulling the Ptt in rllle riskier o-pl'lrgamble down bel;w the PH in the less-0.Sky H-piur gamble) mak&5 a big dffiml:m;:e betVeen a d w i te;t of fanning out Mid an appro::i(lat, Wit.

45. Many of the judgment Cperiments u,ed between-ubj<:,t deisign becaute' they' search bx moon, ie; in jllldgtMlU1 in two settings. The optimal choice of defign $;$ kp,nds on 'he needd for statistical powil: (wtthln ll more powerful thm between), the nbJ:Cis' toleranre fr repeated lflsh [high tolerance permib $;$ Wml, ind 'he ern:in from overstating.rtn<l undentat Ug: ;:onsisten-j' ntlmve to natural Mtmga /the conventional view is 'hat within ovmtflle:, betwun understates, w t f <4 rdl: know of much cvWence on precil:,ely th' point).

46. Earlier 'lildies (e.g., M<O)lcllef ind Nogue 1951) gadiered &uth data by giving ubjt: ts lldenti:al cilicCts many times. The danger in that making so many (chooe:l) may ov:,rwhelm subje,"U and mltY mduu them to use simple role;t; that m;:sqw;:&oo as conformrty to thwiy 'al. tlala in Slovic, Llc'lrtt:nswin, wd Edwatds {1965} sug p).

47. Wblthert t:lpcrnw have subject/ play gambles $;$ not oorirely a product of background or modem rlllS'n c=ioosnt'l'i. Siegel, Coomb . Edwards, Slovic, Tversley, and oilier psych:,ol&:ts often had illh_oon; play gilmbllts in lhl- 1950, llllJ 1960,,. Some felt lhat playing gambkij madl: littk dfferem:e and quit dmng it; othn think playing gam is important.

48. Neilson (1989) wggests- a c(Ner alre'lNidve UeQry in which $lild$ itl:n rise when gamblel hllv! fewer powble out:Ome&, His lhwfy jvids. nonlinear weiShing of ptobabltities entirely. This theoy violates cootinwty. and cannot e:,pijin all the,; dlua, but diser>K further elqlloration.

49. Theit modd use, j Taylor sllic5 $exp:lDfO$ n t lm/!l.f.expected utility by 5<Nerai terms

50. Subjects were ;down an illitit gamble on ettht (the left or rWt triangle edge then b d to name a gamble On the hyp:xenme, wherep...""0, which was tqUivatrn (in preferenc!':J.e. a point m lht w n e rndiffetM(e curve). Subjects then namud a n c=i v e m k r i o f) O i n t s t h i l l . w e n - u a H y p e r f e r r e d u r l t l t h r l e m f i v e p o i n t s o n I s i n g J e i n d i f f e r e n c e C J W e w e r e f o u n d . (S I d k i n o t p l a y a n i g a m b l e J)

51. Real (1991) ;tmlied to behavior of bees "isiting" urtifidial flowert widt differmt probl- btdil di tribution, of fttCtw: {controlti>i by the experim,mler). He disceoverl that b e e a p p l e d " r o W d o o : s n m a t e J W p r o b a b i l i t i e s o f g e t t i n g r e w i n l e i . " w i t h n e c t a r , a n d o v e r e s t i m a t e h i g h p r o b a b i l i t i e s , w h i c h h u r t s a t m i m p o r t a n t c r o s - d i f f e r t l : t e i n p m b a b i l i t y p e r c e p t i o n b e t w e e n a n d p e o p l e .

52. As the twJ probabilities fuU (m;limainmga <mBtant ratio), iflllllllling out is O'W people will 1:cume more TIS seeking and choose value. of A

53. Elislatl alluded to "A large number of niws.. u d e t a b s l u t e l y i a l c o o d i - E o m " w g g e s t i n g a m b i g u i t y a v = i o n i s t h e m a j o r i t y ! o f - c h o i c e .

54. The O D x c r e d w a s c h n l e n b y a o o i n f l i p , a f t e r " t h e g r o u p w a s I s k k d i f t h e r e w a s a o y o b j < O O I l t o f l i p p i n g a c o i n t o d e t e r m i n t t h e w i n n i n g c l W f f u r t h e g a m e t o b e p l a y t d " (6 7) . I n t h e r e x p e r i m e n t s w h o j e c t s c o u l d < w o s e a c o l o r t o b e t o o , o c w e r e M k e d t h e i r c e r t a i n t y e q u i v a l e n t s f o r b o t h o n b o t h t . :) m ;

5;t A t h i r d p t f o c i p . I t i s < m t e M i n v a r i a n c e : p r e f = f o r a o o h j l : t n o t d e p e n d (U t t h e s e t o f c l o i c e s f r o m w h i c h u C M b e p i c k e d . T h e r e a r e s o m e i l l e r t e : l f u f : Y l > l a t i o n s o f t h t ! p r i n c i p l n o o V h u b e r , P a y n e , a n d P u t o 1 9 8 2 : T v e r s k y a n d S i m o n w n 1 9 9 3) . T h e l o t t e r y c o m p l a t i o n I l l i t a t P h s e n - e d i n H l l d k e t o r e g r e t , d m r i b e d i n < i t t i o n H L H 2 l a t e r , a f i : I l l O t h a - e u . m . p k , b e c a u s e t h e y s n o w t h i l l p r e f e r e n c e U T a g a m b l e i n a p a i r w i s e d e p e n d s i n u s u s : m - a t t : w a y o n t h e g a m b l e i f i t i s p a i r e d w i t h s

56. In i l l i l l y c h a n c e e x p e r i m e n t s , t h e c u r w h l t i o n o f o u t c o m e s i , n o t e q u i l d ! l y f u O O . R e g r e t m a y c a u s e E U v i o l a t i o n i f h l v e p a r t i c u l a r < m t : i a o n s i n m i n d .

Baruilio, Kagel, and Jiumyakul (1990), Harless (!Wu}. Camerer (1989a, 19972a). Stamter and Sugdun {1987b) ulik di playi mat rontruUei fur regret effects and still found EU viola- tions; Tversky and KahnMman (1992) controlro.t for regret and fmmid SEU ViolatiMS {using bets on I l l l W l a l e v e n I l l) f l o w - e v e r , i t h l m t t h a t t h e m i n i t s t r i n g : E U i o i a t i l l J S A p p e l T W I l l S U R k e s t h a t d i d n o t c o o r o l f o r t r g e t f (e . g . , P r e l t { 1 9 9 0 , a n d C o n t s k 1 9 3 9) . L o o m t s M d S u g d e n (1 9 8 7 b) , L O O U s { 1 9 8 & ! , a n d S t a r m e r a n d S u g d e n { 1 9 8 9 a) f o o m t W J R g r e l e f f < u l I l l C I l l l l l f o r p u r t o f l h t , c o m m o o r a t i o (l l d o m m o o e o : n e e f f i r , - t , b u t n o t a l l .

57. Notice that violating reduction in mndom-garuble SCUing5 by isolating eadt pair :<cmo- miw; on mootil effort: ro obey tedeuction regllmr, mculplying pmbooilities and dloo,ing & portfolio of gambles to e t h e : < b (i c e o f l l l e , t r e l l < !) : < Q m p l < O l l p o u n d l m t e r y .

58. Si@larly, I l l . e n r i c h e d " p e r m r u i l i k e t h l e m e r w o o t M a d o n n a i s l i k e l y t o h t b o t h m a r r i e d f r e q H e n d y { d , o s e n) a n d d i f r e q u e n t (r e j ; a b u s e & ! w i t h m a n y U T f W l a l f e a t u r e : i r n i g h t l e b e l i k e l y c i n d l a t d e f u r b o t h l i e H l S i t i o n a n l d i v { t t m t .

59. T v e N k y e t a l . a l s o ! M d a p r o o , e o o n : t o I I V < t l e r i t i k m o f B D M , l i k e C o ; a n d E p s (d n ' t , p r o c e : L u r e , i e w l o o h s u t ; j e t s s k a l e d t h a t w e r e n l t d o o z y t o r a n k h e w . A l l d i f f e r e n c e i s t h r u A l b j e c t s p r i c e d h b e t i n a p a i r S C p a r a t t y { r a t t w U m l l y , / 5 i n C o l a n t i B p s . c i n , w h i c h t o m a k e p r i c i n g m r w h l i k e c o o l d . T Y f m l k y e t a l . a j , o i m n - e l a d r a m a t i c p a r e o f r e v e N a s i n u n o v e l i n c r e t e r n p o m d h o i C l ' c o o l e x t . I ' e o f l e m u c h I l ' (l t i m p a t i e r u w h e n d l o m : i n g i t r a a w h e n p r i c i n g : f t e c , a m p l e , 5 7 p t . n : t n t n f 8 U l l o j e m p r e f e r r e d \$ 1 . 6 0 0 i n e i g h t e e n m o r n h s t o S i , \$ 0 0 i n f i v e y - c m , b u t o n l y 1 2 p e r c e m p n c e d t h e \$ 1 . 6 0 0 b o n d l n g h e l -

60. I l k h m s e x p e r i m I M s t r i k e s I M a s a t x m u p l e o f o o . I t t h e b e s t a n d m t w o l m o f e x p e r i m e n t a l e c o n w u c m e t h o d s b r o u g h t t o b e a t o n a y - i n s p i r e d p h e n o m e n o n . T h e e c o o n .

- mists' intuition that people might behave differently if they are highly motivated drove his extraordinary design, creating a test of robustness to incentives that few psychologists would ever conduct. At the same time, his odd choice of objects (Volvo and Opel) violates the basic recipe for producing reversals that has been well-known since the early 1970s: that is, have people choose between two objects that are high and low on opposite attributes, then "price" both objects (or adjust one to indifference) using a rating scale that matches one of the two attributes. A poor choice of objects meant the ingenuity of his design was largely wasted.
61. A lower value of r implies less error in choosing, so more P-bettors will actually choose the P-bet. Since reversals occur when P-bettors price the S-bet higher, an increased number of P-bettors can raise the overall reversal rate.
 62. Such valuations are often called "contingent valuations" because they are designed to measure how consumers value goods as if-or contingent on the assumption that there is a market for the good.
 63. When means are used, WTP = \$3.45 and WTA = \$4.71 in the final auction; when medians are used, WTP = \$1.33 and WTA = \$3.49.
 64. In their auction, bids were not stated directly. Instead, buyers (sellers) were asked whether they would prefer to buy (sell) a mug or not at each price between 0 and \$10 (in \$.25 increments). Their responses were used to construct demand and supply curves; their intersection determined a market price and who would trade. Thus, subjects' bids were constructed from choices between the good and a series of potential market prices. (In psychological terms, bidding is a choice task rather than a "production" or valuation task.) The procedure seems to make bidding one's true valuation more transparent; conversely, it makes strategic underbidding to affect prices more opaque. (Strategic underbidding can be optimal because the risk of failing to buy at a profitable price, because of underbidding, is offset by the potential gain from being the price-setter and lowering the price.)
 65. As discussed in a section III.D methodological digression, the within-subjects design may overstate subjects' consistency (compared to the degree observed in between-subjects designs) if subjects think consistency is demanded of them, or if the equivalence of buying and selling is more transparent and if the buying and selling tasks are conducted right after one another.
 66. Hanemann (1991) points out that the wealth effect can be large if income elasticity is high and the cross-elasticity of other goods with the good being valued is low. These conditions might apply to some environmental goods (such as glorious beachfront property or mountaintop views) but they only apply in experiments with more mundane goods such as mugs under the absurd presumptions that (1) there are no good substitutes for mugs and (2) sellers prefer to spend most of their mug-wealth on mugs. (The condition (2) is empirically indistinguishable from the endowment effect, and can be considered a formal restatement of it.)
 67. The reference-dependent model also explains the finding that cross-price elasticities appear to be asymmetric: high-quality brands gain more market share from low-quality brands by cutting prices than vice versa (because consumers do not want to give up the high quality 'in their endowment').
 68. Kahneman, Knetsch, and Thaler (1990) suggest that endowment effects are, at particularly common in markets where goods are specifically bought for exchange, but Rietz's results suggest otherwise.
 69. Call the optimal reservation price in period t P_t . If the observed price P' is rejected one period, then $P' < P_t$. Since the reservation price should be P in all periods, if $P' < P$, in one period it is less than P , in all periods (since $P_t = P$) and should never be recalled.
 70. One motive for eliciting reservation wages directly is that the percentage of searches ending at the right time is a weak test of whether behavior is optimal. Cox and Oaxaca (1990) calculate that a naive subject who used the mean of the offer distribution as a reservation wage would end searches optimally 75 percent of the time. Compared to this benchmark, the fact that 80 percent of searches ended at the right time is unimpressive.
 71. I am always surprised that economists do not eagerly or routinely turn to the psychological models as viable alternatives, or at least to psychological facts for theoretical inspiration. Of

course, the psychological ideas usually do not come packaged as economists need them. Psychological models are often parameter-heavy or expressed verbally, since models are not usually constrained by severe analytical or econometric demands. Refitting the models is a chore economists must do for themselves. I do not see how psychologists can offer more help than they have already, in papers such as Tversky and Kahneman (1991).

Bibliography

- Abdelloui, M., and Bertrand Munier. 1992. Experimental estimation of indifference curves in the Marschak-Machina triangle: Can decision models be ranked or risk structure related? Working paper. GRID, ENS de Cachan.
- Abelson, Robert P., and A. Levi. 1985. Decision making and decision theory. In *Handbook of social psychology*, 3rd ed., G. Lindzey and E. Aronson, editors, New York: Random House.
- Ainslie, G. 1975. Specious reward: A behavioral theory of impulsiveness and impulse control. *Psychological Bulletin* 82:463-509.
- Ajzen, I. 1977. Intuitive theories of events and the effects of base-rate information on prediction. *Journal of Personality and Social Psychology* 35:303-14.
- Akerlof, G., and J. Yellen. 1985. Can small deviations from rationality make significant differences in economic equilibria? *American Economic Review* 75:708-20.
- Allais, M. 1953. Le comportement de l'homme rationnel devant le risque, critique des postulats et axiomes de l'école américaine. *Econometrica* 21:503-46.
- . 1979. The so-called Allais paradox and rational decisions under uncertainty. In *The Expected Utility Hypothesis and the Allais Paradox*, M. Allais and O. Hagen, editors, Dordrecht: Reidel.
- Allais, M., and Hagen, O. E. editors. 1979. *The expected utility hypothesis and the Allais Paradox*. Dordrecht: Reidel.
- Allen, F. 1987. Discovering personal probabilities when utility functions are unknown. *Management Science* 33:542-4.
- Alloy, L. B., and L. Y. Abramson. 1979. Judgment of contingency in depressed and non-depressed students: Sadder but wiser? *Journal of Experimental Psychology: General* 108:441-85.
- Alpert, M., and H. Raiffa. 1982. A progress report on the training of probability assessors. In *Judgment Under Uncertainty: Heuristics and Biases*, D. Kahneman, A. Slovic, and A. Tversky, editors, Cambridge: Cambridge University Press. 294-305.
- Anderson, M. J., and S. Sunder. 1988. Professional traders as intuitive Bayesians. Working paper 88-89-51. Carnegie-Mellon University.
- Anscombe, F. J., and R. Aumann. 1963. A definition of subjective probability. *Annals of Mathematical Statistics* 34:199-205.
- Antle, R., and A. Smith. 1986. An empirical investigation of the relative performance evaluation of corporate executives. *Journal of Accounting Research* 24:1-39.
- Argote, L., R. Devadas, and N. Melone. 1990. The base-rate fallacy: Contrasting processes and outcomes of group and individual judgment. *Organizational Behavior and Human Decision Processes* 46:296-310.
- Argote, L., M. A. Seabright, and L. Dyer. 1986. Individual versus group use of base-rate and individuating information. *Organizational Behavior and Human Decision Processes* 38:65-75.
- Arkes, H., and Hammond, K. R. 1987. *Judgment and decision making: An interdisciplinary reader*. Cambridge: Cambridge University Press.
- Atneave, F. 1953. Psychological probability as a function of experienced frequency. *Journal of Experimental Psychology* 46:81-6.
- Aumann, R. J. 1962. Utility theory without the completeness axiom. *Econometrica* 30:445-62; 32 (1964): 210-12.

- Ausubel, L. M. 1991. The failure of competition in the credit card market. *American Economic Review* 81:50-81.
- Babcock, Linda, George Loewenstein, Samuel Issacharoff, and Colin Camerer. In press. Biased judgments of fairness in bargaining. *American Economic Review*.
- Bar-Hillel, M. 1973. On the subjective probability of compound events. *Organizational Behavior and Human Performance* 9:396-406.
- . 1980. The base-rate fallacy in probability judgments. *Acta Psychologica* 44:211-33.
- Bar-Hillel, M., and W. A. Wagenaar. In press. Perceptions of randomness. *Advances in Applied Mathematics*.
- Baron, Jonathan. In press. Nonconsequentialist decisions. *Behavioral and Brain Sciences*.
- Baron, J., and J. C. Hershey. 1988. Outcome bias in decision evaluation. *Journal of Personality and Social Psychology* 54:569-79.
- Barro, R., and S. Fischer. 1976. Recent developments in monetary theory. *Journal of Monetary Economics* 2:13-76.
- Battalio, R. C., J. H. Kagel, and K. Jiranyakul. 1990. Testing between alternative models of choice under uncertainty: Some initial results. *Journal of Risk and Uncertainty* 3:25-50.
- Battalio, R. C., J. H. Kagel, and D. N. MacDonald. 1985. Animals' choices over uncertain outcomes: Some initial experimental evidence. *American Economic Review* 75:597-613.
- Bazennan, M. H. 1990. *Judgment in managerial decision making*. New York: Wiley.
- Bazennan, M., G. F. Loewenstein, and S. Blount White. 1992. Reversals of preference in allocation decisions: Judging an alternative versus choosing among alternatives. *Administrative Science Quarterly* 37:220-40.
- Beach, L. R., and L. D. Phillips. 1967. Subjective probabilities inferred from estimates and bets. *Journal of Experimental Psychology* 75:354-9.
- Becker, G. M., M. H. DeGroot, and J. Marschak. 1963. An experimental study of some stochastic models for wagers. *Behavioral Science* 8:41-55.
- . 1964. Measuring utility by a single-response sequential method. *Behavioral Science* 9:226-232.
- Becker, J. L. and R. Sarin. 1987. Gamble dependent utility. *Management Science* 33:1367-82.
- Becker, J. L., and R. K. Sarin. 1990. Economics of ambiguity in probability. Working paper. UCLA Graduate School of Management.
- Becker, S. W., and F. O. Brownson. 1964. What price ambiguity? Or the role of ambiguity in decision-making. *Journal of Political Economy* 72:62-73.
- Belt, D. E. 1982. Regret in decision making under uncertainty. *Operations Research* 30:961-81.
- Bell, David E., Howard Raiffa, and Amos Tversky. 1988. *Decision making: Descriptive, normative, and prescriptive interactions*. Cambridge: Cambridge University Press.
- Benzion, U., A. Rapoport, and J. Yagil. 1989. Discount rates inferred from decisions: All experimental study. *Management Science* 35:270-84.
- Berg, J. E. 1995. The impact of controllability and information on the use of public information in contracting: An experimental investigation. Working paper. University of Iowa, College of Business.
- Berg, J. E., L. Daley, L. Dickhaut, and J. O'Brien. 1985. Preference reversal and arbitrage. In *Research in experimental economics*, vol. 3, V. L. Smith, editor, Greenwich, Conn.: JAI Press. 31-72.
- . 1986. Controlling preferences for gambles on units of experimental exchange. *Quarterly Journal of Economics* 101:281-306.
- Berg, J. E., and J. W. Dickhaut. 1990. Preference reversals: Incentives do matter. Working paper. University of Minnesota Department of Accounting, November.
- Berg, J. E., J. W. Dickhaut, and C. Kanodia. In press. The role of information asymmetry in escalation phenomena: Experimental evidence. *Journal of Economic Behavior and Organization*.
- Berg, J. E., J. W. Dickhaut, and K. McCabe. 1992. Risk preference instability across institutions: A dilemma. Working paper. University of Minnesota, Department of Accounting.
- Bernasconi, Michele. 1992. Different frames for the independence axiom: An experimental

- investigation in individual decision making under risk. *Journal of Risk and Uncertainty* 5, 159-74.
- Bernasconi, M. 1994. Nonlinear preferences and two-stage lotteries: Theories and evidence. *Economic Journal* 104:54-70.
- Bernasconi, M., and G. Loomes. 1992. Failures of the reduction principle in an Ellsberg-type problem. *Theory and Decision* 32:77-100.
- Bernoulli, Daniel. 1738. Specimen theoriae novae de mensura sortis. *Commentarii Academiae Scientiarum Imperialis Petropolitanae* (for 1730 and 1731) 5:175-92.
- Bewley, T. F. 1986. Knightian decision theory: Pan L. Cowles Foundation discussion paper no. 807, New Haven, Conn.
- Blank, R. M. 1991. The effects of double-blind versus single-blind reviewing: Experimental evidence from American Economic Review. *American Economic Review* 81(5): 1041-67.
- Bohm, P. 1990. Behavior under uncertainty without preference reversal: A field experiment. University of Stockholm, Department of Economics.
- Billie, F. 1988. Learning to make good predictions in time series. In *Bounded rational behavior in games and markets*, R. Tietz, W. Albers and R. Sellen, editors, Berlin: Springer-Verlag.
- Bostic, R., R. J. Herrnstein, and R. D. Luce. 1990. The effect on the preference-reversal phenomenon of using choice indifference. *Journal of Economic Behavior and Organization* 13:193-212.
- Bowman, David, Debby Minehart, and Matthew Rabin. 1993. Modeling loss aversion: Some general issues, and an application to a savings model. Working paper. University of California at Berkeley, Department of Economics.
- Boyce, Rebecca R., Thomas C. Brown, Gary H. McClelland, George L. Peterson, and William D. Schulze. 1992. An experimental examination of intrinsic values as a source for the WTA-WTP disparity. *American Economic Review* 82:1366-73.
- Braun, Philip A., and Ilan Yaniv. 1992. A case study of expert judgment: Economists' probabilities versus base-rate model forecasts. *Journal of Behavioral Decision Making* 5:217-31.
- Braunstein, Y. M., and A. Schotter. 1982. Labor market search: An experimental study. *Economic Inquiry* 20: 133-44.
- Brehmer, B. 1980. In one word: Not from experience. *Acta Psychologica* 45:223-41.
- Brier, G. W. 1950. Verification of forecasts expressed in terms of probability. *Monthly Weather Review* 78:1-3.
- Brothers, Alan. 1990. An empirical investigation of some properties that are relevant to generalized expected-utility theory. Unpublished doctoral dissertation. University of California at Irvine.
- Brown, J. N., and R. W. Rosenthal. 1990. Testing the minimax hypothesis: A re-examination of O'Neill's game experiment. *Econometrica* 58(5):1065-81.
- Brown, William O., and Raymond D. Sauer. 1993. Does the basketball market believe in the hot hand? Comment. *American Economic Review* 83 (5):1377-86.
- Bruner, J. S., L. Postman, and J. Rodrigues. 1951. Expectations and the perception of color. *American Journal of Psychology* 64:216-27.
- Busche, K., and C. D. Hall. 1988. An exception to the risk preference anomaly. *Journal of Business* 61:337-46.
- Camerer, C. F. 1987. Do biases in probability judgment matter in markets? Experimental evidence. *American Economic Review* 77:981-97.
- . 1989a. An experimental test of several generalized utility theories. *Journal of Risk and Uncertainty* 2:61-104.
- . 1989b. Does the basketball market believe in the 'hot hand'? *American Economic Review* 79:1257-61.
- . 1990. Do judgmental correct biases in probability judgment? Evidence from judgmental experiments. In *Advances in behavioral economics*, vol. 2, L. Green and J. Kagel, editors, Greenwich, Conn.: JAI Press. 126-72.
- . 1992a. Recent tests of generalized utility theories. In *Utility theories: Measurement and applications*, W. Edwards, editor, Cambridge: Cambridge University Press.

- Camerer, C. P. 1997b. The rationality of preferences and choice in a market. *Organizational Behavior and Human Decision Processes* 51:237-71.
- Camerer, C. F., and T. R. Ho. 1994. Violations of the rationality axioms and nonlinearity in probability. *Journal of Economic Theory* 8:167-96.
- Camerer, C. F., and B. J. Ho. 1991. The process of performance in expert judgment: How can experts know so much and predict so badly? In *Toward a general theory of choice*, ed. R. H. Thaler, K. A. Ericsson, and I. Smith, editors, Cambridge: Cambridge University Press.
- Camerer, C. F., and R. M. K. 1994. Ambiguity and nonadditive probability in non-rational game experiments. In *Decision Experiments*, ed. R. H. Thaler, M. Macrina, and B. Munier, editors, Dordrecht: Kluwer Academic Publishers.
- Camerer, C. F., and H. C. Kunreuther. 1989. Experimental risk and uncertainty. *Journal of Risk and Uncertainty* 2:265-300.
- Camerer, C. F., W. L. Lichtenstein, and M. Weber. 1989. The curse of knowledge in economic settings: An experimental study. *Journal of Economic Theory* 97:123-54.
- Camerer, C. F., and M. W. Weber. 1992. Recent developments in modeling preferences: Uncertainty and ambiguity. *Journal of Risk and Uncertainty* 5:325-70.
- Camerer, C. F., and K. Weigelt. 1990. Bubbles and coexistence in experimental markets for stochastic-lived assets. Working paper 90-37-09-2. University of Pennsylvania Department of Decision Sciences.
- Carbone, Enrica, and John D. Hey. 1994. Discriminating between preference functions: A preliminary Moov Carlo study. *Journal of Risk and Uncertainty* 8:223-42.
- Carlson, J. A. 1967. The reliability of information in experimental market with a supply-response lag. *Swedish Economic Journal* 33:305-21.
- Carlow, Cynthia R., and Richard D. Johnson. 1992. Measuring the rate of time preference as a function of delay: An experimental study. Working paper, University of Alberta Department of Marketing and Economic Analysis.
- Casey, J. T. 1994. Buyer's pricing behavior for risky alternatives: Encoding processes and preference reversal. *Management Science* 40:730-49.
- Casey, J. T. 1991. Reversal of the preference reversal phenomenon. *Organizational Behavior and Human Decision Processes* 51:224-51.
- Caskey, L. 1985. Modeling the formation of price expectations: The Bayesian approach. *American Economic Review* 75:768-79.
- Cattell, N. J., Jr. 1977. Decision making with multiple probabilistic goals. In *Cognitive theory*, ed. N. J. Castellan, D. P. Osherson, and G. R. Potter, editors, Hillsdale, N. J.: Lawrence Erlbaum Associates.
- Chapanis, A. 1953. Random-number guessing behavior. *American Psychologist* 8:332.
- Chew, S. K. 1991. A generalization of the quasi-linear mean with applications to the measurement of income inequality and decision theories. *Journal of Economic Theory* 51:1063-92.
- Chew, S. K. 1985. Implicit wealth and mean-weighted utility theory, estimation, and demand revelation of the multi-unit auction: for an uncertainty model. Working paper 15, Johns Hopkins University Department of Political Economy, June.
- Chew, S. K. 1989. All the utility theories with two betweenness properties. *American Economic Review* 79:273-98.
- Chew, S. K., L. G. Epstein, and U. Segal. 1991. Mixture of simple and quadratic utility. *Journal of Mathematical Economics* 59:139-63.
- Chew, S. H., and K. R. MacCrimmon. 1979. Alpha choice theory: An axiomatization of expected utility. Working paper 669. University of British Columbia Faculty of Commerce.
- Chew, S. H. 1979b. Alpha choice theory, joint compound lotteries and the Allais paradox. Working Paper No. 666. University of British Columbia, Faculty of Commerce and Business Administration.
- Chew, S. R., and W. S. W. 1985. Empirical tests of weight-dependent utility theory. *Journal of Mathematical Psychology* 30:58-72.
- Chipman, J. S. 1960. The foundations of utility. *Econometrica* 28:193-224.
- Chiquet, G. 1955. Theory of choice. *Annales de l'Institut Henri Poincaré* 5:131-29.
- Chitlen, S., and J. J. 1991. The hindsight bias: A meta-analysis. *Organizational Behavior and Human Decision Processes* 46:147-68.
- Chiu, Y.-P., and R. L. Chiu. 1990. The subsistence of preference reversals in simplified market settings: A note. *American Economic Review* 80:1002-11.
- Cicchetti, Charles J., and Jeffrey A. Dubin. 1994. An econometric analysis of risk aversion and the decision to self-insure. *Journal of Political Economy* 102:169-16.
- Cluttfelter, Charles T., and Philip J. Cook. 1993. The "gambler's fallacy" in lottery play. *Management Science* 39:1521-25.
- Cohen, D., and L. L. Knetsch. 1990. Judicial choice and disparity between measures of economic values. Working paper, Simon Fraser University.
- Cohen, L. J. 1981. Can human information be experimentally demonstrated? *Behavioral and Brain Science* 4:317-31.
- Cohen, L. J., and Yves Jaffrey, and T. Said. 1985. Individual behavior under risk and under uncertainty: An experimental study. *Theory and Decision* 18:203-28.
- Cohen, M., and J.-Y. Jaffray. 1988. Preponderance of the reversal effect: probability distortion in decision making under risk. In *Risk, decision, and rationality*, B. R. Munier, editor, Dordrecht, Holland: D. Reidel.
- Combs, B., and P. Slovic. 1979. Causes of death: Biased newspaper coverage and biased judgments. *Journalism Quarterly* 56:837-43, 849.
- Conlisk, J. 1989. Three variants on the Allais surmise. *American Economic Review* 79:392-407.
- Connolly, T., and N. Gilani. 1982. Information search in judgment tasks: A regression model and some preliminary findings. *Organizational Behavior and Human Decision Processes* 30:330-50.
- Connolly, T., and B. K. Thom. 1987. Predecisional information acquisition: Effects of task variables on suboptimal search strategies. *Organizational Behavior and Human Decision Processes* 39:397-416.
- Coombs, C. H., and S. S. Komorita. 1958. Measuring utility of money through decisions. *American Journal of Psychology* 71:383-9.
- Copeland, T., and D. Friedman. 1987. The market value of information: Some experimental results. *Journal of Business* 65:241-66.
- Courtesy, D. L., J. L. Hovis, and W. D. Schulze. 1987. The disparity between willingness to accept and willingness to pay measures of value. *Quarterly Journal of Economics* 102:619-90.
- Cox, J. C., and S. Epstein. 1988. Preference reversals without the independence axiom. *American Economic Review* 79:408-26.
- Cox, J. C., and D. M. Grether. 1985. The preference reversal phenomenon: Response modes, market, and incentives. *Economic Theory*.
- Cox, J. C., and R. M. Isaac. 1986. Experimental economics and experimental psychology: Ever the twain shall meet? In *Economic psychology: Interactions in theory and application*, A. I. Maccoby and H. W. MacFadyen, editors, Amsterdam: North-Holland.
- Cox, J. C., and R. L. Oaxaca. 1989. Laboratory experiments with a finite-horizon job-ladder model. *Journal of Risk and Uncertainty* 2:301-30.
- Cox, J. C. 1990. Diminishing returns of the reservation wage hypothesis. Working paper, University of Alberta, Department of Economics.
- Cox, J. C. 1991. Tests for a reservation wage effect. In *Decision making under risk and uncertainty*, New methods and empirical findings, J. Geweke, editors, Dordrecht, Holland: Kluwer Academic Publishers.
- Cox, J. C., V. I. Smith, and I. M. Walker. 1983. A test that discriminates between models of the double first auction mechanism. *Journal of Economic Behavior and Organizational Theory* 4:W5-19.

- Crawford, V. P. 1988. Slo: hastk chice with qnasicorn: ave prefereflre fo11ction& W\lking paper SS-28. t:niveNity of Oilif<jmU. San Diego. Depmmtn1 cl':&:ooomi=
- Cwmnmg... ft G.. O. S. Bruoirahire. ind W D. Schulu, editors, 1966. vimlint eiwirol'menlaJ goob. Toti>"a. N.J.: lwwmao Wd Allimheld.
- Curley, S. P., and J. E. Yrte\$. J9&5. The anel" and rad8fl of l'he probability interval as factc,n aff ng ambiguii prefmmces. Organi::mi;mai Behaviorami Hwnan De cmym Pmam:1 j6;172:-87.
- Curley, S. P. an,j J. F. Ywes. J989. An empirical evahration of demiptn:e model of 11mbigwty reactions in chnce st1111fion Joiml4i of Mmhemmic;d Psydtdla_y 33:397-427.
- Clirley, S. P., J. F. Yate,, and R. A. Abram,, 1986. PsychologiCal sooces of i11nbiglity avui<J. aoce. Or-gwuwtkma/ &h.iVW, and Hum = Decimm Proce,,ses :!S:B0--56.
- Curran, L. and R. Slli11. J9S\I. Proelje\J Vet'us utility. Mrmn nr Sdence 35:22-4-1.
- Daniels, B. P., and C. R. Plotl. J9&8. [nfiation and eipctwioos in iuperi:nrtal m..rkd.. In &urirkd rarional behavior in experimenwl ga,MS mnd murkela, R. Tietz, W. Albers, and R. Se! 0!5, Berlin: Springer-Verlag.
- Dallielh, R. L and L. R. Kdkr. J990, An experimental evaluariarto ffile de:l:riptive validity of lottery-dependent utility tlloory. J,,urnal of Ril,k and Ujccrain,y 3:115-34.
- Davidwn. D., Suppe,, P., Hld Siegel, S. 1957. V«ition-making: M experim<muui appro=h. Stmford, Calif.; Stantod Unwersity Pr-ss.
- Dawes, R. M. 1971. A care ;:(ody of graduate admissions: Application of three primicipies of hmd1111 decismn making. Jmericim P.lychoiagisr 26: 180-8.
- ... 1988. Rational clwwe in ,m l.mctnain world. San Francisco: Harcoun, Brace, Jovimovich.
- ... 1990. The potential nonfalsity of the false consensus effect. In Ir:sights In d,,ri,m maJdng: A ni/>uie k Hillel J. Einhorn, R. M. Hogarth, editor, Chicago: University of Chicago Press.
- Dawes, R. M., Faust D., and P. E. Meehl. 1989. Clinical versus actuarial judgment. Scmm ce ?43:1668-74.
- de Fmetti, B. 1937. La prevision: Ses loh logiques, ses sources sources \$ubjectives. All/IQ/es de l'IrtMitur Ham- Poim:are 7:1-68. English ITTmslation in H. E. Kyburg, and K E. Smoker, editors (1964), Srudiei in subjecriv" probability. New York: Wiley Publisher. 91- J58.
- ... 1962. Does it make aene to speak of good probability appraisers? l' The "demist pi:lnlalti, I. J. Good. editor, New York: Basic Books. 357-64.
- de Finetti, B. 1977. Probabilites de probabilirie: Areal problem.or amismtt'ersumding?"In New din:c1wns in the appfcnlion lIf Baye\$/an nwthods. A. Aykae and C. BrwnM, editors, Amterdilllll: NOMI-Holland, 1-10.
- De Neufville, ft., al/AP. Delquie, l't8S. A model of the influence of certainty mnd prnbabliiy "effects" on l'he meawrement of lility. In Risk, decision, and ratWnality, B. R. Mumer, editor, Dordrecin, Hullimd: D. Reidel.
- Debreo, O. 19&8. Stochutic chm::e and erutinal wility. Econometrica. 26:440-4.
- Delrel, E. 198fi. An axiomatic ia1ioo of preftten,,7\$ under Wli:flilllly: Weakenig the ndenee IUKJlll. J01J.mal ofEc."1Wf1W; fM.(lrf 40:304-18.
- O<sw, J., and S. Ribein; dll C';sta Werlang. 1992. Unefitainty a\11rsion and the uptm,11 choke oi portfolio, ;:,i:a N I 197-204.
- Dillt R I L and S. Surwer. 1986. Incentiv«, learning, and pr=siu.gofincomwtion in a market environment: An e\aninatitn of the base rate fallacy. IDluoon,fr!:) marke: rvseanh. S. M(riarty, etii,ru. Mlman, Okla,, t:niversity of Oklahoma Pre.s.
- Dwyet, G. P. Jr., A. W. Wlliami, K C. &rullio. and T. 1 MIX(In. 1993. TP&S :if tational expeeunions in a alark lleting Eronmrik Jo<mal J03:556-61J.
- E,jwaM, W. 1951. Probability preferemeli in gambling. American Jmmral ilf Prydwlogy 66:34.
- ... 19546. Probability preferences among bets wllh dtt'fering expected vahies, Ameritar. Jmmwl fl{ f/ydwlcgy 67:56-67.
- ... 1954b, The reliabity of probaoility preft:i= A-rinm Jour:a/ af Psydwlogy 67:68-95.
- HL:Mc. The theory of decisiilll tnakmg. PSJNlt)Wgirul /Jufenn 51:381}4i7.
- ... 19611L Bth.a'l-ioral deciiioo theory Ann,ai Rel-UW e f Pzyd w#lgy 12:473-98.
- ... 1961b. Probability learning in IOQl triats. Mm Q' of Expe J P;y,y,hcio_y 62:385--9<
- ... 1%2i Snb;ective pmtabilitiei inftrred from deci'IMIs. Psydwirigmd Rn 1/2w 69: 109-135.
- ... 1968. CmwervilUtil'll in luk'TM inf-TT'm\ion procts,ing. In Fmm ai l'tph'se.taeim, of 1mmati jmigm.Mi, B. Kldmnuntz, editor. ew YOJt: Wiley. 17-52.
- Eger, C. and I. Okkhau.. 1982. An -xaminJKioo of rite coMervativ in!ormation proo,ming bia;; m an aceooming framcwtrl:. Joumai e f Atr<mnting Remm,i, 1.0:711-23,
- Elnlwm, H. l flno. 11ie use of nonlinear, noucompentatory models in rlb:lion--mllhng. Pn-i:ilwlgual Bulieriti 73:22J-30.
- Einhorn, H. J., Md R. M. Hoglrth. 1978, Confidence in Judgl'llffit: P=i tence ofltw illusion of valid.it; l'4yi:jw/ogical Review 85:395-416.
- ... 1935. Ambiguity M-O uriterdlinty in probabilir;tic infurenne. P,y.:/wlcglcflJ RN iew 92:433-61.
- EllsOOrg, D, 1954, Classic and Cltiffitt notioru. of "measumbie utitiyb Economic Journal M:528---56.
- ... 1961. Rilik, amhg111ty, mnd me Svnge axiQ-ms. QU<1urdl Jnunml <f E.,mgmic s 75:643-69.
- Epstein, L. G. In press. Behllviour under risk: Recent <flivclopments in thsory and applica.trom. Ill Advt11tus in ttimQmir 1Mory: Sixth lb>rid Ctmgre1s of rlie Econ,c,nie;m Society, J. J. LaffOifl, edit(,), Cambridge: Cambridge Univmlt' Pres;.
- Erev, ldo, Thomas WnlBten, and DnvKI V. Oul>9cu. 1992, Simultwemw overconfidence and ronserntism in judgment: Impkations for rcMarch and practic:c. Working paper. Uruvei. sity nf North Carolinn, Departm,mt nf P ychology.
- Evan, Dorin A. 1992. Tt&t; of expccted Utility theory ut the indMdual and market levels, Workitg paper. University of Alaharna at Hunnrville, Department of &:onomics and Finance.
- Fellntt, W. 1961, Dfaroti-011 of subjective pmbahiliie a 11reaction to U<f'ertnity. Quanerly Joumu/ of Economics 75:670-94.
- Fernandez, R., and D. Rodrik. 1991. Resiunnce to refumi: Status quo bits in the preserne of imlividual-specitk uncertainty. Am't'rcan &011nm1r Review &I: 1146-55.
- Ferris, S. P.; R. A. Haugen, M-O A K. Makhija. 1988. Predfoung conpor-d'l' volume 11-1h hiMO'ic volume at differential pri<f levefa: Evidence ,mppc,ning the dispooiion effect. Journal c f fi11a,uc 43:617-97.
- Fischer, O. W. 1982. Scoring-rule feedback Md the oven:onll.deocc syndrome in -ul>jeellve pmbattility tim:cWJtiog, OrgaM'alumal lkhi:wtcr and Humm, P,rfrtmamr 29:)52--69.
- Fiocllbott, B. 1975, Hindiighl 7 foresight: The effect of outcmne knowielg_e on judgment under ty. Journal of Experim'lta! Psydwlogy: Hu""-14 Percqitrm aml Pr,fennant L28S-99,
- ... 1982, Debiasing. In Jwlgmcrn ldtder11nctr1atmly:Himrlstitsand hiuse,, D. Kahoeman, P. Slovic, and A. TvefflY. editot\$, Cam: Cambridge Uni:veaily Press. 42'2,4J.
- ... 1983. Predicting frame, Jaurnai ofExperim.t11.talhu:i,cf('ff'Y; L<lmimg, Memwy, and Cogr11flm 9':103-11Ji
- l'il'clihoi, B., mnd R. lieyth. 1975, "i knew if would happen" h'm:d pmti;bltitm of 0111("JlUre!lungs, Owtgwt;atl;mal Bilrar<V and fl-.m Pnfe?fldffCtJ IJ: J. Hi.
- F'hclooff, 8,, P. Slovic, mnd S, Liebtenikill. J97K Flllllt m'CS: Sen.1tivity of mima'OO failur< probabilities kd problem n: . Jmmwi flM(R'f!/flfl/41Psyrlwlot;c HJtm(lf/ Fer-ctfWV/ amth,fi,l'liWl'e 4:330,4.
- Fisbbttrn, P. C 19&2, M,mtrans'l'ive men.Uffille miiily. Jmmwl P fMmhemtkat Pzydw{ogy 16:31-67,

- Filthum, P. C. 191B. Tnm,juve measurable utility. Journal af Ecmwmi; Tiicory 31: 293-J1?.
- 1984. SSB utility themy: An oonomit perpe-ti-cc- MarhImatic11 Sm:nl Sd,mce 8: 63-94.
- . 19Sa. N()fjiJNw preference and «ility flitary, Baltimore: John; Hopbns Press.
- . 1989. Generalint;.,n,i of expe«cd pultly tm-Jries. A.;urvey of re:znt proposals.)tmuls of()peratk11 Rnearrh 19:3-21.
- . 1990. lwtim-pective on die utility tlleory of voo Neumann and Morgenstern. Working pilpr. AT&T Bell Labs.
- Ffahbum, P. C., and O. A. Knchenbel:ger, 1979. TuQ--cece «Qll Neumann-Mo em utility function., Dnnfrm StItMe1 10:503-18.
- Pnbhum, P. C. and A. Rubiniteln. 1982. Time preference. lrrnernali&w/ Ecmwmi; Review 23:677-94.
- Forsythe, Robert, FOOSt Nebon, fJlorge R. Ne11fili11f, Md Jad.; Wright. 1992. Anatomy Qf a political swck ulllbt. Amfri;on Ecocoomk Review 82:1142-61.
- PrandosI, Robert, Pravcen KulIU, Roland Michelit:cll, and Vemcn Smilh. 1993. Experimental UUU ;f the emittwmenf effot: Wmkiog papcc, Umver;Uy of Ariwrn, partnem uf Ec-i-oomics, Mw:h.
- Frtidridm:m, Barbam L., and Daniel Kahneman. In press. Duration neglect in retroopootive evaluati-On of atk;tive epiOOcs. Jmmwl fj{ Pt sonaiity awi Social P.lyd10Jogy
- French, K. R., and J. M. Poterbl.. 1991. Inve ordiver ifi<:atiOll and international equity markc16. Amark111 Ei: = m k Review l'aper.1 and Procefdingl' 1 81:222..o
- FrieJruan, M. 1953. Mel'wdtJlogy of pc-sljve economics. Chicago: Univorslty ofOic:ago Press.
- Friedman, O. 19\$9. Tho S-shaped value function as acon\$U-ained optimum. American FN>Mmi; Review 79:1243-1.
- Frisch, D., and J. Daron. 1988. Ambiguity and rationabty. Journal of Brhavinral Decis/011 Making 1:149-57.
- Ganguly, A. R., J. H. Kagel, and D. V. Molief. In pres.. 'Ibe dfem of biases in probocility judgmentli in asset market; All clporimental euuninatinn of ln, ba'ie rate fallacy. /Qumw/ of Economic Behavir Md Organiwlio11.
- OltfdenOfli, P., and N.e. Sahlm. 19\$2. Umehabl flObJbiliti nsk taking, and ,J:lsioo mak-ing. Syntheu 53:361--86.
- Garner, A. C. 1982. Experimental evidence ClI rhc rationality of intuitive forecaster\$. In Ju. slanch in r,xperlmrmal eca,wmfrr, vol. 2, V. L. Smith, ONim, Greenwich, Conn.: JAI lflfs.
- Gertner, Robert. 199J. Gttme ioow Md economic behavior: Risk.laking on "Card Sharks." Quarterly Jvumal of Lt'conomfcs. 108:507-21.
- Gigerenur, G., W. HelL and H. Blank. 1981!. Pre000tatku and ontent; The us; ofba,e,ate;, a. a :ominuous wtlabk. Joum;1 afErperimemal Plydwiagy 14(3):51 25.
- G-i, Gctd, U. Hoffrage, and H. Klwlbdlting. 1991. Probabilistic nrenml models: A Brun&wikian theory -Of oonfideUCt: Pzyt;lw(qgknl Reviiwr 98:506--28,
- Gl\$liotti, G., and B. Sopher, l m . Te\$ilng alternative lruar representatiOM -Of behaviVT' omkr uncertainty: fxf>clted utility, weil;hkd linear utilit., tnuuitwe convex and nontraruive convex preferences. Rutgeni Univernt; Department of Econmnia, February.
- Gilboa. L 198f, Expected uulty with purely s11bjc;five n ive probabildi-s. Joi,mttl lIf Mathematical EdRWffWJ 16:65-88.
- Gilboa, L and D, &hmcidkr. 1989, MiUllin expected utility with a lmn-llnique prior. Jm;m,i{ ofbfoth4.mMkal EcQ/Wffies 18; 141-53.
- G'Uboo, ltlhak, and David t In pres;. Case-based di;.,n theory. Quan;rJy Jlxmla(of Econmk-s.
- Qilovfub, T. 1991. Hew we know whal liti w: Fallacitsofhw;n, rowning. Sew York; Free
- Gikwieh, J., R. ValJnnc, Md A. Tversky. Jely 1985. The bnthand in basketball. Oa the misp;.,, ceptkm (lfmndoru scqu Cng/Ulive P;r;dwk,K)t t7:295-J14.
- Gode, Dharan]<J K. and Shyam Sunder. 1993. Allocmve efficiency Qf markeH with l«Q» intelligence ltadors: Market a partial wbstituM for individllll nlonality. J l qf Po/ilic;/ &-mwmy 10:119--37.
- Green, J., ml B. JWtleL 1988. Ordinal indepcndee.e ill llOlllineat utility thoory. Journal o fRisk and Um:mail'iry 1:35:'i-87. (Erratum, 1989. 2, 119)-
- Greeriteeg, M. R., D. Q. Sachsnum, P. M. Sandman, and K. L. Salomon<1 1989. Network e;.,sing ;ews CQHnge of environme:ml risk. R/lt Analysis 9:119--26.
- Gregory. Robin, and Lita Flfbby. 1917_ Auctioos. experiments and cantingent valuation. Public Ckmce 55:273--89.
- Grether, D. M. 1980. Bayes- rule as a descriptive model: Tie representativeness heumtic. Quar-ferly Journal of Ecorwmu; 595:531-57.
- 1981. Financial incentivt cfects and indJvidual decision maki_n g. W<lking ptpCf no. 401. California Institute of Technology.
- 1990. Testing Raye& rule and Ue repre.l'<nmth-ene\$ heuristic: Some experimental evi-dence. Jmmwl af E-cmmmc Beho.vior and Organiw.tio11 17:31-57.
- Grether, D., and C. R. Ptott. 1971. E.ronomie theory of cholee and the preference revenal pheoomenon. Ameri;an Ecmwmc Review 69:623-38.
- Greiber, David M., Alan Schwanz, and Louis L. Wilde. 1988. Uncertainty and shopping behav-iolr: An experimental analyis. Review of Eummnic Studies 60:323--42,
- Grether, D., and L. L. Wilde. Match 1984. An analsis of conjunctive choice: Theory and experiments. Jouma/ of Cullsumtr 11.tuorch 10:373-85.
- Grice, H. P. 1971. Logic and convers&tioo. In The Wgic cf grammar. D. Davidson and G. Harman. editors, Encino, Calif.; Dickenson.
- Griffin, Dale. and Amos Tversky. 1992. The weighing of evidence and thl' dctennittanu of confidence. Cagmive Psychology 24:411-35.
- Griffith, R. M. 1949. Odds adjustment. by American horse-race betton; American J011n/41. of Psychology 62:290--4.
- Gui, F. 1991. A theory of diappointnment in decision making under lllCenainty. & u - e m c a 59:667--86.
- Guler, K., C. R. Prott, and Q. H. Vuong. 1987_ A tudy of iero-out anctions: Experiruemal analysi5 of a process of ailocaling priate rights m the u c of public property. WMing paper 650. Califumia Insb.tute of Technology, Division of Humanitie& and Social Sci:il:rues, August
- Hagen, O, 1991L Decisions under risk: A de-cl'riptive tOOdd and a llllique for deci.ion making. European Journal of?olirlrnl EtOMmy 1:3&1-405.
- Hagen, O. 1%9. Separation on cardiuJ utility and specific utility of mk m theory of cltoies undec uncertainty. S'Ettsokootrisk TUbjkr fl 3.
- HaltWlllger, J., and M. Waldmm. 1985. Rttiorutlt.tpeciilitnU and the limm of rutationality: An analysis of heterogeneity. Affll'rican &-orwmi: kirvi,Jw 75:326-40.
- Hammack J., and Bmw n, G. M., Jr. 1974. Waurfu..i /Jlil wet kmde-: Toward bio <ltMWmk Matysis. Baltimore: Johns. Hopkins l'ul'ersity Pren, fur Resounts> for the Future.
- Handa, J. 1917. Risk, probabilities and a new ltheory ilf carminat ilility. Journal of Pmitiroj Economy 85:97-122.
- flanemann, W. b.t June 1991. Willingues; to pay and willingnessi toa-xept How much can they differ? Amefican Ecorwmur Re-view 81:635--47.
- Hanlie. B. G. S.. E. J. Johnson. and P. S. Fader_ 1993. Modeling XAAa w,ion and reference dependeece effects oo brand chioice. Manetmg Sdmc-t 12:378-94,
- Harless. O, W. 1989. laboratory e,-Jdence on the disparity beiewett willingneU ID pay ;lid compensation demanaed. JollrMI of &mwmi: lkha.l'Wr wuj O, H: 359-79_
- 1992a. Action erill\$ The rlfca of problem representation Oll regret. Amrr-icdn &o-QllUC Kfflew 82:634-49.
- 1992h. Predictions about mdffereoc curves inside the unl - A test of variams of ro:pected OOfity theory. Jmumil of Economic rand Organisation 18.391-4)4.

- Kauodu, C., R. Bu.hman, and I. Dickhaut. 1989. Escalation and the sunk cost effect. *Annals of the New York Academy of Sciences* 580:19-17.
- Karmukar, U. D. 1977. The effect of probabilities on the subjective evaluation of lotteries. Working paper, MIT, Sloan School of Management.
- . 1978. Subjectively Weighted Utility: A descriptive model of elicited utility. *Organizational Behavior and Human Decision Processes* 21:6-72.
- Karol, E., and Z. Safra. 1987. Preference reversal and the observability of preferences by experimental methods. *Econometrica* 55:67-85.
- Kami, E., and O. Schmidt. 1990. Utility theory with uncertainty. In *Handbook of Experimental Economics*, ed. J. H. Kagel and A. Roth, 11-4. Westport, CT: Ablex.
- Keller, R. 1985a. The effects of problem framing on the choice between substitution and principal-agent. *Journal of Economic Theory* 35:38-51.
- . 1985b. Testing of the "Narrowness of compound alternatives" principle. *Omega* 13:49-58.
- Kelley, L. R., U. Segal, and T. Wang. 1993. The Becker-DeGroot-Miller method: A general utility theory. *Journal of Economic Theory* 61:1-11.
- Keren, G. 1987. Pacing uncertainty in the game (if bridge: A calibration study). *Organizational Behavior and Human Decision Processes* 41:98-114.
- . 1988. On the ability of humans to estimate probabilities: Some calibration studies. *Acta Psychologica* 67:95-119.
- Keynes, J. M. 1921. *A Treatise on Probability*. London: Macmillan.
- Knetsch, J. L. 1989. The endowment effect and evidence of nonreversible indifference curves. *American Economic Review* 79:277-84.
- Kolm, J. C., and J. A. Sinden. 1987. The persistence of evaluation disparity. *Quarterly Journal of Economics* 102:691-8.
- . 1984. Willingness to pay and compensation demanded: Experimental evidence of an unexpected disparity in measures of value. *Journal of Economic Theory* 31:317-21.
- Krueger, M., and V. L. Smith. 1987. Hypothetical valuation and preference reversals in the context of utility trading. In *Laborman, Labor, and the Quality of Work*, ed. J. C. Smith, 11-24. Cambridge: Cambridge University Press.
- Kurzban, R., V. L. Smith, and A. W. Williams. 1985. Individual rationality, individual irrationality, and value estimation. *American Economic Review* 75:197-202.
- Knight, F. H. 1921. *Risk, Uncertainty, and Profit*. Boston: Houghton Mifflin.
- Koehler, J. J. 1989. The nontransitivity of base rates in probabilistic judgment. Working paper, University of Texas at Austin, Department of Management.
- Kogut, C. A. 1990. Commercial search behavior and the role of the Journal of Economic Behavior and Organization. *Journal of Economic Behavior and Organization* 14:381-91.
- Krantz, D., R. O. Luce, P. Suppes, and A. Tversky. 1971. Foundations of Expected Utility Theory. *Journal of Mathematical Psychology* 16:1-25.
- Krout, Y., and H. Levy. 1992. Further tests of the separation of the capital and the pricing model. *American Economic Review* 82:664-70.
- Krout, Y., H. Levy, and A. Rapoport. 1988a. Experimental tests of the separation of the capital and the pricing model. *American Economic Review* 78:500-19.
- . 1988b. Experimental tests of the separation of the capital and the pricing model. *Organizational Behavior and Human Decision Processes* 42:388-410.
- Krout, Y., J. A. L. Fm, and D. R. L. Luce. 1990. Conventional conventions, order of information acquisition, and the effect of base rates and individual information on social judgment. *Journal of Experimental Psychology and Social Psychology* 59:1140-5.
- Kuhn, H. 1953. *The Game of Wits*. 2nd ed. Chicago: University of Chicago Press.
- Kunreuther, H., O. Emswiler, W. Drury, and P. Slovic. 1990. Public attitudes toward the siting of a high-level nuclear waste repository. *Risk Analysis* 10:1-10.
- Kunreuther, H., R. Gimpert, L. Miller, P. Sadler, P. Skwicz, B. Borlan, and M. Ketsch. 1978. *Disaster Insurance: Public Policy Issues*. New York: Wiley.
- Lachman, R., J. L. Lichtenstein, and M. C. Buerlied. 1979. Cognitive psychology and behavioral economics. *Journal of Experimental Psychology* 108:1-11.
- Lichtenstein, S., and P. Slovic. 1982. Calibration of probabilities: The state of the art to 1980. In *Judgment Under Uncertainty: Heuristics and Biases*, ed. D. Kahneman, A. Slovic, and A. Tversky, 306-334. Cambridge: Cambridge University Press.
- Lichtenstein, S., and P. Slovic. 1971. Reversals of preference between bids and choices in gambling decisions. *Journal of Experimental Psychology* 89:46-55.
- . 1973. Response-induced reversals of preference in gambling: An extended replication in Las Vegas. *Journal of Experimental Psychology* 101:16-20.
- Lichtenstein, S., P. Slovic, and D. Zink. 1969. Effect of instruction in expected value on optimality of gambling decision. *Journal of Experimental Psychology* 79:236-40.
- Lightman, A., and O. Gingerich. 1990. When do anomalies begin? *Science* 255:690-95.
- Lind, B., and C. Plott. 1991. The winner's curse: Experiments with buyers and with sellers. *American Economic Review* 81:335-46.
- Lindman, H. R. 1971. Inconsistent preferences among gamblers. *Journal of Experimental Psychology* 89:390-7.
- Lippman, S. A., and J. J. McCall. 1976. The economics of job search: A survey (I and II). *Economic Inquiry* 14:155-89; September 1976, 14:347-68.
- Loewenstein, G. 1987. Anticipation and the valuation of delayed consumption. *Economic Journal* 97:666-84.
- . 1988. Frames of mind in intertemporal choice. *Mathematical Economics* 16:200-14.
- Loewenstein, George, and Jon Elster. 1992. *Choice over time*. New York: Russell Sage Foundation Press.
- Loewenstein, G., and D. Kahneman. 1991. Explaining the endowment effect. Working paper, Carnegie Mellon University, Department of Social and Decision Sciences.
- Loewenstein, G., and D. Prelec. 1992. Anomalies in intertemporal choice: Evidence and an interpretation. *Quarterly Journal of Economics* 107:573-97.
- Loewenstein, G., S. Blown, and M. L. L. In press. An introduction to revealed preference for fairness. *Journal of Economic Behavior and Organization*.
- Loomes, G. 1988a. Further evidence of the impact of regret and disappointment in choice under uncertainty. *Econometrica* 55:47-62.
- . 1988b. When actions speak louder than words. *American Economic Review* 78:4(-)-20.
- . 1989. Predicted violations of the invariance principle in choice under uncertainty. *American Economic Review* 79:19.
- . 1991a. Evidence of a new foundation of the independence axiom. *Journal of Risk and Uncertainty* 4:91-108.

- Loomes, G. 1991bs The c- of prefertence!twenal; Cmmcm. l: m v = ty of York, C'mm: for E:qienm!mtal l;;oomm;\$.
- 199k Tmting d=sioo !heories by using "vlll!te equivale11CeS: Oxfi,rd &mwmic Papers 43:644-66.
..... In pre#. Pref re: Fxp!lination& evhlence and implications. In Jmro11Sirve prefert11ce (Am<als of ffilu<JN W7771 ,r,l 20), Gehrlein, edimr.
Loomes, G., C Starmr. MOO. R. Sugden. 1989. Pref;reoce reve-sak Informa!xm pmce&siog clfect or rational noouan!rnve choice? Economic J-ounwl 99:140-:il
--- 1991, Ob!ieMng liowiron m transitivily by erimental. melbooo. &o,,,merrica 59(2):425-39.
--- 1992. Are prefur=s lllhnotoni;? Testing rome predictions of regret lhemv. Ew-mintira 59:17-33.
- G., JMO R. Sugden, 1982. Regret theory: An alternative theory of faulonal choice urulu umertainty. Ei!Oflljmk Jnumal 92:805-25.
----- 1987a Some imp&aoons of a more g=c=a l fonnofregret ltheory. lul<l"MI of Economic Twory 41 :210-37,
----- 1987b, Teting for n;grrt wtd diwppointment in choice under uncervwnty. frmwmir Jcmmoi 97:118-29.
loolnfl>, Graham, and Caron Taylor. 1992. Non-transitive preferences over gains and losses. Ei'f'f1tm1irJoul'MI !02:357-65.
Lovell, M, C. 1986. Tests of the rational expccllltion hypothe;is. American Ec,;nmic Review 76:110-24,
Luca, R. E., Jr. 1986, Adaptive beavivc and ec,;nomk theory. Jourool of Business 59:S401-S426,
Lu,;e, R. b. 19!18. A probabili LIC thoory of ltililil', Econometrica 26:193-224.
--- 1959. lmiil'idual clwlu bfhavior: A r/urxeriCQ/nnoylis. New York: John Wiley and Son.
--- 1990, Rational versm; plausible accounting equivalences in preferencejudgments. Psy.. dwiogiuil Scit'!Ce li22li-34.
Luce, R D sn. Barbiua Mellen; an! ShHie Chang. 199J. fa choice the correct primitive?- On uring certainty equivahmts and reference levels to predict choices among gambles. Jou.mal qfRiskr:md Umt,mil1ty 6;115-43.
Luce. R. D., and l. S. Nwens. 1985, CIMSHleatioo of 1;0111ateration=urement structures au:ording to :cak type. Jrmmal ofMatMmatcal Psychology 29: 1-72.
Luce, R. O., and P. Suppes. 1965. Preference, udllly, and iubjective probability. In Halldbook. Q:fm1Jthemat;,,,jpych(>iCgy, of. flil, R, D, Lu , R. B. Su h, and E, Galutler, editors, New York: Wiley.149--410.
Luce, l Dunean and Detlof Viff Wiruerfeldt. 1994. What oolllllKlll ground exists for de.scriptive, prescriptive, and normal.live utility theQ:res? Mall<lltmt11 W n a 40:263-?9.
Jypny. Gregory J. 1991. An B-perirnenu: uooy of llllllllgerial pay and fum hedging da,isioos. Workin; plijXt. Coooc<tia Univefflity, Department af Finance.
MacCrimmo, K. R. 1965. An M/HTUMftal rtuij, af rM J,c1SWll miki11g behavior of busi>es, &11rUJitf. Unpublmed di terulion. University of CalifWTiill, Los AJ Jc,.
Ma,;Cnm,mou, K. it, and l. S. Lamon. 1979. Ltility theory: Axioms versm; parailous. In The ;cxpwe! urlijil' hyparhests and 1hf A,Jlais P,mtcx. M. Alla» ID.1 O. Hagen, (-0, Oodrodou, Holl/Ind;[]], Rie-del. J33,-409,
Mactxmald, D. N' J. H. Kagcl. ind R. C. Battalk,. i99L Animals' dm,;es O\$N uocenain 001C0me; Further experimental Fe/Ulti. Ee-Ooomic Journal JOJ:1067--34.
Ma,;D,;rmkd. Doo N,, Witlam L. & th, Md Pim! ML Taabe. 199J. GeoraltWJ expected urilny analv5il, and prefarene rev:rsals: Some mitia) R:illllS in the Jou iomUa. .kml7flll of Ecu-nom; Ddiavvrand tioo 11:115-311
Machim, M J. 1%2. "Ex utility" analyii, wilhwtd the independence axiom. Ecmwm;• trleuS0:277-313.
-- 1987. Clmi(e tmdr .m:ertainly: Problm1& solved and unsolved. Joumai Df Enmmamic Perspectives 1:iZl-54.
----- 1989. Comparative statics and non-expected llllility preferences. ItJurnal tJf. &-...n 1c Theory 47:393-40'L
----- 1985. S-tochastic droioe funcoons generatnd from drtmninistic twfim OYer lm-ooritl: fkmwmic Journal 95:515-94.
Malinvaud. E. 195:L Note Ott voo Nt<11nm11vMorgemteru's strong independence axiom. Econo-l'Mtri;a 20:679.
Madowitz, H. 1959. Portfebv : Ej/ciem diwirsificurion of mwsr;rm:rm. J';n, York: Wtley.
-- , 1952, The uliliry ofwealth..Jwmr.itqfPolitical Ecmwmy«i:151-"8,
Marinchak, J. 1968. B;m,mii:s of inquifmi, oommunkating, deciding. rit:PI Et nomic Review 51H-18.
--- 19 n Peoonal prob,bliries of probabilities. 11utury Qffl du-irtmt 6: 121-53.
--- 1950, Rational behavior, WIC.flwW prospeeu, and meaiurabl utility. &OIWI'Ulrrim 18.111-41.
M&rnhlill, l' D" l L. Knetsch, and J. A. Sillden. 1986. Agents' tvaluation Wd th, dl,parily in meas11r:s of econnmk kJ/S JOU11Q/ of EcJf/Wrtic Btfiavior 111d Orgmizatim 7:115-27,
M.u halt R. C.. J.F. Richan!, and G. A. Zhu-kin. 1992. Posterior probabilities of the ir,defen-dence ilXiotn with non-e:qieruaental dilill. JourMI of Bwine,u tJNJ E,;1W111C Statistk.t JO:J1--4,
May. K O. 1954. Intransitivity. utility, and the gregation of preferen-e pal1nms. &onome-lica 22:1-13.
May, R. S. 1986. Inference, subjective probability and frequency of cotreet an,wcr ; Aoonilive approach to the overconfidence pheimmeooo. In New &rrcrlous in research on decision lIWkmt, B. Brehmer. H. Jung1'11lllMO, & Lourens and G. Sew)'n, (OOI, AtnstefOllm: Nonb-fkllillllld,
McClelland, G., W. Scullure, and D. Coumey. 1991. The effect: of framing and the tanw quo an compell;Wng and equivalent •ariation mnmLres of value, Unlversk)' of Colk111Wfi at Boulder. Dtparunent of Psychokogy,
1'-1cCord, M., and R. De f-,euMlle, 196:t Empirical denwnstration lha! expected utlity doo,ion ;ld;TJlji, is oot operational. In Fhlmdations of utility Wjd rirk theory with appcalioru, d. f. Stigwn and F. WMiop, editor:, Dordfechl, Holhm; D. 'Reidel. 181-99.
----- 1986. 1-Query equivalents: Reduction of tht cerodmy effect problem in utility asse,s. ment Mwiagemcnj. Sciffllcc 32S6,-60.
McGlothlin, W. H. 1956. Stability of choices among uncertain alttroative,. Affil<'triam Jvurnr.d ofP:rydwlngy 69:604-15.
McKelvey, R. D. and P. OrdelUOOK. 1985. Elections wnh limited info!tldition: A fulfilled lillioM model U1111 oontcmpwmr:ous pnU and endoctelOOtt data tl iofurrmtioo. sotlr-ei JCUTftll of Ecconmic Thtny JS:5 5.
McKelvey, R. D., and T. e. 1990. Public and private infimatioo: All n:perimantal Study of infonmuio flO(ling. Ecmwne1ricJ 5&:1321-J-9.
McNeil B, J.. S. G. Pauker, Mid A, TWTSky, 19&& On the framing of medical d<;i,mm. lo lkrtskm maQJlg: Descrpth't, iwm111iv, and pritMnptiW! mtenldiom. D. E. Bell, H. Raiffa. md A. Tversky, editors, Ounbrkige: Cambridge- !rriernury l'rellS. 562--ll
Meehl. & ft 19:4. Chnirnl wn,.... stamt<al pl'edktivt: A tmwrojal OMfijili rnd a nn'ew of the e t, Minneapolis, Mum: L' it)' of Mi.MWJt. >ress.
Metfl-\$, B. A. L. D. OrOOilet wd M, H. Dimb&um, 1992. A change of j)ffl(CM !hooy (Of romexual tff<ctS and prekw:..e reverws in Mk)' OrganitUm11al He-havutr- ulld Hllldtffl DtsU,11, Frn-terSfs 57:331-69,
Mookherjee, n,, and B. Sopher. 1994. Lemung behavior in in experir(Kutal mawhng pennie! - Games ind &tmomil: Behavior- 7:61-91.
Mo<m, f: and A. Martin. 1990. &ffi:r htmi,r;l,;s fureoooomic SINTUt &p;-ri1wmtdal and wtluhl- Iron evidence. JOW"flal of fhehaviorol D<L!lon Making 3: 175-93.
M . 1980. The .oUab:mrzioo betw<n Otkar Mo!gemern and loot) Vot Neumann Ot me theory of :. lc<mal of f:rooomic Utewtrw 14:805-16.

- M 0. 1979. Some reifications on illitly theory. In *The i,m utility hyp<tmW and the Aiaais paradox*, M. AUas and O. Hagtn, tdikn, Dmd.mdl: D. keiatL MoikT, D. M., and F. Botgid:L 1984. The ronjurntion !'alllcy: A Wk sp,x:ilk . Pers,mality muj &wiat P_{rc} lwi&gy 8ultti>1 10:2 5:t
- Morri&m. D. G. 1967. On the CJOOK&Qy of ptc in Alla.is' p!llldoo. Bthn it>ral Seim&, J2'J7J.,8J.
- Mod'owia, H. 1974. Effectl. of problem reprownwioo and feedback on raJOOal bclmviur In All,ti and Mor!lIM)'pe pmbleim; DtcIJitm Sdencrs. 5'..U5->tt
- Mosteller, F. and P. Nogu. 1951. An e:qerimental tru:MIJtemenf of utilil)', Jourrw i efP ol'tiad Ero,wmy59:311--404.
- Mun=, B. 191W. New rnodeh of &clslon urultr im.,,rtainty: An mlcrpn:llht - Y. t... n P JGurnal of Opm1rlol'la! Remm:li 38:307-17.
- Murphy, A. H. lo/73. A new v<m r pattitium of the probiwiluy . Jau.rnal of App(ied Mttwr-ology 12:59,5-600.
- Murphy, A. ff. and R. L. Wllllldex. 1970. Si:oring roles in probability l.sse@mtnl and evaluation. Jm Psychcwglca 34:2:7
- 1974. Sub_jel,t,le probablltry ft.recasting e.\pentnenis In mctooroioy: Somr pn:hmimary results. Bi,/ktin v f lhlAttUritan M womlotkaJ Scotty 55: 1206-16.
- au. R. F. 1989. Deci:ron analysis wjh indcwmin111e or inoohterm prmbllilitctS, Anna/\$ of Q,era flms &search 19:375-400.
- Neilson, W. 1989. Prospect theory's diSC(mtinuitis wirtmut rrobability wcight.. Ws)lking paper. Texas A & M, Depwtm,ml (If Economics, Cktober.
- Neilson, W. S. In press. A mixed fan hypothesis and illl impllcatios for behav,w rowarda rik. Jmtmai of &onnmk Beh,wior and Organfuiti,m,
- Nelson, R. G. and D. A. Beler. 1989. Subjective probabilties and ,wring rules: Experimental evidence. Amuican Journal of Agricuini.ro/ EC<momin 71 :363-9.
- Nerlove, M. 1958. Adaptive expeciatiollS and cobweb phenomena. Qlllttrterly J m rnal of Ec,-nomicJ 73:227""40.
- Neuringer, A. 1986. Can pt,Optl behave "randomly"!; 11Vll role of feedback. Jour,w/ of Exper-imental Psycholm., G:lntral 115:62-75.
- Niekerwn, R. S., A. Bllddely. and B. Fteffilm. 1987. Are people' em,natei; of what nUlr people know inilutnte,> by what they them,eel<es knuw? Ada F'syl:wfogita 64:24\$-59.
- Nisbett. R., E. Borgida, k. Crandall, and H. Reed. 1976. Popl.00' ittducion: JllforUMitico is not necusarily informative. In CcgnitWn wul soc/a' behavior, J. S. Cturoll and J. W. Feyne.. edit Hillsdale, N.J.: Erlbaum.
- >:isbett, R. E., G. T. Fons, D. R. Lehman, lld P. W. Cheng. 1937. Teaching reasoolng. Science 238:625--31.
- Norman. D. 1988. The psychofogy of twrydlly tfu"l-l. ew Yiltk: Ba1k Books.
- O'Neill. B. 1987. Nonmetinct! itoflthe mimnuu lheocy ctttwo-pmouteWSUm garne., procu,J. Illgs ufihe NaritmaJ A.tadtmy uf&kro:es, VSA 84:2:106-9.
- Payn, j. W. 1976. Task compelitly Mld cotitlgent procei&ng In decuoio. rualung: Anuiforrua-tiou search and !fl,ill>cot ll,illlysi Organ; ! Bflltrim' and Human hr fom,mue: 22:17-44.
- Payne. J. W., J. Betunau. and E. JCMl<ln, 1992. Bcllavloral idtOO>! tesarob: A ooru,tructiv; pro-sing pm-poori. All'Wi ktvtw of l'sy;:/1-dogy 43:67-131.
- I'et=.m.C. R.,MKIL. R. Bea;h. 1967. ManllSll!.llltuulivcillatist.kian P lJultkin 68:29-46.
- Petasoo, S. P. 1993.Fore,a&ingdyna!ruc!andcoo romurlret faidetxe from upttimerttai - marle Jm.maJ of Ectm-amk and i(m 22:169-84.
- Plott, C.R. 1986. Rati<mal choke m perimeJtal mukcts. Jcl.tmai e flllwhws; 59(4;SJOJ-S327.
- Plott, C. R., and L. Wikie. 1982. Profemmtal di w., wtf i - An eiprtimcnud examination of - special of t5 with unexumuy. In Rtflll'Ch in c,qMri-mtmtal UDIWmKS, V. L. Smith, eairor, Grtcnwroh, Coon.: JAJ .
- Pratt, J. W., D. A. Wiw, and R. Zecldiau-. 1979. Pnte in almo!lt .omf)flitive markets. Quan,uly Jm,mal of Eeotwmit-s 93:189-211.
- Prlec, D. 1990. A noowmtJU:-- effed, and Jts implidilloru for MllC iocel! :uoe; i pecte, ! utlity models . Jmmmai e fRisk lbid Uncrtain'y 3:141-59.
- Pt M. G., and P. B&mL! 1948. An mrotw.mtdy {If the aucuoio value of an urucrtain u.,tcome. Ameri-cml J @mai of P gy 61:163-93.
- Quiggin. I. 1932. A theory of anticip!Ued lltitiry. JDW'114' c f&coiwrn< lkhavk>t arid O . titm 3:3'B-43.
- Ram,ey, F. 1931. Tn.th and prob,d<ility. In The fm,t,wmons of l'mlhcctftltin and other l-ogkll es,rays, F. Ranmy, editor, London: Routledge & Kes,m Paut 158<98 Reprinted m H. E. Kybur g and H. E. Smoldei. ediron, Shidit<i in xubj it< probabllity, New York: Wiley. J-964 61-91.
- Rapoport, A., and R. B. Boebel. 1992. Mied strategies in strictly com_)ctive glUrCE A futher test of the minimu hypotheis. Game, arid &onomie Behav;:r 4-:u.i J.
- Rapoon, A., and D. V. Budesru. 1992. G:neratlon <lf raudom terie. in two-,pemm strkdy comperuive games. Jno,rnml of E,rpmimental Ps wJw,gy: Gentro/ 121:352--03.
- Real, Le.she. 1991. Animal choice bl'havior llld the ew\ution of cognitive ardcutecure. Sderv:e 253:981)...86.
- Rl'delmeier, Don, and Daniel Kahneman. 1993. The pail of illlinvasive medical procedure: Paten.ts' real rime and recrospective evaluations of color)(&sd,py. Wodking papir. Unlver-sity of Catifomia at Berkeley, Dcpartneit of Psychofogy.
- Rietz, T. A. 1991. Asset prices and risk allocation In experimental markets. Working p!pler 109. Northwestern University, Department of Finance, August.
- Ritov, I., and J. Baron. 1990. Reluctance Jvatcinate: Omission lWJ and Jllllbiguity. Jounwl of Be/wuviral Decision Making 3:263-n.
- Roberts, H. V. 1963. Risk, amhlguil', and Savage uioim: C<monent. Quarterly Jouma{ ef Economics 77:327-36.
- Roby, T. B. 1965. Belief states: A preliminary empirical study. Ocasion Scenc< laaondory, L. G. Hascom Field.
- Roell, A. 1987. Risk aversion in Quiggin and Yaari's rank-order model of chauce under Wk-ff-nunty. EcoJijmic Juumal 97:143-59.
- Roman, J. 1973. Effects of some probability displ.tys oo choices. Organka:wMl B-havie r and Human Performam,e 9: 1-15.
- Romis. O. L. and J. E. Vales. 1997. Computeots of prnhability judgment 11tcuncy: Individlial consistency and effect, of subjeec-matter and i,rnent method. Organ,dlwlal l!elunfnr and Human lhcisiOrc Prr;muses 40: 1%-218.
- Rms. B. M., and N. l.e,y. 1951. Patltmed pdwiciOllf Qf chance tvenu by childrtn and adultll, PlijCMJiologi& p&rn 4:111-124.
- Ro%, L. D. Greene, and P. Hoos-. 1977. 'the "false, effe:t": An egocemtic billi in social tion ad attribntkjo. Jow-tw/ ti/ bp,rimtnu.f Sni<4 Psrdwfogy 13:219-301.
- Miclacl. and HK'ie Sicrily. 1982. Egocemric biases in availabil;y ad ettribu.60!l in Judg, mem wtdr umertuinry. Heiuisria tmJ mates, D. Kahl!Mru,i, P. Slovie, and A. Tvd"lky, eairors, Cambridge. Cwtbridge Uoi .
- Roth, Alvin E., and MK'hKI W. K. Maloof. 193, Game thoretic lllXlelt and the role (If information in Nrgllinmg: An experimental: Study. Psy;lwkgu:ai R&vie. 86: 574-94.
- Rubinstein. Anel 1988. Similarity and dccnkirt-makina under rnk (ls Ulert t utility thoory resolutiwi to the All,tis pararlox?). l'W'ni/ c f&tm,JmtC Thoory 46:US--S).
- RUS'liell, T., 1100 R. Th;k-. 1985. The rel't'lllOOof quasi-ratiorudicy ill ma,d,e<. Amtriam lJ.b., rwmk Re"ew 15: 10!1-82.
- Rmro. J. E., and P. J. H. Schoemaker. 198), Dtl'iru>! Imp!: The 4fl barrws U brilliant! .de(l'l-ltllang Md iww w ovcn:ome them. New York
- sma, z., u. Segal, and A. Spiva Sepembe: 1990a. Preference tevt!tlll and m:ill spectro-inility behavior. A.m,rica, Economic RFV!tw im922-JO.

- Safra, Z., U. Segal, and A. Spivak. 1990b. The Becker-DeGroot-Manchak mechanism and non-expected utility: A testable approach. *Journal of Risk and Uncertainty* 3:177-90.
- Samuelson, P. 1952. Probability, utility, and the independence axiom. *Econometrica* 20: 670-8.
- Samuelson, W., and R. Zeckhauser. 1988. Status quo bias in decision making. *Journal of Risk and Uncertainty* 1:7-59.
- Sanders, A. F. 1968. Choice among bets and revision of opinion. *Acta Psychologica* 28:76-83.
- Sarin, R. K., and M. Weber. 1993. The effect of ambiguity in a market setting. *Management Science* 39:602-15.
- Sarin, R. K., and R. L. Winkler. 1992. Ambiguity and decision modeling: A preference-based approach. *Journal of Risk and Uncertainty* 5:389-407.
- Savage, L. J. 1954. *The foundations of statistics*. New York: Wiley.
- Scharfstein, D. S., and J. Stein. 1990. Herd behavior and investment. *American Economic Review* 80:465-79.
- Schelling, T. 1984. Self-command in practice, in policy, and in a theory of rational choice. *American Economic Review* 74: 1-11.
- Schkade, D. A., and E. J. Johnson. June 1989. Cognitive processes in preference reversals. *Organizational Behavior and Human Decision Processes* 44:203-31.
- Schmalensee, R. 1976. An experimental study of expectation formation. *Econometrica* 44:17-41.
- Schmeidler, D. 1989. Subjective probability and expected utility without additivity. *Econometrica* 57:571-17.
- Schoemaker, P. J. H. 1990. Are risk attitudes related across domains and response modes? *Management Science* 36: 1451-63.
- Schoner, A., and Y. M. Braunstein. 1981. Economic search: An experimental study. *Economic Inquiry* 19:1-25.
- Schweitzer, Maurice. In press. Disentangling status quo and omission effects: Experimental analysis. *Organization Behavior and Human Decision Processes*.
- Segal, U. 1987a. The Ellsberg Paradox and risk aversion: An anticipated utility approach. *International Economic Review* 28:175-202.
- . 1987b. Some remarks on Quiggin's anticipated utility theory. *Journal of Economic Behavior and Organization* 8:145-54.
- . 1988. Does the preference reversal phenomenon necessarily contradict the independence axiom? *American Economic Review* 78:233-36.
- . 1989. Anticipated utility: A measure representation approach. *Annals of Operations Research* 19:359-73.
- . 1990. Two-stage lotteries without the reduction axiom. *Econometrica* 58:349-77.
- Sbafir, E. In press. Choosing versus rejecting: Why some options are both better and worse than others. *Cognitive Psychology*.
- Shelley, M. K. In press. Outcome signs, question frames, and discount rates. *Management Science*.
- Simon, H. A. 1955. A behavioral model of rational choice. *Quarterly Journal of Economics* 69:99-118.
- Slovic, P. 1969. Manipulating the attractiveness of a gamble without changing its expected value. *Journal of Experimental Psychology* 79:139-45.
- Slovic, P., and S. Lichtenstein. 1968. The relative importance of probabilities and payoffs in risk taking. *Journal of Experimental Psychology* 78: 1-18.
- Slovic, P., S. Lichtenstein, and W. Edwards. 1965. Boredom induced changes in preferences among bets. *American Journal of Psychology* 78:208-17.
- Slovic, P., and A. Tversky. 1974. Who accepts Savage's axiom? *Behavioral Science* 19:368-73.
- Smith, C. A. B. 1961. Consistency in statistical inference and decision. *Journal of Royal Statistical Society, series B* 23: 1-37.
- Smith, V. L. 1979. Measuring nonmonetary utilities in uncertain choices; The Ellsberg urn. *Quarterly Journal of Economics* 83:324-29.

- . 1991. Rational choice: The contrast between economics and psychology. *Journal of Political Economy* 99:811-91.
- Smith, V. L., G. L. Suchanek, and A. W. Williams. 1988. Bubbles, crashes, and endogenous expectations in experimental spot asset markets. *Econometrica* 56:1119-51.
- Smith, V. L., and J. M. Walker. 1993. Monetary rewards and decision costs in experimental economics. University of Arizona, Department of Economics.
- Snizek, J. A., and T. Buckley. 1991. Confidence depends on level of aggregation. *Journal of Behavioral Decision Making* 4:263-72.
- Snizek, J. A., and R. A. Henry. 1989. Accuracy and confidence in group judgment. *Organizational Behavior and Human Decision Processes* 43:1-28.
- Soll, Jack B. 1993. Determinants of miscalibration and over/under-confidence: The intention between random noise and the ecology. Working paper. University of Chicago, Center for Decision Research.
- Stanner, C. In press. Testing new theories of choice under uncertainty using the common consequence effect. *Review of Economic Studies*.
- Starmer, C., and R. Sugden. 1989a. Violations of the independence axiom in common ratio problems: An experimental test of some competing hypotheses. *Annals of Operations Research* 19:79-102.
- . 1989b. Probability and juxtaposition effects: An experimental investigation of the common ratio effect. *Journal of Risk and Uncertainty* 2: 159-78.
- . 1989c. Experimental evidence of the impact of regret on choice under uncertainty. Discussion paper no. 23. University of East Anglia, Economics Research Centre.
- . 1991. Does the random-lottery incentive system elicit true preferences? An experimental investigation. *American Economic Review* 81(4):971-78.
- . 1993. Testing for juxtaposition and event-splitting effects. *Journal of Risk and Uncertainty* 6:235-54.
- Strotz, R. H. 1953. Cardinal utility. *American Economic Review (Papers and Proceedings)* 43:384-405.
- Sugden, R. 1986. New developments in the theory of choice under uncertainty. *Bulletin of Economic Research* 38: 1-24.
- Sunder, S. 1991. Market for information: Experimental evidence. Working paper. Carnegie-Mellon University, GSIA.
- Taylor, S. E., and J. D. Brown. 1988. Illusion and well-being: A social psychological perspective on mental health. *Psychological Bulletin* 103:193-210.
- Thaler, R. H. 1980. Toward a positive theory of consumer choice. *Journal of Economic Behavior and Organization* 1:39-60.
- . 1981. Some empirical evidence on dynamic inconsistency. *Economic Letters* 8:201-7.
- . 1985. Mental accounting and consumer choice. *Marketing Science* 4:199-214.
- . 1987. The psychology of choice and the assumptions of economics. In *Laboratory experimentation in economics: Six points of view*, A. E. Roth, editor, Cambridge: Cambridge University Press.
- . 1992. *The winner's curse: Anomalies and paradoxes of economic life*. New York: Free Press.
- Thaler, R. H., and E. J. Johnson. 1990. Gambling with the house money and trying to break even: The effects of prior outcomes on risky choice. *Management Science* 36: 643-60.
- Thaler, R. H., and H. M. Shefrin. 1981. An economic theory of self-control. *Journal of Political Economy* 89:392-410.
- Thurstone, L. L., C. Coombs, and R. Davis. 1954. *Decision processes*. New York: Wiley.
- Tobin, J. 1958. Liquidity preference as behavior toward risk. *Review of Economic Studies* 25:65-86.
- Toda, Masanao. 1963. Measurement of subjective probability distribution. Report 3. State College, Pennsylvania, Institute for Research, Division of Mathematical Psychology.

